

1. Why we need manager?

We need manager for below:

- To assign task
- To get new project
- Leave management
- Work management
- Promotions
- Adding more members
- Monitoring

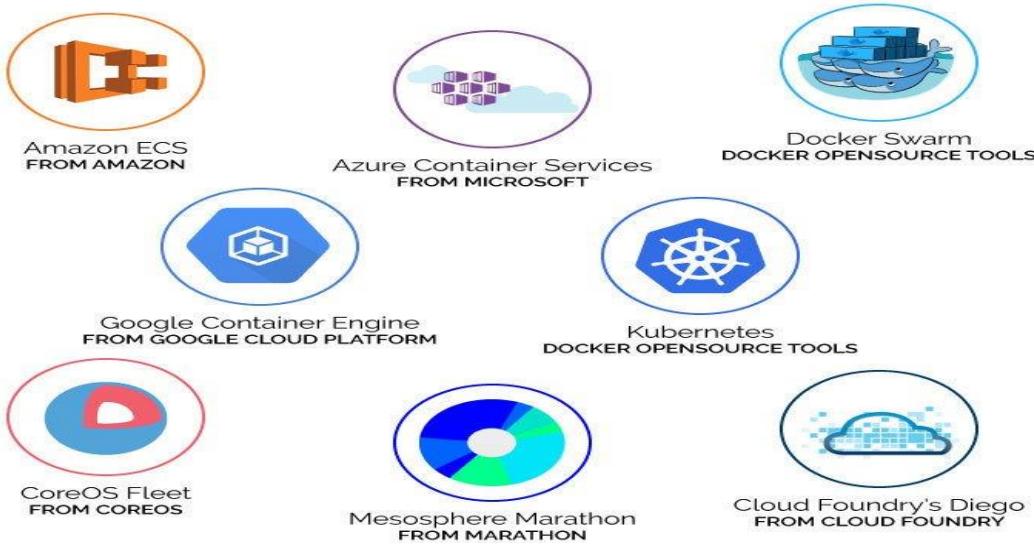
2. What is container orchestration? Why we need container orchestration?

- Container orchestration automatically provisions, deploys, scales, and manages containerized applications without worrying about the underlying infrastructure.
- Container orchestration automates the deployment, management, scaling, and networking of containers.
- Hosting container based application in different environment brings many challenges to infra & devops team such as below:
 - Which nodes suit the container requirement?
 - Maintain the desire state?
 - How to share load (user request)?
 - Scalability?
 - Auto Healing?
 - Auto scaling?
 - Zero downtime upgrades/rollback
 - Container & nodes communication
 - Security—RBAC & certificates
- All the above features are provided by container orchestration tool, that's why we need this.

3. How many container orchestration software's available in market?

- In market below are container orchestration software's available in market
- Rancher, Nomad, Mesos, Docker Swarm, kubernetes.
- Here kubernetes is open source
- Some of the organization taken kubernetes open source code & added their own flavors & released their own kubernetes at enterprises level & gives support for any issue.
- Amazon AWS cloud give you EKS (Elastic Kubernetes Service)
- Azure gives you AKS(Azure Kubernetes service)
- Google cloud gives you GKE(Google Kubernetes Engin)
- VMware gives you VMware Tanzu
- Redhat gives you Openshift

Tools of Container Orchestration

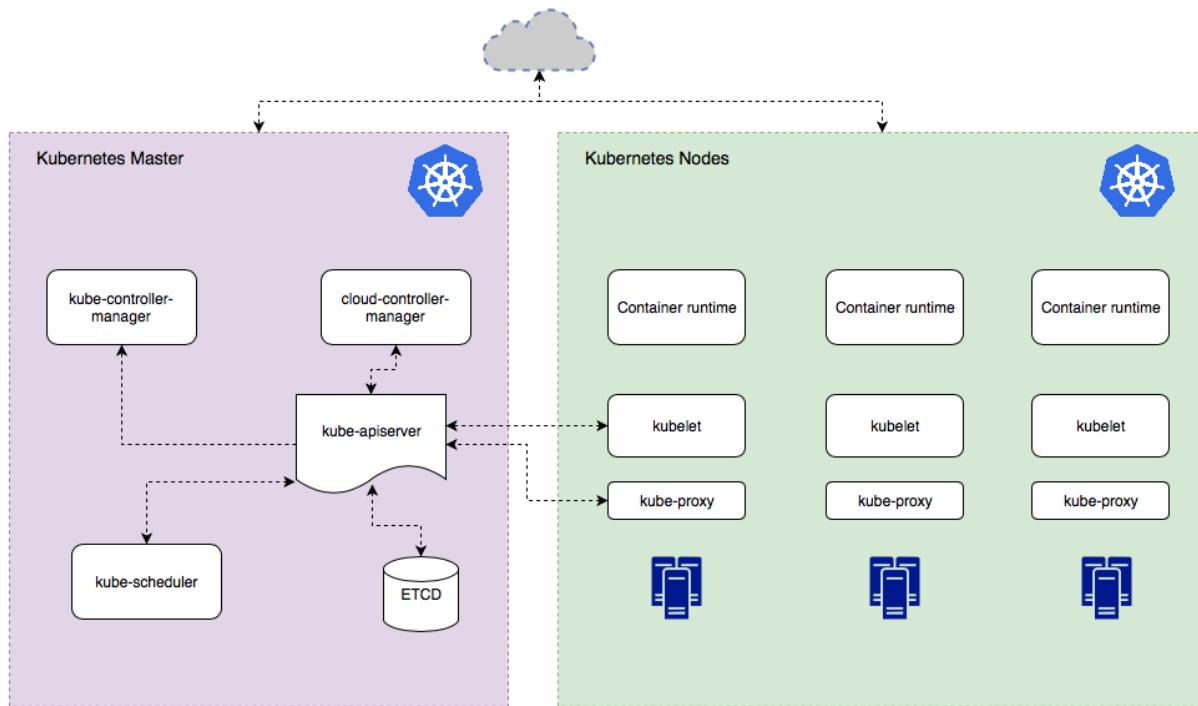


4. What is Kubernetes? What all components are available?

- Kubernetes (K8s in short) is an open-source container orchestration platform introduced by Google in 2014.
- It is a successor of Borg, Google's in-house orchestration system that accumulated over a decade of the tech giant's experience of running large enterprise workloads in production.
- In 2014, Google decided to further container ecosystem by sharing Kubernetes with the cloud native community.
- Kubernetes became the first graduated project of the newly created Cloud Native Community Foundation (CNCF), an organization conceived by Google and the Linux Foundation as the main driver of the emerging cloud native movement.
- Kubernetes work on **Master-Node** model.
- **Kubectl** is keyword which is used in Kubernetes to give instructions to master.
- **Master** has below important components

Component	Description
Etcd	This is a data store used by Kubernetes to store all information about the cluster. It's critical for keeping everything up to date
kube-apiserver	This component exposes Kubernetes API to users allowing them to create API resources, run applications, and configure various parameters of the cluster. It is key components which accept all requests.
kube-scheduler	This component is responsible for scheduling user workloads on the right infrastructure. When scheduling applications, kube-scheduler considers various factors such as available node resources, node health, and availability, as well as user-defined constraints.
kube-controller-manager	The component that manages API objects created by users. It makes sure that the actual state of the cluster always matches the desired state

cloud-controller-manager	This component embraces various controllers, all of which interact with the cloud providers' APIs.
--------------------------	--



- On node (worker node) below components available
- **CRI** : To run container on worker node we need container runtime(CRI), it can be docker or container-d or Crio
- **Kubelet**: It is an agent of API server to get command from Master. And report back to API server.
- If you want to run container on kubernetes below will be steps:
 - We give request through **KUBECTL**
 - **KUBECTL** goes to **API server**, & give request to API server.
 - **API server** create object.
 - **Scheduler** keeps on watching API server & checks any request came to **API server** for container creation.
 - **Scheduler** find which node is appropriate for container creation & it tells to **API server**.
 - **API server** talks to **KUBLET** of worker node to run container.
 - **KUBLET** accept request & uses **CRI**, which will download image, run container & report back to **API Server**.
 - **API server** save that information (image, memory, CPU, number of replica) about container in **ETCD**
 - **Controller manager** keep on watching instances of container and manages desire state.
- We can have **KUBELET & CRI** in master machine also to run container on master machine.

5. How many ways we can install Kubernetes?

- We can install kubernetes in 2 ways
- **Hard way**: In hard way we need to install each & every component one by one by downloading.
- **Kubeadm way**: It is easiest way to install kubernetes, it is utility by using we get all component.

- We need minimum 2 CPU & 4 GB RAM to setup kubernetes.
- It is recommended to install Kubernetes on Linux only.

6. How to install Kubernetes?

- First we need to install any container runtime (CRI) from below

Runtime	Path to Unix domain socket
Containerd	unix:///var/run/containerd/containerd.sock
CRI-O	unix:///var/run/crio/crio.sock
Docker Engine (using cri-dockerd)	unix:///var/run/cri-dockerd.sock

- After that we need to install component kubeadm, kubelet and kubectl.
 - **kubeadm:** the command to bootstrap the cluster.
 - **kubelet:** the component that runs on all of the machines in your cluster and does things like starting pods and containers.
 - **kubectl:** the command line util to talk to your cluster.
- Login into manager/master install below
 - Docker, kubeadm,kubelet,kubectl
 - Initialize master get token
 - Syntax:> sudo kubeadm init --ignore-preflight-errors=all

```
root@manager:/#
root@manager:/# sudo kubeadm init --ignore-preflight-errors=all
I0105 02:24:31.302690 8254/ version.go:255] remote version is much newer: v1.29.0; falling back to: stable-1.23
[root@manager ~]# [init] Using Kubernetes version: v1.23.17
[preflight] Running pre-flight checks
    [WARNING Port-6443]: Port 6443 is in use
    [WARNING Port-10259]: Port 10259 is in use
    [WARNING Port-10257]: Port 10257 is in use
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:
kubeadm join 10.128.0.2:6443 --token cbg23d.gzrcn0pqsm1jqbh9 \
  --discovery-token-ca-cert-hash sha256:fc12deeeaa5818555ea3fb4c41a34101f86436bb6022f947124a7a0f67e47686
root@manager:/#
```

- Setup API server home path, so that kubectl can communicate with API server
- In kubernetes home directory there is file “/etc/kubernetes/admin.conf” in which all API server details are available , so if get invoke get node command without this file path it will give error
- Syntax:> kubectl get nodes
- So to resolve this we need to execute below command
 - Syntax:>kubectl get nodes --kubeconfig /etc/kubernetes/admin.conf

```
root@manager:/# kubectl get nodes
The connection to the server localhost:8080 was refused - did you specify the right host or port?
root@manager:/#
```

Classification: Public

```
root@manager:/# kubectl get nodes --kubeconfig /etc/kubernetes/admin.conf
NAME      STATUS    ROLES           AGE     VERSION
manager   NotReady control-plane,master   13h    v1.23.10
root@manager:/#
```

- o To start using your cluster, you need to run the following as a regular user:

```
>mkdir -p $HOME/.kube
>sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
>sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
root@manager:/# mkdir -p $HOME/.kube
root@manager:/# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
root@manager:/# sudo chown $(id -u):$(id -g) $HOME/.kube/config
root@manager:/# kubectl get nodes
NAME      STATUS    ROLES           AGE     VERSION
manager   NotReady control-plane,master   13h    v1.23.10
root@manager:/#
```

```
devops@manager:/# kubectl get nodes
The connection to the server localhost:8080 was refused - did you specify the right host or port?
devops@manager:/# mkdir -p $HOME/.kube
devops@manager:/# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
devops@manager:/# sudo chown $(id -u):$(id -g) $HOME/.kube/config
devops@manager:/# kubectl get nodes
NAME      STATUS    ROLES           AGE     VERSION
manager   NotReady control-plane,master   13h    v1.23.10
devops@manager:/#
```

- o Alternatively, if you are the root user, you can run:
- o Syntax >export KUBECONFIG=/etc/kubernetes/admin.conf
- o Kubernetes does not give you networking feature, we need to install separately
- o Due to networking not enabled master is showing in not ready state.
- o You should need to install networking driver -- Weave/flannel/canal/calico etc... by using below command
- o > kubectl apply -f <https://raw.githubusercontent.com/projectcalico/calico/v3.24.1/manifests/calico.yaml>
- o Validate same using command: >kubectl get nodes

```
root@manager:/# kubectl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.24.1/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/irpresaervations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
root@manager:/# kubectl get nodes
NAME      STATUS    ROLES           AGE     VERSION
manager   NotReady control-plane,master   23h    v1.23.10
root@manager:/# kubectl get nodes
NAME      STATUS    ROLES           AGE     VERSION
manager   Ready    control-plane,master   23h    v1.23.10
root@manager:/#
```

Classification: Public

- Now we can join any number of worker nodes by running the following on each as root:
- Login into worker/node install below
 - Docker, kubeadm,kubelet,kubectl
 - Apply token by using below command
 - > kubeadm join 10.128.0.2:6443 --token cbg23d.gzrcn0pqsm1jqbh9 --discovery-token-ca-cert-hash sha256:fc12deea5818555ea3fb4c41a34101f86436bb6022f947124a7a0f67e47686

```
root@worker-1:/#
root@worker-1:# kubeadm join 10.128.0.2:6443 --token cbg23d.gzrcn0pqsm1jqbh9 --discovery-token-ca-cert-hash sha256:fc12deea5818555ea3fb4c41a34101f86436bb6022f947124a7a0f67e47686
[preflight] Running pre-flight checks
    [WARNING SystemVerification]: this Docker version is not on the list of validated versions: 24.0.7. Latest
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
W0105 14:49:57.098984 67736 utils.go:69] The recommended value for "resolvConf" in "KubeletConfiguration" is: /run/systemd/resolve/resolv.conf
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

root@worker-1:#

root@worker-2:/#
root@worker-2:# kubeadm join 10.128.0.2:6443 --token cbg23d.gzrcn0pqsm1jqbh9 --discovery-token-ca-cert-hash sha256:fc12deea5818555ea3fb4c41a34101f86436bb6022f947124a7a0f67e47686
[preflight] Running pre-flight checks
    [WARNING SystemVerification]: this Docker version is not on the list of validated versions: 24.0.7. Latest
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
W0105 14:51:11.591326 66565 utils.go:69] The recommended value for "resolvConf" in "KubeletConfiguration" is: /run/systemd/resolve/resolv.conf
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

root@worker-2:#
```

- We can check worker who has join cluster using below command
- Syntax> kubectl get nodes

```
root@manager:/# kubectl get nodes
NAME     STATUS   ROLES      AGE     VERSION
manager  Ready    control-plane,master  23h    v1.23.10
worker-1 Ready    <none>    2m54s   v1.23.10
worker-2 Ready    <none>    100s   v1.23.10
root@manager:/# sudo kubeadm token create --print-join-command
kubeadm join 10.128.0.2:6443 --token j2iva5.dpobern3v30xrfw3 --discovery-token-ca-cert-hash sha256:fc12deea5818555ea3fb4c41a34101f86436bb6022f947124a7a0f67e47686
root@manager:/#
```

- Run Below on Master Node to get join token
- Syntax> sudo kubeadm token create --print-join-command

7. What is Pod in Kubernetes?

- In kubernetes we don't run container independently, instead we run container inside POD.
- Pod runs our container in kubernetes
- Pod maintains resources, IP address, storege etc of container.
- It makes sure that container will be up & running always.
- Pod is security layer on top of container.
- If something happened to container pod tries to bring up that container.
- In kubernetes anything we create is called as object.
- Pod is first & smallest object, that we run in kubernetes.

8. What is use of kubectl? What are its different commands?

- In order to perform any task in kubernetes we use client command “**kubectl**”
- Command to start pod
 - Syntax :> kubectl run <pod name> --image <image name>
 - Ex:> kubectl run pod1 --image nginx
- Command to check pods
 - Syntax :> kubectl get pods
 - Command to check pods with wider output, which gives you details like name, status, ip, node etc.
 - Syntax :> kubectl get pods -o wide

```
root@manager:/#
root@manager:/# kubectl run pod1 --image nginx
pod/pod1 created
root@manager:/# kubectl get pods
NAME READY STATUS RESTARTS AGE
pod1 1/1 Running 0 64s
root@manager:/# kubectl get pods -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
pod1 1/1 Running 0 2m27s 192.168.133.193 worker-2 <none> <none>
root@manager:/#
```

- Command to check details of a pod, which gives detail like name, namespace, node where it is running, its container details, volumes, events, etc....
- Syntax :> kubectl describe pod <pod name>
- Syntax :> kubectl describe pod pod1

```
root@manager:/#
root@manager:/# kubectl get pods -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
pod1 1/1 Running 0 2m27s 192.168.133.193 worker-2 <none> <none>
root@manager:/# kubectl describe pod pod1
Name: pod1
Namespace: default
Priority: 0
Node: worker-2/10.128.0.5
Start Time: Sun, 07 Jan 2024 06:53:54 +0000
Labels: run=nginx
Annotations: cnr.projectcalico.org/containerID: 5b6c3ce9a28eeecb96bb205ae54a3defefe3529041d1c3e0dbe09f743d35ba229
            cnr.projectcalico.org/podIP: 192.168.133.193/32
            cnr.projectcalico.org/podIPs: 192.168.133.193/32
Status: Running
IP: 192.168.133.193
IPs:
  IP: 192.168.133.193
Containers:
  pod1:
    Container ID: docker://70392234487b8e3cddf0ba6e87a7a24206ac77eaf97ea83d53f2f7a48ca6506c
    Image: nginx
    Image ID: docker-pullable://nginx@sha256:2bdca9f2f8ae8d8dc50ed00f2ee56d00385c6f8bc8a8b320d0a294d9e3b49026
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Sun, 07 Jan 2024 06:53:56 +0000
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-vgqfv (ro)
Conditions:
  Type Status
  Initialized True
  Ready True
  ContainersReady True
  PodsScheduled True
```

Classification: Public

```
c1 root@manager: /  
Volumes:  
  kube-api-access-vgqfv:  
    Type:          Projected (a volume that contains injected data from multiple sources)  
    TokenExpirationSeconds: 3607  
    ConfigMapName:   kube-root-ca.crt  
    ConfigMapOptional: <nil>  
    DownwardAPI:    true  
  QoS Class:      BestEffort  
  Node-Selectors: <none>  
  Tolerations:    node.kubernetes.io/not-ready:NoExecute op=Exists for 300s  
                  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s  
Events:  
  Type  Reason  Age   From            Message  
  ----  -----  --  --  --  
  Normal Scheduled  13m  default-scheduler  Successfully assigned default/pod1 to worker-2  
  Normal Pulling   13m  kubelet         Pulling image "nginx"  
  Normal Pulled    13m  kubelet         Successfully pulled image "nginx" in 488.167684ms  
  Normal Created   13m  kubelet         Created container pod1  
  Normal Started   13m  kubelet         Started container pod1  
root@manager:/#
```

- Command to delete single pod
 - Syntax :> kubectl delete pod <pod name>
 - Ex :> kubectl delete pod pod1
- Command to delete all pods
 - Syntax :> kubectl delete pods --all

```
c1 root@manager: /  
root@manager:/# kubectl get pods  
NAME    READY  STATUS    RESTARTS   AGE  
pod1   1/1    Running   0          31m  
root@manager:/# kubectl delete pod pod1  
pod "pod1" deleted  
root@manager:/# kubectl get pods  
No resources found in default namespace.  
root@manager:/# kubectl run pod2 --image alpine  
pod/pod2 created  
root@manager:/# kubectl get pods -o wide  
NAME    READY  STATUS    RESTARTS   AGE     IP           NODE    NOMINATED NODE  READINESS GATES  
pod2   0/1    CrashLoopBackOff 1 (6s ago) 10s   192.168.133.194  worker-2  <none>        <none>  
root@manager:/# kubectl get pods -o wide  
NAME    READY  STATUS    RESTARTS   AGE     IP           NODE    NOMINATED NODE  READINESS GATES  
pod2   0/1    CrashLoopBackOff 1 (8s ago) 12s   192.168.133.194  worker-2  <none>        <none>  
root@manager:/# kubectl get pods -o wide  
NAME    READY  STATUS    RESTARTS   AGE     IP           NODE    NOMINATED NODE  READINESS GATES  
pod2   0/1    CrashLoopBackOff 1 (11s ago) 15s   192.168.133.194  worker-2  <none>        <none>  
root@manager:/# kubectl get pods -o wide  
NAME    READY  STATUS    RESTARTS   AGE     IP           NODE    NOMINATED NODE  READINESS GATES  
pod2   0/1    CrashLoopBackOff 2 (3s ago) 19s   192.168.133.194  worker-2  <none>        <none>  
root@manager:/# kubectl delete pods --all  
pod "pod2" deleted  
root@manager:/# kubectl get pods -o wide  
No resources found in default namespace.  
root@manager:/# kubectl get pods  
No resources found in default namespace.  
root@manager:/#
```

- Command to enter inside POD
 - Syntax :> kubectl exec -it <pod name> --bash
 - Ex :> kubectl exec -it pod3 -- bash
- Command to exit from POD > exit
- We can use CTRL+PQ Command to exit from POD

```
c1 root@manager: /  
root@manager:/# kubectl exec -it pod3 -- bash  
root@pod3: # exit  
exit  
root@manager:/#
```

Classification: Public

- Command to check log of pod
- Syntax :> kubectl logs <pod name>
- Ex :> kubectl logs pod3
- Syntax :> kubectl logs -f <pod name>
- Ex :> kubectl logs -f pod3

```
root@manager:/# kubectl logs pod3
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/01/11 12:55:06 [notice] 1#1: using the "epoll" event method
2024/01/11 12:55:06 [notice] 1#1: nginx/1.25.3
2024/01/11 12:55:06 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/01/11 12:55:06 [notice] 1#1: OS: Linux 5.15.0-1047-gcp
2024/01/11 12:55:06 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/01/11 12:55:06 [notice] 1#1: start worker processes
2024/01/11 12:55:06 [notice] 1#1: start worker process 29
2024/01/11 12:55:06 [notice] 1#1: start worker process 30
root@manager:/# kubectl logs -f pod3
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2024/01/11 12:55:06 [notice] 1#1: using the "epoll" event method
2024/01/11 12:55:06 [notice] 1#1: nginx/1.25.3
2024/01/11 12:55:06 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2024/01/11 12:55:06 [notice] 1#1: OS: Linux 5.15.0-1047-gcp
2024/01/11 12:55:06 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2024/01/11 12:55:06 [notice] 1#1: start worker processes
2024/01/11 12:55:06 [notice] 1#1: start worker process 29
2024/01/11 12:55:06 [notice] 1#1: start worker process 30
```

- If we run OS container image in pod it's status will go in "CrashLoopBackOff", because in plain OS container there is no back ground running process available, so it is getting stopped & pod is again & again trying to bring it up.
- So in kubernetes we never run OS container in pods.
- Command to get all resources available on kubernetes
- Syntax:> kubectl api-resources

```
root@manager:/# kubectl api-resources
NAME          SHORTNAMES   APIVERSION      NAMESPACED   KIND
bindings      v1           v1              true         Binding
componentstatuses   cs          v1              false        ComponentStatus
configmaps     cm          v1              true         ConfigMap
endpoints      ep          v1              true         Endpoints
events         ev          v1              true         Event
limits         limits      v1              false        LimitRange
namespaces     ns          v1              false        Namespace
nodes          no          v1              false        Node
persistentvolumeclaims   pvc         v1             true        PersistentVolumeClaim
persistentvolumes    pv          v1             false        PersistentVolume
pods          po          v1              true         Pod
podtemplates   rc          v1              true         PodTemplate
replicationcontrollers   rc          v1             true        ReplicationController
resourcequotas    quota      v1              true         ResourceQuota
secrets        secrets     v1              true         Secret
serviceaccounts sa          v1              true         ServiceAccount
services       svc         v1              true         Service
mutatingwebhookconfigurations   admissionregistration.k8s.io/v1   false        MutatingWebhookConfiguration
validatingwebhookconfigurations   admissionregistration.k8s.io/v1   false        ValidatingWebhookConfiguration
customresourcedefinitions      crd,crds   apiextensions.k8s.io/v1   false        CustomResourceDefinition
apiservices      apiregistration.k8s.io/v1   false        APIService
controllerrevisions   controllerrevision   apps/v1      true         ControllerRevision
daemonsets      ds          apps/v1      true         DaemonSet
deployments      deploy      apps/v1      true         Deployment
replicasetss    rs          apps/v1      true         ReplicaSet
statefulsets     sts         apps/v1      true         StatefulSet
tokenreviews     authentication.k8s.io/v1   false        TokenReview
localsubjectaccesreviews   authorization.k8s.io/v1   true         LocalSubjectAccessReview
selfsubjectaccesreviews   authorization.k8s.io/v1   false        SelfSubjectAccessReview
selfsubjectrulereviews   authorization.k8s.io/v1   false        SelfSubjectRulesReview
subjectaccesreviews   authorization.k8s.io/v1   false        SubjectAccessReview
horizontalpodautoscalers   hpa         autoscaling/v2  true         HorizontalPodAutoscaler
cronjobs        cj          batch/v1     true         CronJob
jobs            batch       batch/v1     true         Job
certificatesigningrequests   csr          certificates.k8s.io/v1   false        CertificateSigningRequest
leases          coordination.k8s.io/v1   true         Lease
bgpconfigurations   crd.projectcalico.org/v1   false        BGPConfiguration
bgppeers        crd.projectcalico.org/v1   false        BGPPeer
blockaffinities   crd.projectcalico.org/v1   false        BlockAffinity
```

- Command to get all resources api version in kubernetes

- Syntax:> kubectl api-versions

```
root@manager:/# kubectl api-versions
admissionregistration.k8s.io/v1
apiextensions.k8s.io/v1
apiregistration.k8s.io/v1
apps/v1
authentication.k8s.io/v1
authorization.k8s.io/v1
autoscaling/v1
autoscaling/v2
autoscaling/v2beta1
autoscaling/v2beta2
batch/v1
batch/v1beta1
certificates.k8s.io/v1
coordination.k8s.io/v1
crd.projectcalico.org/v1
discovery.k8s.io/v1
discovery.k8s.io/v1beta1
events.k8s.io/v1
events.k8s.io/v1beta1
flowcontrol.apiserver.k8s.io/v1beta1
flowcontrol.apiserver.k8s.io/v1beta2
networking.k8s.io/v1
node.k8s.io/v1
node.k8s.io/v1beta1
policy/v1
policy/v1beta1
rbac.authorization.k8s.io/v1
scheduling.k8s.io/v1
storage.k8s.io/v1
storage.k8s.io/v1beta1
v1
root@manager:/#
```

- Command to get explanation of any resources/ component of kubernetes
- Syntax:> kubectl explain <component/component name>
- Ex:> kubectl explain pods
- Ex:> kubectl explain namespace

```
root@manager:/# kubectl explain pods
KIND: Pod
VERSION: v1

DESCRIPTION:
Pod is a collection of containers that can run on a host. This resource is
created by clients and scheduled onto hosts.

FIELDS:
apiVersion <string>
APIVersion defines the versioned schema of this representation of an
object. Servers should convert recognized schemas to the latest internal
value, and may reject unrecognized values. More info:
https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources

kind <string>
Kind is a string value representing the REST resource this object
represents. Servers may infer this from the endpoint the client submits
requests to. Cannot be updated. In CamelCase. More info:
https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds

metadata <Object>
Standard object's metadata. More info:
https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#metadata

spec <Object>
Specification of the desired behavior of the pod. More info:
https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status

status <Object>
Most recently observed status of the pod. This data may not be up to date.
Populated by the system. Read-only. More info:
https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status
```

Classification: Public

```
root@manager:/# kubectl explain namespace
KIND:     Namespace
VERSION:  v1
DESCRIPTION:
  Namespace provides a scope for Names. Use of multiple namespaces is
  optional.
FIELDS:
  apiVersion <string>
    APIVersion defines the versioned schema of this representation of an
    object. Servers should convert recognized schemas to the latest internal
    value, and may reject unrecognized values. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#resources

  kind <string>
    Kind is a string value representing the REST resource this object
    represents. Servers may infer this from the endpoint the client submits
    requests to. Cannot be updated. In CamelCase. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#types-kinds

  metadata <Object>
    Standard object's metadata. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#metadata

  spec <Object>
    Spec defines the behavior of the Namespace. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status

  status <Object>
    Status describes the current status of a Namespace. More info:
    https://git.k8s.io/community/contributors/devel/sig-architecture/api-conventions.md#spec-and-status
```

- We can get inner object detail also using explain command
- Ex:> kubectl explain pod.spec
- Ex:> kubectl explain pod.spec.containers
- EX:> kubectl explain pod.spec.containers.resources

```
root@manager:/# kubectl explain pod.spec.containers.resources
KIND:     Pod
VERSION:  v1
RESOURCE: resources <Object>

DESCRIPTION:
  Compute Resources required by this container. Cannot be updated. More info:
  https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/
  ResourceRequirements describes the compute resource requirements.

FIELDS:
  limits      <map[string]string>
    Limits describes the maximum amount of compute resources allowed. More
    info:
    https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/
  requests    <map[string]string>
    Requests describes the minimum amount of compute resources required. If
    Requests is omitted for a container, it defaults to Limits if that is
    explicitly specified, otherwise to an implementation-defined value. More
    info:
    https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/
```

- Command to edit pod
- Syntax :> kubectl edit <pod name>
- Ex :> kubectl edit pod3
- It will open YML file in which we can edit & change image to new image

```
c:~ root@manager: /
root@manager:/# kubectl get pods
NAME    READY    STATUS    RESTARTS   AGE
pod1   1/1     Running   1 (35m ago)  42h
pod2   1/1     Running   1 (5h31m ago) 34h
pod3   1/1     Running   1 (35m ago)  10h
root@manager:/# kubectl edit pod3
error: the server doesn't have a resource type "pod3"
root@manager:/# kubectl edit pod pod3
pod/pod3 edited
root@manager:/# kubectl get pods
NAME    READY    STATUS    RESTARTS   AGE
pod1   1/1     Running   1 (37m ago)  42h
pod2   1/1     Running   1 (5h33m ago) 34h
pod3   1/1     Running   2 (6s ago)   10h
root@manager:/#
```

Classification: Public

```
node.kubernetes.io/unreachable:NoExecute op=EXISTS for 300s
Events:
Type    Reason     Age   From      Message
----    -----    --   ----      -----
Normal  SandboxChanged  36m (x2 over 37m)  kubelet  Pod sandbox changed, it will be killed and re-created.
Normal  Pulling      36m   kubelet  Pulling image "nginx"
Normal  Pulled       36m   kubelet  Successfully pulled image "nginx" in 380.315077ms
Normal  Killing      77s   kubelet  Container pod3 definition changed, will be restarted
Normal  Created      76s (x2 over 36m)  kubelet  Created container pod3
Normal  Started      76s (x2 over 36m)  kubelet  Started container pod3
Normal  Pulling      76s   kubelet  Pulling image "httpd"
Normal  Pulled       76s   kubelet  Successfully pulled image "httpd" in 416.25632ms
root@manager:/#
```

- Command to delete everything i.e. all pods, controllers & services from kubernetes
- Syntax:> kubectl delete all --all

```
root@manager-1:/# kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
rs1-2w85d  1/1     Running   0          10h
rs1-j4f98  1/1     Running   0          10h
rs1-wrwjv  1/1     Running   0          10h
root@manager-1:/# kubectl get svc
NAME        TYPE      CLUSTER-IP      EXTERNAL-IP   PORT(S)    AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP    11h
svc1        ClusterIP  10.106.194.48 <none>        2345/TCP   6h59m
svc2        ClusterIP  10.97.106.220 <none>        2345/TCP   6h54m
root@manager-1:/# kubectl get rc
No resources found in default namespace.
root@manager-1:/# kubectl get rc
No resources found in default namespace.
root@manager-1:/# kubectl get rs
NAME    DESIRED   CURRENT   READY   AGE
rs1     3          3          3      10h
root@manager-1:/# kubectl delete --all
error: at least one resource must be specified to use a selector
root@manager-1:/# kubectl delete all --all
pod "rs1-2w85d" deleted
pod "rs1-j4f98" deleted
pod "rs1-wrwjv" deleted
service "kubernetes" deleted
service "svc1" deleted
service "svc2" deleted
replicaset.apps "rs1" deleted
```

9. How to run kubectl in any machine which will connect to API server which runs on master machine?

- In enterprise we don't have access to master machine of kubernetes cluster.
- So we can install kubectl in any machine & connect to API server of master machine using below steps:
- Create new machine → install Linux OS → login in
- Install kubectl using below commands
 - >apt-get update
 - >apt-get install -y apt-transport-https ca-certificates curl
 - >apt-get update
 - > sudo curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo gpg --dearmor -o /usr/share/keyrings/kubernetes-archive-keyring.gpg
 - >apt-get update
 - >sudo wget https://raw.githubusercontent.com/lerndevops/labs/master/scripts/installK8S-v1-23.sh -P /tmp
 - >apt-get install -y kubectl=1.23.10-00

Classification: Public

```
root@test:/# apt-get update
Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu focal InRelease
Hit:2 http://us-central1.gce.archive.ubuntu.com/ubuntu focal-updates InRelease
Hit:3 http://us-central1.gce.archive.ubuntu.com/ubuntu focal-backports InRelease
Get:4 https://packages.cloud.google.com/apt kubernetes-xenial InRelease [8993 B]
Hit:5 http://security.ubuntu.com/ubuntu focal-security InRelease
Err:4 https://packages.cloud.google.com/apt kubernetes-xenial InRelease
  The following signatures couldn't be verified because the public key is not available: NO_PUBKEY B53DC80D13EDEF05
Reading package lists... Done
W: GPG error: https://packages.cloud.google.com/apt kubernetes-xenial InRelease: The following signatures couldn't be verified because the public key is not available
  UBKEY B53DC80D13EDEF05
E: The repository 'https://apt.kubernetes.io kubernetes-xenial InRelease' is not signed.
N: Updating from such a repository can't be done securely, and is therefore disabled by default.
N: See apt-secure(8) manpage for repository creation and user configuration details.
root@test:/# apt-get install -y apt-transport-https ca-certificates curl
Reading package lists... Done
Building dependency tree
Reading state information... Done
ca-certificates is already the newest version (20230311ubuntu0.20.04.1).
curl is already the newest version (7.68.0-1ubuntu2.21).
apt-transport-https is already the newest version (2.0.10).
The following packages were automatically installed and are no longer required:
  libatasmart4 libblockdev-fs2 libblockdev-loop2 libblockdev-part-err2 libblockdev-part2 libblockdev-swap2 libblockdev-utils2 libblockdev2 libmbim-glib4 libmbim-proto4 libmbim-glib5 libnsspr4 libnss3 libnupart libparted-fs-resize0 libqmi-glib5 libqmi-proxy libudisks2-0 libxmlb2 usb-modeswitch usb-modeswitch-data
Use 'sudo apt autoremove' to remove them.
@ upgraded, 0 newly installed, 0 to remove and 11 not upgraded.
root@test:/# curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo gpg --dearmor -o /usr/share/keyrings/kubernetes-archive-keyring.gpg
File '/usr/share/keyrings/kubernetes-archive-keyring.gpg' exists. Overwrite? (y/N) y
root@test:/# sudo apt-get update
Hit:1 http://us-central1.gce.archive.ubuntu.com/ubuntu focal InRelease
Hit:2 http://us-central1.gce.archive.ubuntu.com/ubuntu focal-updates InRelease
Hit:3 http://us-central1.gce.archive.ubuntu.com/ubuntu focal-backports InRelease
Hit:4 https://security.ubuntu.com/ubuntu focal-security InRelease
Get:5 https://packages.cloud.google.com/apt kubernetes-xenial InRelease [8993 B]
Get:6 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 Packages [69.9 kB]
Fetched 78.9 kB in 1s (86.9 kB/s)
Reading package lists... Done
root@test:/# sudo apt-get install -y kubectl=1.23.10-00
Reading package lists... Done
Building dependency tree
Reading state information... Done
root@test:/# 
root@test:/# sudo apt-get install -y kubectl=1.23.10-00
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  libatasmart4 libblockdev-fs2 libblockdev-loop2 libblockdev-part-err2 libblockdev-part2 libblockdev-swap2 libblockdev-utils2 libblockdev2 libmbim-glib4 libmbim-proto4 libmbim-glib5 libnsspr4 libnss3 libnupart libparted-fs-resize0 libqmi-glib5 libqmi-proxy libudisks2-0 libxmlb2 usb-modeswitch usb-modeswitch-data
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  kubectl
@ upgraded, 1 newly installed, 0 to remove and 11 not upgraded.
Need to get 8930 kB of archives.
After this operation, 46.6 MB of additional disk space will be used.
Get:1 https://packages.cloud.google.com/apt kubernetes-xenial/main amd64 kubectl amd64 1.23.10-00 [8930 kB]
Fetched 8930 kB in 1s (14.0 MB/s)
Selecting previously unselected package kubectl.
(Reading database ... 62268 files and directories currently installed.)
Preparing to unpack .../kubectl_1.23.10-00_amd64.deb ...
Unpacking kubectl (1.23.10-00) ...
Setting up kubectl (1.23.10-00) ...
root@test:/# apt-get install -y apt-transport-https ca-certificates curl ^c
root@test:/# ^C
root@test:/# kubectl version
Client Version: version.Info{Major:"1", Minor:"23", GitVersion:"v1.23.10", GitCommit:"7e54d50d3012cf3389e43b096ba35300f36e0817", GitTreeState:"clean", GoVersion:"go1.17.13", Compiler:"gc", Platform:"linux/amd64"}
The connection to the server localhost:8080 was refused - did you specify the right host or port?
root@test:/#
```

- Copy admin.conf container from manager machine & create config file in new machine using below steps
 - Login manager machine open admin.conf & copy content
 - > cat /etc/kubernetes/admin.conf
 - Login new machine execute below command
 - > mkdir .kube
 - > cd .kube/
 - > vi config
 - Paste file content & update permission
 - > sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config

```
C:\ root@test: /  
root@test:/# sudo su  
root@test:/# sudo mkdir -p $HOME/.kube  
root@test:/# sudo vi $HOME/.kube/config  
root@test:/# sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Now check in new machine using below command
 - > kubectl get nodes

Classification: Public

```
root@client:/#
root@client:/# kubectl get nodes
NAME      STATUS   ROLES          AGE     VERSION
manager   Ready    control-plane,master   28h    v1.23.10
worker-1  Ready    <none>        27h    v1.23.10
worker-2  Ready    <none>        27h    v1.23.10
root@client:/#
```

10. What is use of label? How we can use label?

- Label is use to identify nodes.
- We can attach label to pod so that we can identify it easily.
- Syntax:> kubectl run <pod name> --image <image name> --labels='<key value labels>'
- Ex:> kubectl run pod2 --image httpd --labels='own=san,env=prod'
- We can check label of pod using below command
- Syntax:> kubectl get pods --show-labels
- We can show pods based on some label condition using below command
- Syntax:> kubectl get pods -l <condition>
- Ex:> kubectl get pods -l own=san

```
root@manager:/#
root@manager:/# kubectl run pod1 --image httpd --labels='own=san,env=prod'
Error from server (AlreadyExists): pods "pod1" already exists
root@manager:/# kubectl run pod2 --image httpd --labels='own=san,env=prod'
pod/pod2 created
root@manager:/# kubectl get pods
NAME  READY  STATUS  RESTARTS  AGE
pod1  1/1    Running  0          8h
pod2  1/1    Running  0          2m35s
root@manager:/# kubectl get pods --show-labels
NAME  READY  STATUS  RESTARTS  AGE  LABELS
pod1  1/1    Running  0          8h    run=prod
pod2  1/1    Running  0          5m25s  env=prod,own=san
root@manager:/# kubectl get pods -l own=san
NAME  READY  STATUS  RESTARTS  AGE
pod2  1/1    Running  0          10m
```

11. How we can use environment variable?

- We can add environment variable to run pod using below command
- Syntax:>kubectl run <pod name> --image <image name> --env='<key value env variables>'
- EX:> kubectl run pod3 --image nginx --env='DBNAME=www.sa.com'
- We can check environment variable using below command
- Ex:> kubectl describe pod3
- We can check environment variable by entering into pod using below command
- Syntax:> kubectl exec -it <pod name> -- bash
- Ex:> kubectl exec -it pod3 - bash
 > env

```

root@manager: / 
pod2 1/1 Running 0 24h
root@manager:/# kubectl run pod3 --image nginx --env='DBNAME=www.sa.com'
pod/pod3 created
root@manager:/# kubectl get pods
NAME READY STATUS RESTARTS AGE
pod1 1/1 Running 0 32h
pod2 1/1 Running 0 24h
pod3 1/1 Running 0 15s
root@manager:/# kubectl describe pod3
error: the server doesn't have a resource type "pod3"
root@manager:/# kubectl describe pod pod3
Name: pod3
Namespace: default
Priority: 0
Node: worker-2/10.128.0.5
Start Time: Thu, 11 Jan 2024 03:08:21 +0000
Labels: run=pod3
Annotations: cni.projectcalico.org/containerID: 92848e7ef2af67c99ba388dba60c4b658d1facc2272e55f27dde808541e56f1
            cni.projectcalico.org/podIP: 192.168.133.195/32
            cni.projectcalico.org/podIPs: 192.168.133.195/32
Status: Running
IP: 192.168.133.195
IPs:
  IP: 192.168.133.195
Containers:
  pod3:
    Container ID: docker://d43b71d04a4fd76317e488fd98eccfde4764ded54b05839a4ae9e0fc53b7bdc
    Image: nginx
    Image ID: docker-pullable://nginx@sha256:2bdc49f2f8ae8d8dc50ed00f2ee56d00385c6f8bc8a8b320d0a294d9e3b49026
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Thu, 11 Jan 2024 03:08:23 +0000
    Ready: True
    Restart Count: 0
    Environment:
      DBNAME: www.sa.com
    Mounts:
      .. /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-9sh2d (ro)
root@manager:/# kubectl exec -it pod3 -- bash
root@pod3:/# env
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_SERVICE_PORT=443
HOSTNAME=pod3
DBNAME=www.sa.com
PWD=/
PKG_RELEASE=1~bookworm
HOME=/root
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
NJS_VERSION=0.8.2
TERM=xterm
SHLVL=1
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
KUBERNETES_SERVICE_HOST=10.96.0.1
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP_PORT=443
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
NGINX_VERSION=1.25.3
_=/usr/bin/env
root@pod3:/# echo JDBC:ODBS:$DBNAME
JDBC:ODBS:
root@pod3:/# echo JDBC:ODBS:$DBNAME
JDBC:ODBS:www.sa.com
root@pod3:#

```

12. How to create pod using YML file?

- We can give request to kubernetes to create pos using YML file.
- To create any pos using YML file we need to add below tags
- **API version:** APIVersion defines the versioned schema of this representation of an object
- **Kind:** Kind is a string value representing the REST resource this object represents
- **Metadata:** Standard object's metadata
- **spec:** Specification of the desired behavior of the pod.
- We can create YML file like below: p.yml

```

apiVersion: v1
kind: Pod
metadata:

```

Classification: Public

```
    name: mypod
    labels:
        env: dev
  spec:
    containers:
```

```
      - name: one
        image: tomcat
```

- We can create pod using below command

```
kubectl create -f p.yml
```

```
root@manager:/#
root@manager:/# vim p.yml
```

```
root@manager:/#
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: mypod
```

```
  labels:
```

```
    env: dev
```

```
spec:
```

```
  containers:
```

```
    - name: one
```

```
      image: tomcat
```

```
root@manager:/# kubectl create -f p.yml
```

```
pod/mypod created
```

```
root@manager:/# kubectl get pods
NAME     READY   STATUS    RESTARTS   AGE
mypod   1/1     Running   0          5m29s
pod1    1/1     Running   1 (29h ago)  2d23h
pod2    1/1     Running   1 (34h ago)  2d15h
pod3    1/1     Running   2 (28h ago)  39h
```

- We can check YML file during running pod using below command

```
Syntax>kubectl run <pod name> --image <image name> --dry-run=client -o yaml
EX:> kubectl run pod4 --image nginx --dry-run=client -o yaml
```

```
root@manager:/#
```

```
root@manager:/# kubectl run pod4 --image nginx --dry-run=client -o yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  creationTimestamp: null
```

```
  labels:
```

```
    run: pod4
```

```
  name: pod4
```

```
spec:
```

```
  containers:
```

```
    - image: nginx
```

```
      name: pod4
```

```
      resources: {}
```

```
  dnsPolicy: ClusterFirst
```

```
  restartPolicy: Always
```

```
status: {}
```

- We can check JSON file during running pod using below command

```
Syntax>kubectl run <pod name> --image <image name> --dry-run=client -o json
EX:> kubectl run pod4 --image nginx --dry-run=client -o json
```

```
root@manager:/# kubectl run pod4 --image nginx --dry-run=client -o json
```

```
{
```

```
  "kind": "Pod",
```

```
  "apiVersion": "v1",
```

```
  "metadata": {
```

```
    "name": "pod4",
```

```
    "creationTimestamp": null,
```

```
    "labels": {
```

```
      "run": "pod4"
```

```
    }
```

```
  },
```

```
  "spec": {
```

```
    "containers": [
```

```
      {
```

```
        "name": "pod4",
```

```
        "image": "nginx",
```

```
        "resources": {}
```

```
      }
```

```
    ],
```

```
    "restartPolicy": "Always",
```

```
    "dnsPolicy": "ClusterFirst"
```

```
  },
```

```
  "status": {}
```

Classification: Public

- Pod can run multiple containers, but cannot be identical i.e. cannot have same image, port.
- Multiple container use case is tomcat & logger container , tomcat & cache container
- When we have tight dependence of 2 containers, those should run together & die together in pod.
- All containers in pod having same IP address.
- We can create multiple container using YAML file **p.yml** as below

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    env: dev
spec:
  containers:
    - name: one
      image: tomcat
    - name: two
      image: nginx
```

- Here we can denote multiple container with **hyphen** -
- Now create container using below command
- Syntax> kubectl create -f p.yml

```
root@manager:/
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    env: dev
spec:
  containers:
    - name: one
      image: tomcat
    - name: two
      image: nginx
~
```

```
root@manager:/
root@manager:/# kubectl create -f p.yml
pod/mypod created
root@manager:/# kubectl get pods
NAME    READY  STATUS   RESTARTS   AGE
mypod  2/2    Running  0          10s
root@manager:/# kubectl get pods -o wide
NAME    READY  STATUS   RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
mypod  2/2    Running  0          16s   192.168.133.201  worker-2  <none>        <none>
root@manager:/# kubectl describe pod mypod
Name:           mypod
Events:
  Type    Reason     Age   From           Message
  ----  -----     --   --            --
  Normal  Scheduled  72s  default-scheduler  Successfully assigned default/mypod to worker-2
  Normal  Pulling    71s  kubelet         Pulling image "tomcat"
  Normal  Pulled    70s  kubelet         Successfully pulled image "tomcat" in 406.85446ms
  Normal  Created    70s  kubelet         Created container one
  Normal  Started    70s  kubelet         Started container one
  Normal  Pulling    70s  kubelet         Pulling image "nginx"
  Normal  Pulled    70s  kubelet         Successfully pulled image "nginx" in 413.340772ms
  Normal  Created    70s  kubelet         Created container two
  Normal  Started    70s  kubelet         Started container two
  Normal  Pulling    70s  kubelet         Pulling image "nginx"
  Normal  Pulled    70s  kubelet         Successfully pulled image "nginx" in 413.340772ms
  Normal  Created    70s  kubelet         Created container three
  Normal  Started    70s  kubelet         Started container three
```

13. What is request & limit in kubernetes?

- In prod environment we need to restrict/limit resource i.e. CPU & RAM. usage by container, because extreme resource usage by one container may impact on other container running in system. So to avoid this situation request & limits came into picture.
- **Request** is minimum resource we need to run pod/container & perform its task.
- When we give request to API Server via kubectl client to run pod using minimum resource optionally, **scheduler** check that minimum number of resource is available in which node, in that node on pod gets run.
- **Limit** is maximum number of usage, where to stop occupying resource by pod/container.
- When you specify a resource limit for a container to run, the **kubelet** enforces those limits so that the running container is not allowed to use more of that resource than the limit you set.
- If the node where a Pod is running has enough of a resource available, it's possible (and allowed) for a container to use more resource than its **request** for that resource specifies. However, a container is not allowed to use more than its resource **limit**.
- In kubernetes we measure CPU(Computer Processing Unit) in core & RAM(Random Access Memory) in binary system Mebibyte (MiB)
- 1 core = 1000milli CPU
- In binary 1kb =1024byte in decimal 1kb=1000 byte
- We can check CPU using command:> lscpu
- We can check RAM using command:> free -h
- We can create container with request & limit as below YML file l.yml

```

apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
    - name: app
      image: images.my-company.example/app:v4
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "300m"
    - name: log-aggregator
      image: images.my-company.example/log-aggregator:v6
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "300m"
  
```

- Now run container using below command
- Syntax:> kubectl create -f l.yml
- We can modify YML file to change image & update existing container using below command
- Change existing YML file l.yml using command > vim l.yml

```

apiVersion: v1
kind: Pod
metadata:
  name: frontend
  
```

```

spec:
  containers:
    - name: app
      image: tomcat
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "300m"
    - name: log-aggregator
      image: nginx
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "300m"

```

- Update container using below command
- Syntax:> kubectl apply -f <changed file name>
- Ex:> kubectl apply -f 1.yml

```

root@manager:/#
root@manager:/# vim 1.yml
root@manager:/# kubectl create -f 1.yml
pod/frontend created
root@manager:/# kubectl get pods
NAME     READY   STATUS    RESTARTS   AGE
frontend 0/2     ImagePullBackOff 0          16s
mypod    2/2     Running   0          126m
root@manager:/# vim 1.yml
root@manager:/# kubectl get pods
NAME     READY   STATUS    RESTARTS   AGE
frontend 0/2     ImagePullBackOff 0          4m13s
mypod    2/2     Running   0          130m
root@manager:/# vim 1.yml
[44L, 418C written]
root@manager:/# kubectl apply -f 1.yml
Warning: resource pods/frontend is missing the kubectl.kubernetes.io/last-applied-configuration annotation which is required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or kubectl apply. The missing annotation will be patched automatically.
pod/frontend configured
root@manager:/# kubectl get pods
NAME     READY   STATUS    RESTARTS   AGE
frontend 2/2     Running   0          8m31s
mypod    2/2     Running   0          134m
root@manager:/# vim 1.yml
root@manager:/# root@manager:#

```

- **Request** of resource for container should always less than **limit**, else it cannot create container, it will give compilation error while creating container

```

root@manager:/
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
    - name: app
      image: tomcat
      resources:
        requests:
          memory: "64Mi"
          cpu: "3"
        limits:
          memory: "128Mi"
          cpu: "300m"
    - name: log-aggregator
      image: nginx
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "300m"

```

- If we try to create container using command, give compilation error
- Ex:>kubectl create -f l.yml

```
root@manager:/#
root@manager:/# vim l.yml
root@manager:/# kubectl create -f l.yml
The Pod "frontend" is invalid: spec.containers[0].resources.requests: Invalid value: "3": must be less than or equal to cpu limit
root@manager:/#
```

- If we give over limit of resource for container than system capacity it accept request but cannot create container
- Change limit in YML file > vi l.yml

```
root@manager: /
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
    - name: app
      image: tomcat
      resources:
        requests:
          memory: "64Mi"
          cpu: "2.5"
        limits:
          memory: "128Mi"
          cpu: "3"
    - name: log-aggregator
      image: nginx
      resources:
        requests:
          memory: "64Mi"
          cpu: "200m"
        limits:
          memory: "128Mi"
          cpu: "300m"
~
```

- If we try to create container using command, it creates container but not able to run container because of limited resource
- Ex:>kubectl create -f l.yml

```
root@manager:/#
root@manager:/# vim l.yml
root@manager:/# kubectl create -f l.yml
pod/frontend created
root@manager:/# kubectl get pods
NAME      READY  STATUS    RESTARTS   AGE
frontend  0/2    Pending   0          12s
root@manager:/# kubectl describe pod frontend
Name:           frontend
Events:
  Type     Reason            Age     From           Message
  ----   -----           ----   ----           -----
  Warning FailedScheduling  20s    default-scheduler  0/3 nodes are available: 1 node(s) had taint {node-role.kubernetes.io/master: }, that the pod didn't tolerate, 2 Insufficient cpu.
root@manager:/#
```

- If container's RAM limit is reached container will restart automatically
- If container's CPU limit is reached it reduces performance

14. What is static pod in kubernetes?

- We can create pod on worker node using kubelet running on it, such pods are called as static pods.
- These pods are managed by kubelet directly
- We can check kubelt service running status using below command
- Syntax:> service kubelet status
- Kubelet is install in location /var/lib/kubelet/
- We can check **config.yaml** file in above location, in that file we can find location where static pod's YML file to be placed.
- The name of tag to find location is **staticPodPath: /etc/kubernetes/manifests**
- Kubelet keeps watching on static post path i.e. **/etc/kubernetes/manifests**, when it find any YML file it create that POD
- We can also change this location in **config.yaml** file

```
root@worker-1:/var/lib/kubelet
root@worker-1:/# kubectl get nodes
The connection to the server localhost:8080 was refused - did you specify the right host or port?
root@worker-1:/# service kubelet status
● kubelet.service - The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
             └─10-kubeadm.conf
     Active: active (running) since Thu 2024-01-11 12:53:15 UTC; 1 day 20h ago
       Docs: https://kubernetes.io/docs/home/
 Main PID: 510 (kubelet)
   Tasks: 17 (limit: 1134)
  Memory: 88.7M
    CGroup: /system.slice/kubelet.service
            └─510 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml
root@worker-1:/# cd /var/lib/kubelet/
root@worker-1:/var/lib/kubelet# ls -lth
total 40K
-rw-r--r-- 1 root root 93 Jan 8 14:53 kubeadm-flags.env
-rw-r--r-- 1 root root 1015 Jan 8 14:53 config.yaml
drwxr-xr-x 2 root root 4.0K Jan 8 14:53 pkcs
drwxr-xr-x 2 root root 4.0K Jan 8 14:53 plugins_registry
drwxr-xr-x 2 root root 4.0K Jan 8 14:53 plugins
-rw-r----- 1 root root 62 Jan 8 14:53 cpu_manager_state
-rw-r----- 1 root root 61 Jan 8 14:53 memory_manager_state
drwxr-xr-x 4 root root 4.0K Jan 8 14:53 pods
drwxr-xr-x 2 root root 4.0K Jan 11 12:53 pod-resources
drwxr-xr-x 2 root root 4.0K Jan 11 12:53 device-plugins
root@worker-1:/var/lib/kubelet# cat config.yaml
apiVersion: kubelet.config.k8s.io/v1beta1
nodeStatusReportFrequency: 0s
nodeStatusUpdateFrequency: 0s
resolvConf: /run/systemd/resolve/resolv.conf
rotateCertificates: true
runtimeRequestTimeout: 0s
shutdownGracePeriod: 0s
shutdownGracePeriodCriticalPods: 0s
staticPodPath: /etc/kubernetes/manifests
streamingConnectionIdleTimeout: 0s
syncFrequency: 0s
volumeStatsAggPeriod: 0s
```

- Got to location **/etc/kubernetes/manifests**, create YML file in that folder > vi p.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    env: dev
spec:
  containers:
    - name: one
      image: tomcat
    - name: two
      image: nginx
```

- Got to location **/etc/kubernetes/manifests**, create YML file in that folder > vi p.yml

```
root@worker-1:/etc/kubernetes/manifests
root@worker-1:/var/lib/kubelet# cd /etc/kubernetes/manifests
root@worker-1:/etc/kubernetes/manifests# ls -lth
total 0
root@worker-1:/etc/kubernetes/manifests# vi p.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    env: dev
spec:
  containers:
    - name: one
      image: tomcat
    - name: two
      image: nginx
```

- Now we can check pod is running or not on manager using below command
- >kubectl get nodes
- >kubectl get nodes -o wide

```
root@manager:/# kubectl get pods
NAME     READY   STATUS    RESTARTS   AGE
frontend 0/2     Pending   0          103m
mypod-worker-1 2/2     Running  2 (62s ago) 117s
root@manager:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS   GATES
frontend 0/2     Pending   0          112m  <none>       <none>  <none>        <none>
mypod-worker-1 2/2     Running  2 (9m45s ago) 10m  192.168.226.66  worker-1  <none>        <none>
```

- So whenever we static pod on any of the worker node, its name will be shown like <pod name>-<node name> i.e. mypod-worker-1
- Any pod having node name as extension is called as static pods
- If we try to delete this static pod it will be auto created again, that means static pod are **not managed by API server**, it is **managed by kubelet**
- Ex:> kubectl delete pod mypod-worker-1
- Ex:> kubectl get pods -o wide

```
root@manager:/# kubectl delete pod mypod-worker-1
pod "mypod-worker-1" deleted
root@manager:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS   GATES
frontend 0/2     Pending   0          130m  <none>       <none>  <none>        <none>
mypod-worker-1 0/2     Pending   0          3s    <none>       worker-1  <none>        <none>
root@manager:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS   GATES
frontend 0/2     Pending   0          130m  <none>       <none>  <none>        <none>
mypod-worker-1 2/2     Running  2 (27m ago) 7s   192.168.226.66  worker-1  <none>        <none>
root@manager:/#
```

- If we want to delete static pod we need to remove YML file from worker

```
root@worker-1:/etc/kubernetes/manifests# rm -rf p.yaml
root@worker-1:/etc/kubernetes/manifests#
```

```
root@manager:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS   GATES
frontend 0/2     Pending   0          135m  <none>       <none>  <none>        <none>
root@manager:/#
```

15. What is namespace in kubernetes? Why we need namespace?

- Suppose in big house if there are different-different rooms are available, so one room person cannot handle resource of another room.
- Similarly in big kubernetes cluster, which is shared across multiple teams, there should have restriction to access resource/pod of one team from another team directly.
- So to provide this kind of restriction to resource/pod in kubernetes **namespace** is used.
- We can create separate spaces to run different pod in different space, which is called as **namespace**.
- It is a way to isolate workloads.
- We can create namespace using below command
- Syntax:> kubectl create ns <name space name>
- Ex:> kubectl create ns ns1
- Ex:> kubectl create ns ns2
- We can create pod to run in above name space using below command
- Syntax:> kubectl run <pod name> --image <image name> -n <namespace>

- Ex:> kubectl run pod1 --image nginx -n ns1
- Ex:> kubectl run pod2 --image nginx -n ns2
- Here we have creates 2 pods which run in 2 different namespaces i.e. ns1:pod1 & ns2:pod2
- Now if we try to access these pods it is give error "No resources found in default namespace."

```
root@manager:/#
root@manager:/# kubectl create ns ns1
namespace/ns1 created
root@manager:/# kubectl create ns ns2
namespace/ns2 created
root@manager:/# kubectl run pod1 --image nginx -n ns1
pod/pod1 created
root@manager:/# kubectl run pod2 --image nginx -n ns2
pod/pod2 created
root@manager:/# kubectl get pods
No resources found in default namespace.
root@manager:/#
```

- So these resource/pods are available in their respective namespaces.
- In this way we can create different-different namespaces & give access to teams, so that they can run their own pods in their respective namespace only.
- Manager can access these pods using below command
- Syntax:> kubectl get pods -n <namespace>
- Ex:> kubectl get pods -n ns1
- Ex:> kubectl get pods -n ns2

```
root@manager:/# kubectl get pods -n ns1
NAME    READY  STATUS   RESTARTS  AGE
pod1   1/1    Running   0          17m
root@manager:/# kubectl get pods -n ns2
NAME    READY  STATUS   RESTARTS  AGE
pod2   1/1    Running   0          17m
root@manager:/#
```

- Manager can delete namespace using below command
- Syntax:> kubectl delete ns <namespace>
- Ex:> kubectl delete ns ns1

```
root@manager:/# kubectl delete ns ns1
namespace "ns1" deleted
```

- Manager can get all namespace using below command
- Syntax:> kubectl get ns

```
root@manager:/# kubectl get ns
NAME        STATUS  AGE
default     Active  4d20h
kube-node-lease  Active  4d20h
kube-public   Active  4d20h
kube-system   Active  4d20h
ns2          Active  27m
```

- In kubernetes we have below namespaces inbuilt
- Default:** If we don't specify any namespace kubectl will create resource in default namespace.
- Kube-node-lease:** This namespace is for system health checkup pods, like heart beat messages & all
- Kube-public:** If we create some resource in public namespace it will be available for all. For ex monitoring pods.
- Kube-system:** It holds all system level resources, we can get all system level resource using below command
- Syntax:>kubectl get pods -n kube-system

```

root@manager:/# kubectl get pods -n kube-system -o wide
NAME          READY   STATUS    RESTARTS   AGE      IP           NODE   NOMINATED NODE   READINESS GATES
calico-kube-controllers-66966888c4-phmft  1/1    Running   1 (2d4h ago)  4d21h  192.168.102.69  manager   <none>        <none>
calico-node-f9cm  1/1    Running   1 (47h ago)   4d21h  10.128.0.4   worker-1  <none>        <none>
calico-node-gs22w  1/1    Running   1 (47h ago)   4d21h  10.128.0.5   worker-2  <none>        <none>
calico-node-vhlnd  1/1    Running   1 (2d4h ago)  4d21h  10.128.0.7   manager   <none>        <none>
coredns-64897985d-7t5qw  1/1    Running   1 (2d4h ago)  4d21h  192.168.102.68  manager   <none>        <none>
coredns-64897985d-slbqx  1/1    Running   1 (2d4h ago)  4d21h  192.168.102.70  manager   <none>        <none>
etcd-manager    1/1    Running   1 (2d4h ago)  4d21h  10.128.0.7   manager   <none>        <none>
kube-apiserver-manager  1/1    Running   1 (47h ago)   4d21h  10.128.0.7   manager   <none>        <none>
kube-controller-manager-manager  1/1    Running   1 (47h ago)   4d21h  10.128.0.5   worker-2  <none>        <none>
kube-proxy-52xmf  1/1    Running   1 (47h ago)   4d21h  10.128.0.5   worker-2  <none>        <none>
kube-proxy-tmktl  1/1    Running   1 (47h ago)   4d21h  10.128.0.7   manager   <none>        <none>
kube-proxy-wrcv6  1/1    Running   1 (47h ago)   4d21h  10.128.0.4   worker-1  <none>        <none>
kube-scheduler-manager  1/1    Running   1 (2d4h ago)  4d21h  10.128.0.7   manager   <none>        <none>
root@manager:/# cd /etc/kubernetes/manifests
root@manager:/etc/kubernetes/manifests# ls -ltrh
total 16K
-rw----- 1 root root 1.5K Jan  8 14:45 kube-scheduler.yaml
-rw----- 1 root root 3.2K Jan  8 14:45 kube-controller-manager.yaml
-rw----- 1 root root 3.8K Jan  8 14:45 kube-apiserver.yaml
-rw----- 1 root root 2.2K Jan  8 14:45 etcd.yaml

```

- Here we can see **etcd, apiserver, controller-manager & scheduler** all are running as static pods.
- Kubernetes software itself running as container & it manages our application containers

16. What is Namespace Quota in kubernetes?

- When we have multiple namespaces in shared clusters, it is our responsibility to restrict resource usage of namespaces; else there might be chance of consuming all resources by single namespace.
- To avoid such situation Namespace Quota came into picture.
- Whenever any new application come at kuebenates admin for hosting in shared cluster, generally they ask for below details
 - **Object count:** how many deployment/pods/controller & service will run?
 - **Resource utilization:** how much CPU & RAM require for them?
- Based on above question they put limit on namespace, that limit is called as resource quota.
- A resource quota, defined by a ResourceQuota object, provides constraints that limit aggregate resource consumption per namespace.
- It can limit the quantity/number of objects that can be created in a namespace by type, as well as the total amount of compute resources that may be consumed by resources in that project.

Object count quota:

- We can limit the quantity/number of objects that can be created in a namespace by below type : pods, services, replicationcontrollers, configmaps, persistentvolumeclaims, resourcequotas, services.loadbalancers, services.nodeports, secrets
- Let's take example below YML(objquota.yml) in which we have specified limit on object as below on our namespace "prod"

```

apiVersion: v1
kind: Namespace
metadata:
  name: prod
  labels:
    name: production
---
apiVersion: v1
kind: ResourceQuota
metadata:
  name: object-counts
  namespace: prod
spec:
  hard:
    pods: "4"
    replicationcontrollers: "2"

```

Classification: Public

```
services: "10"
count/deployments.apps: "2"
count/replicasets.apps: "2"

root@manager-1: ~
apiVersion: v1
kind: Namespace
metadata:
  name: prod
  labels:
    name: production
---
apiVersion : v1
kind: ResourceQuota
metadata:
  name: object-counts
  namespace: prod
spec:
  hard:
    pods: "4"
    replicationcontrollers: "2"
    services: "10"
    count/deployments.apps: "2"
    count/replicasets.apps: "2"
```

- We can create this quota using below command
- Syntax: > kubectl create -f objquota.yml
- We can check quota of our namespace using below command
- Syntax: > kubectl get quota -n <name of namespace>
- Ex :> kubectl get quota -n prod

```
root@manager-1: ~
root@manager-1:~# kubectl create -f objquota.yml
namespace/prod created
resourcequota/object-counts created
root@manager-1:~# kubectl get quota -n prod
NAME     AGE   REQUEST           LIMIT
object-counts  4m11s  count/deployments.apps: 0/2, count/replicasets.apps: 0/2, pods: 0/4, replicationcontrollers: 0/2, services: 0/10
```

- Now we can check by deployment over given limit i.e. 2 using below command
- Syntax:> kubectl create deploy dep1 --image nginx -n prod
- Syntax:> kubectl create deploy dep2 --image nginx -n prod
- We can check quota after creating above 2 deployment, it show expected & allowed number of deployment count are equal i.e. 2
- Now if we create one more deployment, then we got error for quota limit exceed
- Syntax:> kubectl create deploy dep3 --image nginx -n prod

```
root@manager-1: ~
root@manager-1:~# kubectl get quota -n prod
NAME     AGE   REQUEST           LIMIT
object-counts  11m  count/deployments.apps: 0/2, count/replicasets.apps: 0/2, pods: 0/4, replicationcontrollers: 0/2, services: 0/10
root@manager-1:~# kubectl create deploy dep1 --image nginx -n prod
deployment.apps/dep1 created
root@manager-1:~# kubectl get quota -n prod
NAME     AGE   REQUEST           LIMIT
object-counts  14m  count/deployments.apps: 1/2, count/replicasets.apps: 1/2, pods: 1/4, replicationcontrollers: 0/2, services: 0/10
root@manager-1:~# kubectl create deploy dep2 --image nginx -n prod
deployment.apps/dep2 created
root@manager-1:~# kubectl get quota -n prod
NAME     AGE   REQUEST           LIMIT
object-counts  14m  count/deployments.apps: 2/2, count/replicasets.apps: 2/2, pods: 2/4, replicationcontrollers: 0/2, services: 0/10
root@manager-1:~# kubectl create deploy dep3 --image nginx -n prod
error: failed to create deployment: "deployments.apps/dep3" is forbidden: exceeded quota: object-counts, requested: count/deployments.apps=1, used: count/deployments.apps=2, limited: count/deployments.apps=2
root@manager-1:~#
```

- Here we can also increase number of pod replica to check limit exceed, we have limit of pod as 4 & if we increase this to 5, using below command than it will create only up to 4 pods.
- Syntax:> kubectl scale deploy dep1 --replicas 5 -n prod

```
root@manager-1: ~
root@manager-1:~# kubectl get quota -n prod
NAME     AGE   REQUEST           LIMIT
object-counts  36m  count/deployments.apps: 2/2, count/replicasets.apps: 2/2, pods: 2/4, replicationcontrollers: 0/2, services: 0/10
root@manager-1:~# kubectl scale deploy dep1 --replicas 5 -n prod
deployment.apps/dep1 scaled
root@manager-1:~# kubectl get quota -n prod
NAME     AGE   REQUEST           LIMIT
object-counts  38m  count/deployments.apps: 2/2, count/replicasets.apps: 2/2, pods: 4/4, replicationcontrollers: 0/2, services: 0/10
root@manager-1:~# kubectl get deploy -n prod
NAME      READY  UP-TO-DATE  AVAILABLE  AGE
dep1      3/5     3          3          24m
dep2      1/1     1          1          24m
```

Classification: Public

- Here we can also change quota limit using below command
- Syntax:> kubectl edit quota -n <name of namespace>
- Ex:> kubectl edit quota -n prod
- After executing above command we can change quota limit, save & check for limit change

```
root@manager-1:~#
root@manager-1:# kubectl get quota -n prod
NAME      AGE   REQUEST           LIMIT
object-counts  59m   count/deployments.apps: 2/2, count/replicasets.apps: 2/2, pods: 4/4, replicationcontrollers: 0/2, services: 0/10
root@manager-1:# kubectl edit quota -n prod

# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
kind: ResourceQuota
metadata:
  creationTimestamp: "2024-01-28T16:56:26Z"
  name: object-counts
  namespace: prod
  resourceVersion: "850283"
  uid: a307fdf4-c487-41a0-9bbc-f8bb31c5f7fa
spec:
  hard:
    count/deployments.apps: "2"
    count/replicasets.apps: "2"
    pods: "6"
    replicationcontrollers: "2"
    services: "10"
status:
  hard:
    count/deployments.apps: "2"
    count/replicasets.apps: "2"
    pods: "4"
    replicationcontrollers: "2"
    services: "10"
  used:
    count/deployments.apps: "2"
    count/replicasets.apps: "2"
    pods: "4"
    replicationcontrollers: "0"
    services: "0"
...
resourcequota/object-counts edited
root@manager-1:# kubectl get quota -n prod
NAME      AGE   REQUEST           LIMIT
object-counts  65m   count/deployments.apps: 2/2, count/replicasets.apps: 2/2, pods: 4/6, replicationcontrollers: 0/2, services: 0/10
root@manager-1:#
```

Resource utilization Quota

- We can limit total amount of compute resources that may be consumed by resources in that project.
- Let's take example below YML(resquota.yml) in which we have specified limit on CPU as **cpu: "2"** & RAM as **memory: 2Gi** on namespace "team1"

```
apiVersion: v1
kind: Namespace
metadata:
  name: team1
  labels:
    name: development
---
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources-demo
  namespace: team1
spec:
  hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
```

```

limits.memory: 2Gi
root@manager-1: ~
apiVersion: v1
kind: Namespace
metadata:
  name: team1
  labels:
    name: development
---
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources-demo
  namespace: team1
spec:
  hard:
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi

```

- We can create this namespace using below command
- Syntax:> kubectl create -f rsquota.yml
- We can check quota liming using below command
- Syntax:> kubectl get quota -n team1

```

root@manager-1: ~
root@manager-1:~# vim rsquota.yml
root@manager-1:~# kubectl create -f rsquota.yml
namespace/team1 created
resourcequota/compute-resources-demo created
root@manager-1:~# kubectl get quota -n team1
NAME          AGE     REQUESTS   LIMITS
compute-resources-demo  2m43s  requests.cpu: 0/1, requests.memory: 0/1Gi  limits.cpu: 0/2, limits.memory: 0/2Gi
root@manager-1:~

```

17. What is Service in kubernetes?

- When we try to access any container IP from external machine, kubernetes will not allow accessing it from outside of kuberentes cluster, but it can be accessible from any node of kubernetes cluster i.e. Manager/worker.
- Pod IP & port will not work from outside, but it will work within kubernetes cluster because its scope is within kubernetes network.

```

root@manager-1: /#
root@manager-1:/# kubectl run pod1 --image nginx
pod/pod1 created
root@manager-1:/# kubectl get pods -o wide
NAME    READY  STATUS    RESTARTS   AGE      IP           NODE    NOMINATED NODE   READINESS GATES
pod1   1/1    Running   0          23s     192.168.226.65   worker-1   <none>        <none>
root@manager-1:#

```

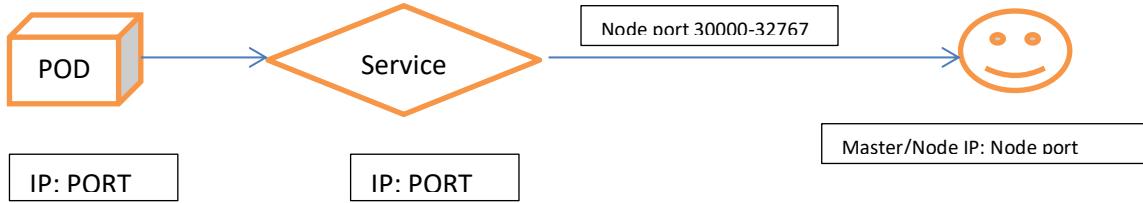


```
root@manager-1:/# curl 192.168.226.65
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

- So to overcome above situation service concept came into picture in kubernetes.
- Service is an object, which provides you feature to access container from outside kubernetes or from inside kubernetes cluster.
- When we create service, it maps to POD.
- Service is having its own IP & port, but this also works internal to kubernetes only.
- There are 2 type of service available
 - NodePort service
 - ClusterIP service

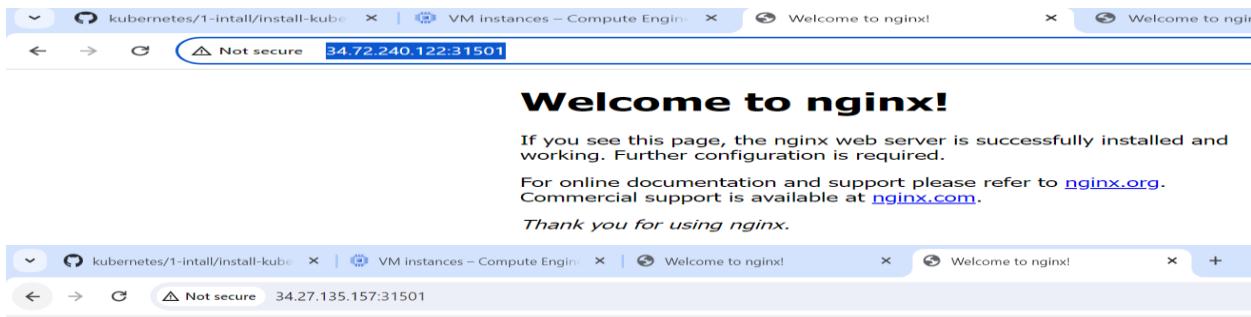
NodePort Service:

- Service publishes **node port**; using node port & IP of master or worker, end user can access application outside of kubernetes.
- This node port ranges from 30000 to 32767



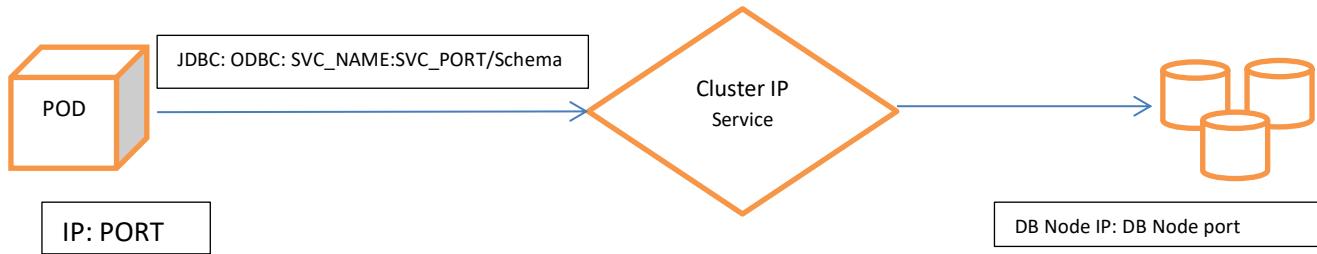
- We can use below command to create service
- Syntax:>kubectl expose pod <pod name> --name <service name> --port <service port> --target-port <port of container of POD> --type NodePort
- Ex:> kubectl expose pod pod1 --name svc1 --port 1234 --target-port 80 --type NodePort
- We can use below command to list services
- Synatx:> kubectl get svc
- Synatx:> kubectl get service
- We can access application from browser using Master/Node(worker) IP : NodePort
- Manager:> http://34.72.240.122:31501/
 - Worker:> http:// 34.27.135.157:31501

```
root@manager-1:/#
root@manager-1:/# kubectl expose pod pod1 --name svc1 --port 1234 --target-port 80 --type NodePort
service/svc1 exposed
root@manager-1:/# kubectl get svc
NAME      TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)        AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP       100m
svc1      NodePort   10.97.80.109  <none>        1234:31501/TCP 114s
root@manager-1:/# kubectl get service
NAME      TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)        AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP       100m
svc1      NodePort   10.97.80.109  <none>        1234:31501/TCP 119s
root@manager-1:/#
```



ClusterIP Service:

- When there is requirement to connect one application with another application within same cluster like web application connects to DB application which running on different pods, we need to use ClusterIP service.
- To connect to DB from application we use connection string which is having DB IP & DB port, in general we give HOST Name instead of IP, but if DB application is running in different pods, it is difficult to give DB IP/Name. So to deal with this issue we need to use cluster IP service.



- We can create cluster IP service using below command.
- Syntax:> kubectl expose rs <name of replica set> --name <service name> --port <service port> --target-port <application PORT> --type ClusterIP
- Ex:> kubectl expose rs rs1 --name svc1 --port 2345 --target-port 3000 --type ClusterIP
- If we do not give any type by default also it will create cluster IP service only
- Syntax:> kubectl expose rs <name of replica set> --name <service name> --port <service port> --target-port <application PORT>
- Ex:> kubectl expose rs rs1 --name svc2 --port 2345 --target-port 3000

```
root@manager-1:/#
root@manager-1:/# kubectl expose rs rs1 --name svc1 --port 2345 --target-port 3000 --type ClusterIP
service/svc1 exposed
root@manager-1:/# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>          443/TCP      4h5m
svc1       ClusterIP  10.106.194.48  <none>          2345/TCP     15s
root@manager-1:/# kubectl expose rs rs1 --name svc3 --port 2345 --target-port 3000
service/svc3 exposed
root@manager-1:/# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>          443/TCP      4h9m
svc1       ClusterIP  10.106.194.48  <none>          2345/TCP     4m46s
svc2       ClusterIP  10.97.106.220   <none>          2345/TCP     27s
svc3       ClusterIP  10.108.154.124  <none>          2345/TCP     3s
```

- Now we can directly access this service from kubernetes cluster
- Syntax:> while true; do curl 10.97.106.220:2345; echo ""; sleep 1; done

```
root@manager-1:/# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes   ClusterIP  10.96.0.1    <none>        443/TCP     4h13m
svc1       ClusterIP  10.106.194.48  <none>        2345/TCP     8m2s
svc2       ClusterIP  10.97.106.220  <none>        2345/TCP     3m43s
root@manager-1:/# while true; do curl 10.97.106.220:2345; echo ""; sleep 1; done
this is my V1 app - Container ID: rs1-wrwjv
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-wrwjv
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-2w85d
```

18. How can we give our own NodePort to service?

- We can give our own NodePort to access application by specifying in YML file for service creation using below steps.
- Support we want to deploy on python application (lerndevops/samplepyapp) which run on 3000 port, so first we need to run pod of that image
- Ex:> kubectl run pod2 --image lerndevops/samplepyapp:v1
- Now using dry-run we can get YML file for service creation by using below command
- Syntax:> kubectl expose pod <pod name> --name <service name> --port <service port> --target-port <application port> --type NodePort --dry-run=client -o yaml
- Ex:> kubectl expose pod pod2 --name svc2 --port 2345 --target-port 3001 --type NodePort --dry-run=client -o yaml
- Now just copy YML output & create new YML file

```
root@manager-1:/# kubectl run pod2 --image lerndevops/samplepyapp:v1
pod/pod2 created
root@manager-1:/# kubectl get pods
NAME      READY  STATUS      RESTARTS  AGE
pod2     0/1   ContainerCreating  0          4s
root@manager-1:/# kubectl get pods
NAME      READY  STATUS      RESTARTS  AGE
pod1     1/1   Running  0          87m
pod2     0/1   ContainerCreating  0          8s
root@manager-1:/# kubectl get pods
NAME      READY  STATUS      RESTARTS  AGE
pod1     1/1   Running  0          88m
pod2     1/1   Running  0          18s
root@manager-1:/# kubectl expose pod pod2 --name svc2 port 2345 --target-port 3001 --type NodePort --dry-run=client -o yaml
couldn't find port via --port flag or introspection
See 'kubectl expose -h' for help and examples
Error from server:NotFound: pods "port" not found
Error from server:NotFound: pods "2345" not found
root@manager-1:/# kubectl expose pod pod2 --name svc2 --port 2345 --target-port 3001 --type NodePort --dry-run=client -o yaml
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    run: pod2
    name: svc2
spec:
  ports:
  - port: 2345
    protocol: TCP
    targetPort: 3001
  selector:
    run: pod2
  type: NodePort
status:
  loadBalancer: {}
```

- In YML file we can specify our own highlighted node port as below & create YML File (s.yml)

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    run: pod2
    name: svc2
spec:
  ports:
  - port: 2345
    protocol: TCP
    targetPort: 3000
    nodePort: 30001
```

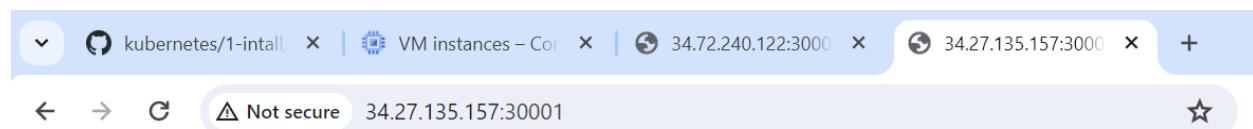
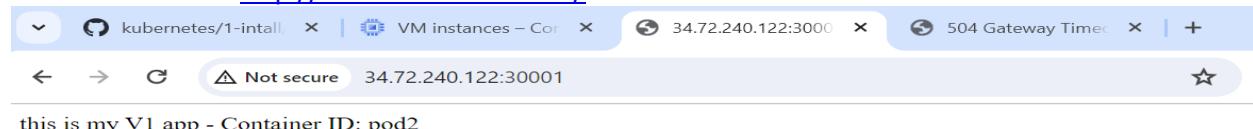
Classification: Public

```
selector:  
  run: pod2  
type: NodePort  
status:  
  loadBalancer: {}
```

- Now create service using below command
- Ex:> kubectl create -f s.yml

```
root@manager-1:/#  
apiVersion: v1  
Kind: Service  
metadata:  
  creationTimestamp: null  
  labels:  
    run: pod2  
  name: svc2  
spec:  
  ports:  
  - port: 2345  
    protocol: TCP  
    targetPort: 3000  
    nodePort: 30001  
  selector:  
    run: pod2  
  type: NodePort  
  
root@manager-1:/# kubectl create -f s.yml  
service/svc2 created  
root@manager-1:/# kubectl get svc -o wide  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE      SELECTOR  
kubernetes   ClusterIP   10.96.0.1    <none>        443/TCP     3h26m    <none>  
svc1       NodePort    10.97.80.109  <none>        1234:31501/TCP  107m    run=pod1  
svc2       NodePort    10.109.2.55   <none>        2345:30001/TCP  11s     run=pod2
```

- Now we can access this application using our own port 30001 & IP of manager/worker(nodes)
- Manager : → <http://34.72.240.122:30001/>
- Worker:→ <http://34.27.135.157:30001/>



- We can access this application from manager/worker(node) by using service IP & port

```
service/svc2 created  
root@manager-1:/# kubectl get svc -o wide  
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE      SELECTOR  
kubernetes   ClusterIP   10.96.0.1    <none>        443/TCP     3h26m    <none>  
svc1       NodePort    10.97.80.109  <none>        1234:31501/TCP  107m    run=pod1  
svc2       NodePort    10.109.2.55   <none>        2345:30001/TCP  11s     run=pod2  
root@manager-1:/# vi s.yml  
root@manager-1:/# curl 10.109.2.55:2345  
this is my V1 app - Container ID: pod2root@manager-1:/#  
  
root@worker-1:/#  
root@worker-1:/# curl 10.109.2.55:2345  
this is my V1 app - Container ID: pod2root@worker-1:/#
```

Classification: Public

19. What is controller in kubernetes? Why do we need controller?

- To achieve high availability & work load manage single pod is not sufficient. So we need to run multiple pods at a time.
- This is where controller came into picture.
- Controller helps us to run multiple containers, achieve workload management, high availability, scale up-scale down containers, etc.
- There are many controller in kubernetes
 - Replication controller(RC)
 - Replica set
 - Deployment
 - Daemon set
 - Job
 - Cron job
 - State full set
 - HPA

20) What is ReplicationController in Kuberntes?

- A ReplicationController ensures that a specified number of pod replicas are running at any one time.
- In other words, a ReplicationController makes sure that a pod or a homogeneous set of pods is always up and available.
- If there are too many pods, the ReplicationController terminates the extra pods. If there are too few, the ReplicationController starts more pods.
- Unlike manually created pods, the pods maintained by a ReplicationController are automatically replaced if they fail, are deleted, or are terminated.
- Now days this controller is deprecated.
- We can create replication controller as below rc.yml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: k-admins
spec:
  replicas: 4
  selector:
    skill: kube
  template:
    metadata:
      name: nginx
      labels:
        skill: kube
    spec:
      containers:
        - name: one
          image: nginx
```

- Here we have created ReplicationController having name "k-admins", having 4 number of pods, with selection criteria who have label skill: kube, ReplicationController tries to find pod with matching criteria, if any pod with this criteria is not available ReplicationController will create pod with given criteria.
- Now we can create by using below command:
- Syntax:> kubectl create -f rc.yml

```
root@manager-1:/#
root@manager-1:/# kubectl create -f rc.yml
replicationcontroller/k-admins created
root@manager-1:/# kubectl get rc
NAME      DESIRED   CURRENT   READY   AGE
k-admins   4         4         4       37s
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-47sbv  1/1    Running   0          75s    192.168.226.70  worker-1  <none>        <none>
k-admins-q7xbc  1/1    Running   0          75s    192.168.226.72  worker-1  <none>        <none>
k-admins-qzsqq  1/1    Running   0          75s    192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          75s    192.168.226.69  worker-1  <none>        <none>
```

- If we try to delete any one of pod using below command, replication controller will recreate another pod by using desired state
- Syntax:> kubectl delete pod k-admins-47sbv

```
root@manager-1:/#
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-47sbv  1/1    Running   0          14m    192.168.226.70  worker-1  <none>        <none>
k-admins-q7xbc  1/1    Running   0          14m    192.168.226.72  worker-1  <none>        <none>
k-admins-qzsqq  1/1    Running   0          14m    192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          14m    192.168.226.69  worker-1  <none>        <none>
root@manager-1:/# kubectl delete pod k-admins-47sbv
pod "k-admins-47sbv" deleted
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-jr4jb 1/1    Running   0          55s   192.168.226.73  worker-1  <none>        <none>
k-admins-q7xbc  1/1    Running   0          15m   192.168.226.72  worker-1  <none>        <none>
k-admins-qzsqq  1/1    Running   0          15m   192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          15m   192.168.226.69  worker-1  <none>        <none>
```

- We can change number of replicas on the fly i.e. we can increase or decrease number of replicas using below command
- Syntax: kubectl scale rc <replication controller name> --replicas <count>
- Ex:> kubectl scale rc k-admins --replicas 2
- Ex:> kubectl scale rc k-admins --replicas 6

```
root@manager-1:/#
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-jr4jb 1/1    Running   0          7m43s  192.168.226.73  worker-1  <none>        <none>
k-admins-q7xbc  1/1    Running   0          22m   192.168.226.72  worker-1  <none>        <none>
k-admins-qzsqq  1/1    Running   0          22m   192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          22m   192.168.226.69  worker-1  <none>        <none>
root@manager-1:/# kubectl scale rc k-admins --replicas 2
replicationcontroller/k-admins scaled
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-qzsqq 1/1    Running   0          23m   192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          23m   192.168.226.69  worker-1  <none>        <none>
root@manager-1:/# kubectl scale rc k-admins --replicas 6
replicationcontroller/k-admins scaled
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-26rkq 0/1    ContainerCreating   0      4s    <none>        worker-1  <none>        <none>
k-admins-dgsj2  0/1    ContainerCreating   0      4s    <none>        worker-1  <none>        <none>
k-admins-hzkfj  0/1    ContainerCreating   0      4s    <none>        worker-1  <none>        <none>
k-admins-kb8vv  0/1    ContainerCreating   0      4s    <none>        worker-1  <none>        <none>
k-admins-qzsqq  1/1    Running   0          25m   192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          25m   192.168.226.69  worker-1  <none>        <none>
root@manager-1:/# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
k-admins-26rkq 1/1    Running   0          28s   192.168.226.75  worker-1  <none>        <none>
k-admins-dgsj2  1/1    Running   0          28s   192.168.226.76  worker-1  <none>        <none>
k-admins-hzkfj  1/1    Running   0          28s   192.168.226.77  worker-1  <none>        <none>
k-admins-kb8vv  1/1    Running   0          28s   192.168.226.74  worker-1  <none>        <none>
k-admins-qzsqq  1/1    Running   0          25m   192.168.226.71  worker-1  <none>        <none>
k-admins-td5lb  1/1    Running   0          25m   192.168.226.69  worker-1  <none>        <none>
```

21) What is ReplicaSet in kubernetes?

- It is advanced version of ReplicationController, which is having all features of ReplicationController with some extra features.
- In ReplicationController we can give only single match criteria, but in ReplicaSet we can have multiple selection criteria which can be defined under `matchLabels`, `matchExpression` tag.
- We can give complex expression in selection criteria using `matchExpression` tag.
- We can create ReplicaSet by using below YAML file i.e. rs.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rsl
spec:
  replicas: 3
  selector:
    matchLabels:
```

```

        app: sanapp
template:
  metadata:
    labels:
      app: sanapp
spec:
  containers:
    - name: abc
      image: lerndeops/samplepyapp:v1

```

- Here we have created ReplicaSet with name “rs1”, having 3 number of replica with selection criteria as label app: sanapp
- Now we can create ReplicaSet using below command
- Syntax:> kubectl create -f rs.yaml

```

root@manager-1:/#
root@manager-1:/# kubectl create -f rs.yaml
replicaset.apps/rs1 created
root@manager-1:/# kubectl get rs
NAME   DESIRED  CURRENT  READY   AGE
rs1     3         3         3       16s
root@manager-1:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
rs1-2vgxd  1/1    Running   0          26s   192.168.226.80  worker-1  <none>        <none>
rs1-c47f7  1/1    Running   0          26s   192.168.226.78  worker-1  <none>        <none>
rs1-xnks5  1/1    Running   0          26s   192.168.226.79  worker-1  <none>        <none>
root@manager-1:/# kubectl describe rs rs1
Name:           rs1
Namespace:      default
Selector:       app=sanapp
Labels:         <none>
Annotations:   <none>
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=sanapp
  Containers:
    abc:
      Image:      lerndeops/samplepyapp:v1
      Port:       <none>
      Host Port: <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
Events:
  Type  Reason     Age   From           Message
  ----  ----
  Normal SuccessfulCreate 68s  replicaset-controller  Created pod: rs1-c47f7
  Normal SuccessfulCreate 68s  replicaset-controller  Created pod: rs1-2vgxd
  Normal SuccessfulCreate 68s  replicaset-controller  Created pod: rs1-xnks5

```

- Now if we delete all pods using below command, ReplicaSet will automatically create number of pod as defined in replicas
- Syntax:>kubectl delete pods -all

```

root@manager-1:/#
root@manager-1:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
rs1-2vgxd  1/1    Running   0          9m2s  192.168.226.80  worker-1  <none>        <none>
rs1-c47f7  1/1    Running   0          9m2s  192.168.226.78  worker-1  <none>        <none>
rs1-xnks5  1/1    Running   0          9m2s  192.168.226.79  worker-1  <none>        <none>
root@manager-1:/# kubectl delete pods --all
pod "rs1-2vgxd" deleted
pod "rs1-c47f7" deleted
pod "rs1-xnks5" deleted
root@manager-1:/# kubectl get pods -o wide
NAME     READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
rs1-2w85d  1/1    Running   0          6s    192.168.226.81  worker-1  <none>        <none>
rs1-j4f98  1/1    Running   0          6s    192.168.226.82  worker-1  <none>        <none>
rs1-wrwjv  1/1    Running   0          7s    192.168.226.83  worker-1  <none>        <none>
root@manager-1:/#

```

- Now we can expose this replica set using service concept with below command

- Syntax:> kubectl expose rs <name of replica set> --name <service name> --port <new port> --target-port <application port> --type NodePort
- Ex:> kubectl expose rs rs1 --name svc1 --port 1234 --target-port 3000 --type NodePort

```
root@manager-1:/# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
rs1-2w85d  1/1    Running   0          24m    192.168.226.81  worker-1  <none>        <none>
rs1-j4f98  1/1    Running   0          24m    192.168.226.82  worker-1  <none>        <none>
rs1-wrwjv  1/1    Running   0          24m    192.168.226.83  worker-1  <none>        <none>
root@manager-1:/# kubectl expose rs rs1 --name svc1 --port 1234 --target-port 3000 --type NodePort
service/svc1 exposed
root@manager-1:/# kubectl get svc
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1   <none>       443/TCP     50m
svc1      NodePort   10.99.61.14  <none>       1234:32073/TCP 15s
root@manager-1:/#
```

- We can access application using service port & any of IP of master/worker(node)
- Master :→ <http://34.72.240.122:32073/>



this is my V1 app - Container ID: rs1-wrwjv

- Worker:→ <http://34.27.135.157:32073/>



this is my V1 app - Container ID: rs1-j4f98

- We can do scale update & down for ReplicaSet using below command
- Syntax: kubectl scale rs <name of replicaset > --replicas <count>
- Ex:> kubectl scale rs rs1 --replicas 6
- Ex:> kubectl scale rs rs1 --replicas 2

```
root@manager-1:/# kubectl scale rs rs1 --replicas 6
replicaset.apps/rs1 scaled
root@manager-1:/# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
rs1-2w85d  1/1    Running   0          40m    192.168.226.81  worker-1  <none>        <none>
rs1-bjlng  1/1    Running   0          9s     192.168.226.84  worker-1  <none>        <none>
rs1-bnts7  1/1    Running   0          9s     192.168.226.86  worker-1  <none>        <none>
rs1-j4f98  1/1    Running   0          40m    192.168.226.82  worker-1  <none>        <none>
rs1-qw6sn  1/1    Running   0          9s     192.168.226.85  worker-1  <none>        <none>
rs1-wrwjv  1/1    Running   0          40m    192.168.226.83  worker-1  <none>        <none>
root@manager-1:/# kubectl scale rs rs1 --replicas 3
replicaset.apps/rs1 scaled
root@manager-1:/# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
rs1-2w85d  1/1    Running   0          40m    192.168.226.81  worker-1  <none>        <none>
rs1-j4f98  1/1    Running   0          40m    192.168.226.82  worker-1  <none>        <none>
rs1-wrwjv  1/1    Running   0          40m    192.168.226.83  worker-1  <none>        <none>
root@manager-1:/#
```

- We can also access application from kubernetes cluster i.e. master/worker node using service IP & service port as well

```

root@manager-1:/# kubectl get svc
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1   <none>        443/TCP     75m
svc1       NodePort   10.99.61.14  <none>        1234:32073/TCP 24m
root@manager-1:/# while true; do curl 10.99.61.14:1234; echo ""; sleep 1; done
this is my V1 app - Container ID: rs1-wrwjv
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-2w85d
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-wrwjv

root@worker-1:/
root@worker-1:/# while true; do curl 10.99.61.14:1234; echo ""; sleep 1; done
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-j4f98
this is my V1 app - Container ID: rs1-2w85d
this is my V1 app - Container ID: rs1-wrwjv

```

22) How we can update application without down time? What is use of deployment controller?

- By using deployment controller we can achieve 0 down time deployment.
- You can define Deployments to create new ReplicaSets, or to remove existing Deployments and adopt all their resources with new Deployments.
- There are 2 deployment models we can use
 - Blue green deployment
 - Rolling updates deployment

23) How we can update application with blue-green deployment?

- To explain this first create an application & deploy that using deployment controller as below
- Create deployment with name **mydeployment** using below YML (dc1.yml)

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: mydeployment
  labels:
    app: myapp
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp-count
          image: lerndevops/samplepyapp:v1
          ports:
            - containerPort: 3000

```

- Do deployment using below command:
- Syntax:> kubectl create -f dc.yml
- We can also give repository raw URL path of YML file as below

- Syntax:> kubectl create -f <https://raw.githubusercontent.com/lerndevops/kubernetes/master/3-controllers/deployments/deployment-ex1.yaml>
- Here even if delete any pod it will comes back, i.e. It will also provide you high availability

```
root@manager-1:~#
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mydeployment
  labels:
    app: myapp
spec:
  replicas: 3
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp-count
          image: lerndevops/samplepyapp:v1
          ports:
            - containerPort: 3000
root@manager-1:~#
root@manager-1:~# vim dc.yaml
root@manager-1:~# kubectl create -f dc.yaml
deployment.apps/mydeployment created
root@manager-1:~# kubectl get deployment
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mydeployment   3/3     3           3          14s
root@manager-1:~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
mydeployment-b48f5b864-4shfx   1/1    Running   0          33s   192.168.226.94   worker-1   <none>       <none>
mydeployment-b48f5b864-qk9jp   1/1    Running   0          33s   192.168.226.95   worker-1   <none>       <none>
mydeployment-b48f5b864-vfd17  1/1    Running   0          33s   192.168.226.96   worker-1   <none>       <none>
root@manager-1:~# kubectl delete pod mydeployment-b48f5b864-4shfx
pod "mydeployment-b48f5b864-4shfx" deleted
root@manager-1:~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
mydeployment-b48f5b864-qk9jp   1/1    Running   0          59s   192.168.226.95   worker-1   <none>       <none>
mydeployment-b48f5b864-vfd17  1/1    Running   0          59s   192.168.226.96   worker-1   <none>       <none>
mydeployment-b48f5b864-xvzcf  1/1    Running   0          7s    192.168.226.97   worker-1   <none>       <none>
```

- Now expose this deployment through service with below command
- Syntax:> kubectl expose deploy <name of deployment> --name <service name> --port <new port> --target-port <application port>
- Syntax:> kubectl expose deploy mydeployment --name svc1 --port 1234 --target-port 3000

```
root@manager-1:~#
root@manager-1:~# kubectl get deployment
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mydeployment   3/3     3           3          4m43s
root@manager-1:~# kubectl expose deploy mydeployment --name svc1 --port 1234 --target-port 3000
service/svc1 exposed
root@manager-1:~# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes   ClusterIP   10.96.0.1      <none>        443/TCP      9h
svc1      ClusterIP   10.96.202.254   <none>        1234/TCP      23s
```

- We can access this application using master/node with below command

```
root@worker-1:~#
root@worker-1:~# while true; do curl 10.96.202.254:1234; sleep 1; echo ""; done
this is my V1 app - Container ID: mydeployment-b48f5b864-vfd17
this is my V1 app - Container ID: mydeployment-b48f5b864-qk9jp
this is my V1 app - Container ID: mydeployment-b48f5b864-qk9jp
this is my V1 app - Container ID: mydeployment-b48f5b864-vfd17
this is my V1 app - Container ID: mydeployment-b48f5b864-qk9jp
this is my V1 app - Container ID: mydeployment-b48f5b864-qk9jp
this is my V1 app - Container ID: mydeployment-b48f5b864-vfd17
this is my V1 app - Container ID: mydeployment-b48f5b864-xvzcf
```

- We can see load balancing is happening on all pods.
- Now end user is accessing this application happily & we want to update this to new version, to do this, we create one more deployment with name green-version using another YAML file (dc2.yaml)

```
apiVersion: apps/v1
```

Classification: Public

```
kind: Deployment
metadata:
  name: green-version
  labels:
    app: green
spec:
  replicas: 3
  selector:
    matchLabels:
      app: green
  template:
    metadata:
      labels:
        app: green
    spec:
      containers:
        - name: myapp-cont2
          image: lerndevops/samplepyapp:v2
          ports:
            - containerPort: 3000
```

```
root@manager-1: ~
apiVersion: apps/v1
kind: Deployment
metadata:
  name: green-version
  labels:
    app: green
spec:
  replicas: 3
  selector:
    matchLabels:
      app: green
  template:
    metadata:
      labels:
        app: green
    spec:
      containers:
        - name: myapp-cont2
          image: lerndevops/samplepyapp:v2
          ports:
            - containerPort: 3000
```

- Create deployment using below command
- Syntax:> kubectl create -f dc2.yml
- We can see that both deployment pods are running

```
root@manager-1: ~
root@manager-1: ~# vim dc2.yml
root@manager-1: ~# kubectl create -f dc2.yml
deployment.apps/green-version created
root@manager-1: ~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
green-version-5d6f77c778-bdv9r   1/1   Running   0          22s   192.168.226.98   worker-1   <none>   <none>
green-version-5d6f77c778-lgt8v   1/1   Running   0          22s   192.168.226.99   worker-1   <none>   <none>
green-version-5d6f77c778-pk68b   1/1   Running   0          22s   192.168.226.100  worker-1   <none>   <none>
mydeployment-b48f5b864-qk9jp   1/1   Running   0          48m   192.168.226.95   worker-1   <none>   <none>
mydeployment-b48f5b864-vfd17   1/1   Running   0          48m   192.168.226.96   worker-1   <none>   <none>
mydeployment-b48f5b864-xvzcf   1/1   Running   0          47m   192.168.226.97   worker-1   <none>   <none>
root@manager-1: ~#
```

- Now we can edit service to update new application version deployment as below command
- Syntax:> kubectl edit svc svc1

```
root@manager-1: ~#
root@manager-1: ~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
green-version-5d6f77c778-bdv9r   1/1   Running   0          49m   192.168.226.98   worker-1   <none>   <none>
green-version-5d6f77c778-lgt8v   1/1   Running   0          49m   192.168.226.99   worker-1   <none>   <none>
green-version-5d6f77c778-pk68b   1/1   Running   0          49m   192.168.226.100  worker-1   <none>   <none>
mydeployment-b48f5b864-qk9jp   1/1   Running   0          97m   192.168.226.95   worker-1   <none>   <none>
mydeployment-b48f5b864-vfd17   1/1   Running   0          97m   192.168.226.96   worker-1   <none>   <none>
mydeployment-b48f5b864-xvzcf   1/1   Running   0          96m   192.168.226.97   worker-1   <none>   <none>
root@manager-1: ~# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP      11h
svc1      ClusterIP  10.96.202.254  <none>        1234/TCP      92m
root@manager-1: ~# kubectl edit svc svc1
32L, 732C written
service/svc1 edited
```

Classification: Public

```
root@manager-1: ~
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: "2024-01-27T04:06:20Z"
  labels:
    app: myapp
  name: svc1
  namespace: default
  resourceVersion: "659957"
  uid: 90bcd45d-1704-472f-8f64-984a368a0f61
spec:
  clusterIP: 10.96.202.254
  clusterIPs:
  - 10.96.202.254
  internalTrafficPolicy: Cluster
  ipFamilies:
  - IPv4
  ipFamilyPolicy: SingleStack
  ports:
  - port: 1234
    protocol: TCP
    targetPort: 3000
  selector:
    app: green
  sessionAffinity: None
  type: ClusterIP
status:
loadBalancer: {}
```

- Once we update service we can see those new upcoming requests are going to new deployment pods i.e. green-version pods

```
root@worker-1: ~
this is my V1 app - Container ID: mydeployment-b48f5b864-qk9jp
this is my V1 app - Container ID: mydeployment-b48f5b864-vfd17
this is my V1 app - Container ID: mydeployment-b48f5b864-xvzcf
this is my V1 app - Container ID: mydeployment-b48f5b864-vfd17
this is my V1 app - Container ID: mydeployment-b48f5b864-qk9jp
this is my V1 app - Container ID: mydeployment-b48f5b864-vfd17
this is my V2 app - Container ID: green-version-5d6f77c778-bdv9r
this is my V2 app - Container ID: green-version-5d6f77c778-bdv9r
this is my V2 app - Container ID: green-version-5d6f77c778-pk68b
this is my V2 app - Container ID: green-version-5d6f77c778-pk68b
```

- Now we can delete old deployment as below command

```
root@manager-1:~# kubectl get deploy
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
green-version   3/3     3           3          65m
mydeployment   3/3     3           3          114m
root@manager-1:~# kubectl delete deploy mydeployment
deployment.apps "mydeployment" deleted
```

- Drawback of this blue-green deployment model is that at some point of time we need to install both deployment version pods, which will consume more (double) resources of cluster.
- To overcome this rolling-update deployment model came into picture.

24) How we can update application with rolling-update deployment model?

- In rolling-update deployment model we slowly add new version pods & remove old version pods.
- In this deployment at some point of time request is going in both old & new version pods.
- Here also we don't see any down time for deployment.
- In rolling-update at some point of time we cross desire state of pods, i.e. current version running pods + few more pods, this is called as **surge**.
- In rolling-update we need to defined **strategy** under which we need to mention below
- We can mention max **surge number** also, i.e. number of extra pods with new version to be run, by using keyword **maxSurge**

- Also we can mention max number of old version which we want to delete by using keyword **maxUnavailable**.
- For adding new version pods, we need to specify time gap between 2 pods, so that pod will get properly up & ready to server request, this we can specify using key word **minReadySeconds**.
- Let's create new deployment with name "kubeserve" & service with name "kubeserve-svc" using single YML file(deploy-ru.yml)

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubeserve
spec:
  replicas: 4
  minReadySeconds: 10
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
      maxSurge: 2
  selector:
    matchLabels:
      app: kubeserve
  template:
    metadata:
      name: kubeserve
      labels:
        app: kubeserve
    spec:
      containers:
        - image: lerndevops/kubeserve:v1
          name: app
---
kind: Service
apiVersion: v1
metadata:
  name: kubeserve-svc
spec:
  type: NodePort
  selector:
    app: kubeserve
  ports:
    - port: 80
      targetPort: 80

```

- Now execute below command to create this
- Syntax:> kubectl create -f deploy-ru.yml

Classification: Public

```
root@manager-1: ~
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubeserve
spec:
  replicas: 4
  minReadySeconds: 10
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxUnavailable: 1
      maxSurge: 2
  selector:
    matchLabels:
      app: kubeserve
  template:
    metadata:
      name: kubeserve
      labels:
        app: kubeserve
    spec:
      containers:
        - image: leaddevops/kubeserve:v1
          name: app
---
kind: Service
apiVersion: v1
metadata:
  name: kubeserve-svc
spec:
  type: NodePort
  selector:
    app: kubeserve
  ports:
    - port: 80
      targetPort: 80
```

```
root@manager-1: ~
root@manager-1:~# vim deploy-ru.yml
root@manager-1:~# kubectl create -f deploy-ru.yml
deployment.apps/kubeserve created
service/kubeserve-svc created
root@manager-1:~# kubectl get deploy
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
kubeserve  0/4     4           0           13s
root@manager-1:~# kubectl get pods
NAME            READY   STATUS      RESTARTS   AGE
kubeserve-6dc4d8bd-2psg5  0/1     ContainerCreating   0          23s
kubeserve-6dc4d8bd-2x276  0/1     ContainerCreating   0          23s
kubeserve-6dc4d8bd-p6mwz  0/1     ContainerCreating   0          23s
kubeserve-6dc4d8bd-wr5z5  0/1     ContainerCreating   0          23s
root@manager-1:~# kubectl get pods -o wide
NAME            READY   STATUS      RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
kubeserve-6dc4d8bd-2psg5  0/1   ContainerCreating   0      31s  <none>       worker-1  <none>        <none>
kubeserve-6dc4d8bd-2x276  0/1   ContainerCreating   0      31s  <none>       worker-1  <none>        <none>
kubeserve-6dc4d8bd-p6mwz  0/1   ContainerCreating   0      31s  <none>       worker-1  <none>        <none>
kubeserve-6dc4d8bd-wr5z5  0/1   ContainerCreating   0      31s  <none>       worker-1  <none>        <none>
root@manager-1:~# kubectl get pods -o wide
NAME            READY   STATUS      RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
kubeserve-6dc4d8bd-2psg5  1/1   Running      0          49s  192.168.226.125  worker-1  <none>        <none>
kubeserve-6dc4d8bd-2x276  1/1   Running      0          49s  192.168.226.124  worker-1  <none>        <none>
kubeserve-6dc4d8bd-p6mwz  1/1   Running      0          49s  192.168.226.122  worker-1  <none>        <none>
kubeserve-6dc4d8bd-wr5z5  1/1   Running      0          49s  192.168.226.123  worker-1  <none>        <none>
root@manager-1:~# kubectl get svc
NAME        TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)        AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP       3m27s
kubeserve-svc  NodePort  10.110.126.146  <none>        80:30782/TCP  2m36s
```

- Here 4 pods are running, we can access these by using service IP & service port as below
- Syntax:> while true; do curl 10.110.126.146:80; sleep 1; echo ""; done

```
root@worker-1: ~
root@worker-1:~# while true; do curl 10.110.126.146:80; sleep 1; echo ""; done
This is v1 pod kubeserve-6dc4d8bd-2psg5
This is v1 pod kubeserve-6dc4d8bd-p6mwz
This is v1 pod kubeserve-6dc4d8bd-2x276
This is v1 pod kubeserve-6dc4d8bd-wr5z5
```

- Here we can see request is going on each pods, so load balancing is working properly.
- Here we can check deployment status by using below command
- Syntax:> kubectl rollout status deploy <name of deployment>

Classification: Public

- Ex:> kubectl rollout status deploy kubeserve
- We can check rollout history as well using below command, which will give how many version we have deploy so far.
- Syntax:> kubectl rollout history deploy <name of deployment>
- Ex:> kubectl rollout history deploy kubeserve

```
root@manager-1: ~# kubectl get deploy
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
kubeserve  4/4     4           4           18m
root@manager-1: ~# kubectl rollout status deploy kubeserve
deployment "kubeserve" successfully rolled out
root@manager-1: ~# kubectl rollout history deploy kubeserve
deployment.apps/kubeserve
REVISION  CHANGE-CAUSE
1         <none>
1
```

- Here we have deployed just 1st version so it is showing REVISION as 1
- Now we can deploy 2nd version either by editing YML file or we can use below command to do same
- Syntax> kubectl set image deploy <name of deployment> app=<image>:<new version>
- Ex:> kubectl set image deploy kubeserve app=leaddevops/kubeserv:v2

```
root@manager-1: ~# kubectl rollout status deploy kubeserve
deployment "kubeserve" successfully rolled out
root@manager-1: ~# kubectl rollout history deploy kubeserve
deployment.apps/kubeserve
REVISION  CHANGE-CAUSE
1         <none>

root@manager-1: ~# kubectl set image deploy kubeserve app=leaddevops/kubeserve:v2
deployment.apps/kubeserve image updated
root@manager-1: ~# kubectl rollout status deploy kubeserve
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
deployment "kubeserve" successfully rolled out
root@manager-1: ~# kubectl rollout history deploy kubeserve
deployment.apps/kubeserve
REVISION  CHANGE-CAUSE
1         <none>
2 |         <none>
```

- We can see at end use window customer was getting response from v1 pods, than gradually v2 pods are also giving response in between, after some v1 pods are completely removed & v2 only started serving request

```
root@worker-1: ~
This is v1 pod kubeserve-6dcd44d8bd-zsjm2
This is v1 pod kubeserve-6dcd44d8bd-zsjm2
This is v1 pod kubeserve-6dcd44d8bd-9nxsz
This is v1 pod kubeserve-6dcd44d8bd-pc29k
This is v2 pod kubeserve-7674cb6d4f-4dhqz
This is v1 pod kubeserve-6dcd44d8bd-pc29k
This is v2 pod kubeserve-7674cb6d4f-2pt7h
This is v2 pod kubeserve-7674cb6d4f-4dhqz
```

```
This is v1 pod kubeserve-6dcd44d8bd-8k5fd
This is v2 pod kubeserve-7674cb6d4f-4dhqz
This is v2 pod kubeserve-7674cb6d4f-2pt7h
This is v2 pod kubeserve-7674cb6d4f-2pt7h
This is v2 pod kubeserve-7674cb6d4f-2pt7h
```

- Now consider a scenario, suppose accidentally we have deployed some wrong version pod & it started giving error, so in that situation we need to do rollback to previous version or any correct version.
- We have updated wrong version **v3** image in our application as below

```
root@manager-1:~# kubectl set image deploy kubeserve app=leaddevops/kubeserve:v3
deployment.apps/kubeserve image updated
root@manager-1:~# kubectl rollout status deploy kubeserve
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
deployment "kubeserve" successfully rolled out
```

- This image started giving error to end user

```
root@worker-1: ~
This is v2 pod kubeserve-7674cb6d4f-4dhqz
This is v3 pod kubeserve-5cdf574594-t25z6
This is v2 pod kubeserve-7674cb6d4f-4dhqz
This is v2 pod kubeserve-7674cb6d4f-n6xw4
Some internal error has occurred! This is pod kubeserve-5cdf574594-brsrb
Some internal error has occurred! This is pod kubeserve-5cdf574594-brsrb
This is v3 pod kubeserve-5cdf574594-f81mx
```

- In this case using below command we can roll back to previous version
- Syntax :> kubectl rollout undo deploy <deployment name>
- Ex:> kubectl rollout undo deploy kubeserve

```
root@manager-1:~# kubectl rollout undo deploy kubeserve
deployment.apps/kubeserve rolled back
root@manager-1:~# kubectl rollout status deploy kubeserve
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
deployment "kubeserve" successfully rolled out
```

- Now we can see rollback of previous version happening at end user window

```
root@worker-1: ~
Some internal error has occurred! This is pod kubeserve-5cdf574594-t75tj
Some internal error has occurred! This is pod kubeserve-5cdf574594-brsrb
Some internal error has occurred! This is pod kubeserve-5cdf574594-brsrb
This is v2 pod kubeserve-7674cb6d4f-rwcmf
This is v2 pod kubeserve-7674cb6d4f-rwcmf
Some internal error has occurred! This is pod kubeserve-5cdf574594-f81mx
Some internal error has occurred! This is pod kubeserve-5cdf574594-kplqx
```

```
This is v2 pod kubeserve-7674cb6d4f-9qsj9
This is v2 pod kubeserve-7674cb6d4f-9qsj9
This is v2 pod kubeserve-7674cb6d4f-ws5tx
This is v2 pod kubeserve-7674cb6d4f-s755k
This is v2 pod kubeserve-7674cb6d4f-s755k
This is v2 pod kubeserve-7674cb6d4f-s755k
```

- We can roll-back to specific version also using below command
- Syntax :> kubectl rollout undo deploy <deployment name> --to-revision=<deployment revision>
- Ex:> kubectl rollout undo deploy kubeserve --to-revision=1

```
root@manager-1:~# kubectl rollout undo deploy kubeserve --to-revision=1
deployment.apps/kubeserve rolled back
root@manager-1:~# kubectl rollout status deploy kubeserve
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 3 out of 5 new replicas have been updated...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "kubeserve" rollout to finish: 1 old replicas are pending termination...
deployment "kubeserve" successfully rolled out
```

- We can see at end user window roll-back to specific version i.e. v1 happening slowly

```
root@worker-1: ~
```

```
This is v2 pod kubeserve-7674cb6d4f-rwcmf
This is v2 pod kubeserve-7674cb6d4f-rwcmf
This is v2 pod kubeserve-7674cb6d4f-ws5tx
This is v1 pod kubeserve-6dc44d8bd-9x7gn
This is v1 pod kubeserve-6dc44d8bd-9x7gn
This is v1 pod kubeserve-6dc44d8bd-9x7gn
This is v1 pod kubeserve-6dc44d8bd-wjtsk
This is v1 pod kubeserve-6dc44d8bd-m6nql
This is v1 pod kubeserve-6dc44d8bd-m6nql
This is v1 pod kubeserve-6dc44d8bd-wjtsk
```

- We can see revision history also, it will show new revision number

```
root@manager-1:~# kubectl rollout history deploy kubeserve
deployment.apps/kubeserve
REVISION  CHANGE-CAUSE
3          <none>
4          <none>
5          <none>
```

- Deployment controller internally uses **replica set** only to maintain revision of changes.
- Whenever we do changes in any deployment controller's application version, kubernetes internally creates replica set for that version.
- Whenever we update new image to deployment, a new replica set will be created, that replica set only manages pods.
- By default in kubernetes 10 number of revision history can be remember, we can increase or decrease in that.
- So we can see that we have change our deployment controller with 3 different version of images i.e. v1, v2 & v3, so internally kubernetes has create 3 different replica set.
- As per our desired state, 5 numbers of pods of version v1 are running.

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
kubeserve-5cdf574594	0	0	0	65m	app	leaddevops/kubeserve:v3	app=kubeserve,pod-template-hash=5cdf574594
kubeserve-6dcd44d8bd	5	5	5	89m	app	leaddevops/kubeserve:v1	app=kubeserve,pod-template-hash=6dcd44d8bd
kubeserve-7674cb6d4f	0	0	0	85m	app	leaddevops/kubeserve:v2	app=kubeserve,pod-template-hash=7674cb6d4f

- If we rollout to any of specific version, suppose v2, that replica set version of pods will become in ready state gradually.

```
root@manager-1:~# kubectl get rs -o wide
NAME      DESIRED  CURRENT  READY   AGE     CONTAINERS   IMAGES
kubeserve-5cdf574594  0        0        0       85m    app         leaddevops/kubeserve:v3
kubeserve-6dcd44d8bd  5        5        5       89m    app         leaddevops/kubeserve:v1
kubeserve-7674cb6d4f  0        0        0       85m    app         leaddevops/kubeserve:v2
root@manager-1:~# kubectl rollout history deploy kubeserve
deployment.apps/kubeserve
REVISION  CHANGE-CAUSE
3          <none>
4          <none>
5          <none>
root@manager-1:~# kubectl rollout undo deploy kubeserve --to-revision=4
deployment.apps/kubeserve rolled back
root@manager-1:~# kubectl get rs -o wide
NAME      DESIRED  CURRENT  READY   AGE     CONTAINERS   IMAGES
kubeserve-5cdf574594  0        0        0       80m    app         leaddevops/kubeserve:v3
kubeserve-6dcd44d8bd  4        4        4       104m   app         leaddevops/kubeserve:v1
kubeserve-7674cb6d4f  3        3        3       101m   app         leaddevops/kubeserve:v2
root@manager-1:~# kubectl get rs -o wide
NAME      DESIRED  CURRENT  READY   AGE     CONTAINERS   IMAGES
kubeserve-5cdf574594  0        0        0       81m    app         leaddevops/kubeserve:v3
kubeserve-6dcd44d8bd  1        1        1       106m   app         leaddevops/kubeserve:v1
kubeserve-7674cb6d4f  5        5        5       102m   app         leaddevops/kubeserve:v2
root@manager-1:~# kubectl get rs -o wide
NAME      DESIRED  CURRENT  READY   AGE     CONTAINERS   IMAGES
kubeserve-5cdf574594  0        0        0       82m    app         leaddevops/kubeserve:v3
kubeserve-6dcd44d8bd  0        0        0       106m   app         leaddevops/kubeserve:v1
kubeserve-7674cb6d4f  5        5        5       102m   app         leaddevops/kubeserve:v2
```

25) What is DaemonSet in kubernetes? What are its usages?

- When there is requirement to mandatorily run a pod in all nodes of kubernetes cluster, than at that time we use Daemon controller.
- DaemonSet ensures that a copy of the Pod is running on all (or some) nodes.
- When a node joins the cluster, a Pod will be added to them.
- When a node is removed from the cluster, these Pods will also be recycled.
- Deleting a DaemonSet will delete all Pods it created.
- Some typical uses of a DaemonSet are:
 - running a cluster storage daemon on every node
 - running a logs collection daemon on every node
 - running a node monitoring daemon on every node
- Create DaemonSet using YML (daemonSet.yml)

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: mydaemonset
  labels:
    app: test
    env: prod
spec:
  selector:
    matchLabels:
      app: test
  template:
    metadata:
      name: mydaemonsetpod
      labels:
        app: test
  spec:
    containers:
      - name: mydaemonsetpod
        image: lerndevops/tomcat:8.5.47
        ports :
```

Classification: Public

```
- containerPort: 8080
resources:
  limits:
    memory : 200Mi
  requests:
    cpu: 100m
    memory : 200Mi
root@manager-1: ~
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: mydaemonset
  labels:
    app: test
    env: prod
spec:
  selector:
    matchLabels:
      app: test
  template:
    metadata:
      name: mydaemonsetpod
      labels:
        app: test
    spec:
      containers:
        - name: mydaemonsetpod
          image: lerndevops/tomcat:8.5.47
          ports :
            - containerPort: 8080
          resources:
            limits:
              memory : 200Mi
            requests:
              cpu: 100m
              memory : 200Mi
```

- Now execute below command
- Syntax:> kubectl create -f daemonSet.yml

```
root@manager-1: ~
root@manager-1:~# vim daemonSet.yml
root@manager-1:~# kubectl create -f daemonSet.yml
daemonset.apps/mydaemonset created
root@manager-1:~# kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
mydaemonset-8flbv   1/1    Running   0          39s
mydaemonset-m94zm   1/1    Running   0          39s
root@manager-1:~# kubectl get nodes
NAME     STATUS   ROLES      AGE   VERSION
manager-1  Ready    control-plane,master  6d10h  v1.23.10
worker-1  Ready    <none>    6d9h  v1.23.10
worker-2  Ready    <none>    101m  v1.23.10
root@manager-1:~# kubectl get ds
NAME   DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
mydaemonset  2         2         2         2           2           <none>       98s
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
mydaemonset-8flbv   1/1    Running   0        2m6s  192.168.226.120  worker-1  <none>        <none>
mydaemonset-m94zm   1/1    Running   0        2m6s  192.168.133.195  worker-2  <none>        <none>
root@manager-1:~#
```

- Here in our cluster as there are 2 worker nodes running, so there are only 2 DaemonSet pods are running.
- We can delete all DaemonSet or specific DaemonSet using below command
- Syntax:> kubectl delete ds --all
- Syntax:> kubectl delete ds --<name of DaemonSet>

```
root@manager-1: ~
root@manager-1:~# kubectl create -f daemonSet.yml
daemonset.apps/mydaemonset created
root@manager-1:~# kubectl get ds
NAME   DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE
mydaemonset  2         2         2         2           2           <none>       8s
root@manager-1:~# kubectl delete ds --all
daemonset.apps "mydaemonset" deleted
root@manager-1:~# kubectl delete ds mydaemonset
Error from server (NotFound): daemonsets.apps "mydaemonset" not found
root@manager-1:~#
```

Classification: Public

26) What is Job Controller in kubernetes? What are its usages?

- When we run a pod it continuously in running state, but when there is requirement to complete a task/ job & die after completion, in that situation we use Job Controller.
- It helps us to perform one single activity.
- Job controller ensures that one or more pods execute their commands and exit successfully.
- In Job controller to stop containers to bring it up again & again, we use key word **restartPolicy: Never** in specification.
- If Job controller pod is not getting start, kubernetes retries number of times, that we can specify using keyword **backoffLimit: 4**
- We can create job controller using below YML(job.yml)

```
apiVersion: batch/v1
kind: Job
metadata:
  name: countdown
spec:
  backoffLimit: 4
  template:
    metadata:
      name: countdown
    spec:
      containers:
        - name: counter
          image: centos:7
          command:
            - "bin/bash"
            - "-c"
            - "for i in 9 8 7 6 5 4 3 2 1 ; do echo $i ; done"
  restartPolicy: Never
```

```
root@manager-1: ~
apiVersion: batch/v1
kind: Job
metadata:
  name: countdown
spec:
  backoffLimit: 4
  template:
    metadata:
      name: countdown
    spec:
      containers:
        - name: counter
          image: centos:7
          command:
            - "bin/bash"
            - "-c"
            - "for i in 9 8 7 6 5 4 3 2 1 ; do echo $i ; done"
  restartPolicy: Never
```

- By running below command we can execute job controller
- Syntax> kubectl create -f job.yml
- We can check Job using below command
- Syntax> kubectl get jobs
- We can check job execution status by checking log of pod
- Syntax>kubectl get pods -o wide
- Syntax> kubectl logs <pod name>
- Ex> kubectl logs countdown-pxwqr
- We can use below command to delete all job
- Syntax>kubectl delete job --all
- We can use below command to specific job

- Syntax>kubectl delete job <job name>

```
root@manager-1: ~
root@manager-1:~# vim job.yml
root@manager-1:~# kubectl create -f job.yml
job.batch/countdown created
root@manager-1:~# kubectl get job
NAME      COMPLETIONS   DURATION   AGE
countdown  1/1          4s         11s
root@manager-1:~# kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
countdown-pxwqr  0/1    Completed   0          32s   192.168.226.121   worker-1   <none>   <none>
root@manager-1:~# kubectl logs countdown-pxwqr
9
8
7
6
5
4
3
2
1
root@manager-1:~# kubectl delete job --all
job.batch "countdown" deleted
root@manager-1:~# kubectl delete job countdown
Error from server (NotFound): jobs.batch "countdown" not found
root@manager-1:~#
```

27) What is Cron Job Controller in kubernetes? What are its usages?

- Suppose there is requirement to perform some task on regular interval, for example take backup of etcd, scale up some replica at some point of time, report generation etc. in that case we use CronJob controller.
- CronJob creates a schedule based on repeated intervalsJob.
- We have created a CronJob which run every 1 minute & print date-time with some message using YML(cron.yml)

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          name: mypod
          labels:
            app: myapp
        spec:
          restartPolicy: OnFailure
          containers:
            - name: hello
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
```

- Using below command we can create CronJob controller
- Syntax> kubectl create -f cron.yml
- We can get all CronJob controller using below command
- Syntax> kubectl get cj
- We can check log of pod to confirm CronJob controller has executed or not using below
- Syntax:> kubectl get pods
- Syntax:> kubectl logs <pod ID>
- We can delete all or specific CronJob controller using below command
- Syntax> kubectl delete cj --all
- Syntax> kubectl delete cj <CroneJob Controller name>

```
root@manager-1: ~
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        metadata:
          name: mypod
          labels:
            app: myapp
        spec:
          restartPolicy: OnFailure
          containers:
            - name: hello
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
```

```
root@manager-1: ~
root@manager-1: # vim cron.yml
root@manager-1: # kubectl create -f cron.yml
cronjob.batch/hello created
root@manager-1: # kubectl get cj
NAME   SCHEDULE   SUSPEND   ACTIVE   LAST SCHEDULE   AGE
hello  */1 * * *  False     0         <none>   30s
root@manager-1: ~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP          NODE   NOMINATED NODE   READINESS GATES
hello-28440245-cfjbm  0/1   Completed   0       23s  192.168.133.199  worker-2  <none>        <none>
root@manager-1: # kubectl logs hello-28440245-cfjbm
Sun Jan 28 04:05:01 UTC 2024
Hello from the Kubernetes cluster
root@manager-1: ~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP          NODE   NOMINATED NODE   READINESS GATES
hello-28440245-cfjbm  0/1   Completed   0       59s  192.168.133.199  worker-2  <none>        <none>
root@manager-1: ~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP          NODE   NOMINATED NODE   READINESS GATES
hello-28440245-cfjbm  0/1   Completed   0       84s  192.168.133.199  worker-2  <none>        <none>
hello-28440246-nnhzk  0/1   Completed   0       24s  192.168.226.65   worker-1  <none>        <none>
root@manager-1: ~# kubectl logs hello-28440246-nnhzk
Sun Jan 28 04:06:02 UTC 2024
Hello from the Kubernetes cluster
root@manager-1: ~# vim cron.yml
root@manager-1: ~# kubectl delete cj --all
cronjob.batch "hello" deleted
root@manager-1: ~# kubectl delete cj hello
Error from server (NotFound): cronjobs.batch "hello" not found
root@manager-1: ~#
```

28) What is HPA (Horizontal Pod Auto Scaler) Controller in kubernetes? What are its usages?

- When application received more traffic, pod get in hang state, because of high CPU & memory usage, so at that time HPA (HorizontalPodAutoscaler) automatically updates a workload resource, by increasing number of pods.
- HPA can monitor CPU & memory utilization intermittently & check if CPU/RAM usage crosses particular % (mark) new pods get created & when CPU/RAM usage goes normal, it automatically delete extra pods.
- In HPA we need to configure auto scaling matrix i.e. min & max number of pods to run depending on resource usage condition, & the target value of raising condition.
- We can set HPA running interval by the `--horizontal-pod-autoscaler-sync-period` parameter to the `kube-controller-manager`
- Job of HPA to monitor Deployment or ReplicaSet, so that auto creation or deletion of pods will happen depending on load condition
- HPA can work only with number of pods, not with number of nodes.
- If we want node base auto scaling we need to go with GKE,AKS etc.
- Horizontal scaling means that the response to increased load is to deploy more Pods.

- So to know memory or CPU utilization matrix of pod we need to install external software, which is by default not available in kubernetes.
- We can use below command to check CPU utilization of pod, but without installing matrix software we cannot get that details.
- Syntax:> kubectl top pods -n <name space>
- Ex:> kubectl top pods -n kube-system

```
root@manager-1:~# kubectl get pods -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE
calico-kube-controllers-66966888c4-fkfkfks  1/1    Running   0          7d2h
calico-node-jfnlr                  1/1    Running   0          7d2h
calico-node-n7598                  1/1    Running   0          17h
calico-node-q5pwm                  1/1    Running   0          7d2h
coredns-64897985d-lsv6p            1/1    Running   0          7d2h
coredns-64897985d-zz55h            1/1    Running   0          7d2h
etcd-manager-1                    1/1    Running   0          7d2h
kube-apiserver-manager-1           1/1    Running   0          7d2h
kube-controller-manager-manager-1   1/1    Running   0          7d2h
kube-proxy-82x71                  1/1    Running   0          7d2h
kube-proxy-jjcqw                  1/1    Running   0          7d2h
kube-proxy-sm5qc                  1/1    Running   0          17h
kube-scheduler-manager-1           1/1    Running   0          7d2h
root@manager-1:~# kubectl top pods -n kube-system
error: Metrics API not available
root@manager-1:~#
```

- So 1st step to set up HPA is install matrix server using below YML file
- Syntax:> kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/3-controllers/hpa/metric-server_1.23.yml
- Now we can get CPU & memory utilization of pods using below command
- Syntax:> kubectl top pods -n kube-system
- We can get CPU & memory utilization of nodes using below command
- Syntax:> kubectl top nodes

```
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/3-controllers/hpa/metric-server_1.23.yml
serviceaccount/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegate created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
service/metrics-server created
deployment.apps/metrics-server created
apiservice.apiregistration.k8s.io/vbeta1.metrics.k8s.io created
root@manager-1:~# kubectl top pods -n kube-system
NAME                               CPU(cores)   MEMORY(bytes)
calico-kube-controllers-66966888c4-fkfkfks  4m          23Mi
calico-node-jfnlr                  46m         89Mi
calico-node-n7598                  33m         78Mi
calico-node-q5pwm                  41m         98Mi
coredns-64897985d-lsv6p            3m          13Mi
coredns-64897985d-zz55h            2m          12Mi
etcd-manager-1                    19m         53Mi
kube-apiserver-manager-1           60m        261Mi
kube-controller-manager-manager-1   17m         56Mi
kube-proxy-82x71                  1m          13Mi
kube-proxy-jjcqw                  5m          15Mi
kube-proxy-sm5qc                  5m          12Mi
kube-scheduler-manager-1           4m          21Mi
metrics-server-7cb5455c67-8wjdv   65m         13Mi
root@manager-1:~# kubectl top nodes -n kube-system
NAME      CPU(cores)   CPU%    MEMORY(bytes)   MEMORY%
manager-1  289m        14%    1363Mi        35%
worker-1   148m        7%     1531Mi        40%
worker-2   146m        7%     1923Mi        50%
root@manager-1:~# kubectl top nodes
NAME      CPU(cores)   CPU%    MEMORY(bytes)   MEMORY%
manager-1  224m        11%    1360Mi        35%
worker-1   139m        6%     1530Mi        40%
worker-2   133m        6%     1925Mi        50%
root@manager-1:~#
```

- Now we can create HPA using below YML, in which we have created
 - simple nginx container having name **nginx** with minimum 3 number of replica & given 20m (mili) CPU
 - Service with name **nginx-svc** for above deployment
 - HPA with name **nginx-hpa** for above deployment, with minReplicas: 3 & maxReplicas: 10, & matrix we used as CPU utilization i.e. when average of CPU utilization of 3 pod reached to 10% of allocated CPU 20m = 2m it will create new pods.

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      name: nginxpod
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          resources:
            limits:
              cpu: 20m
---


apiVersion: v1
kind: Service
metadata:
  name: nginx-svc
spec:
  type: ClusterIP ## this is default if we do not type in service definition
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
---


apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: nginx-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: nginx
  minReplicas: 3
  maxReplicas: 10
  metrics:

```

- type: Resource
resource:
name: cpu
target:
type: Utilization
averageUtilization: 10
- Now run below command to create above application
- Syntax:> kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/3-controllers/hpa/hpa_1.23.yml
- Here we have created HPA, Service, deployment, pods, we can check all using below command
- Syntax:> kubectl get deploy,hpa,svc,pods
- We can also use below command, but it will not give detail about secrets & config map
- Syntax:> kubectl get all

```
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/3-controllers/hpa/hpa_1.23.yml
deployment.apps/nginx created
service/nginx-svc created
horizontalpodautoscaler.autoscaling/nginx-hpa created
root@manager-1:~# kubectl get deploy,hpa,svc,pods
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/nginx   3/3     3           3           2m13s

NAME          REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
horizontalpodautoscaler.autoscaling/nginx-hpa   Deployment/nginx   0%/10%    3          10         3           2m13s

NAME      TYPE    CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP     5h54m
service/nginx-svc  ClusterIP  10.110.197.198  <none>        80/TCP      2m13s

NAME      READY   STATUS   RESTARTS   AGE
pod/nginx-66f66f54b5-887rq  1/1     Running   0          2m13s
pod/nginx-66f66f54b5-gcjn8  1/1     Running   0          2m13s
pod/nginx-66f66f54b5-v6jgm  1/1     Running   0          2m13s
root@manager-1:~# kubectl get all
NAME           READY   STATUS   RESTARTS   AGE
pod/nginx-66f66f54b5-887rq  1/1     Running   0          7m39s
pod/nginx-66f66f54b5-gcjn8  1/1     Running   0          7m39s
pod/nginx-66f66f54b5-v6jgm  1/1     Running   0          7m39s

NAME      TYPE    CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP     6h
service/nginx-svc  ClusterIP  10.110.197.198  <none>        80/TCP      7m39s

NAME      READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/nginx   3/3     3           3           7m39s

NAME      DESIRED   CURRENT   READY   AGE
replicaset.apps/nginx-66f66f54b5  3       3       3       7m39s

NAME          REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
horizontalpodautoscaler.autoscaling/nginx-hpa   Deployment/nginx   0%/10%    3          10         3           7m39s
root@manager-1:~#
```

- Now by increasing load we can see 10 pods get auto created, so to do that we use load runner software.
- We can use Apache Bench (ab), it is a load testing and benchmarking tool for Hypertext Transfer Protocol (HTTP) server. It can be run from command line and it is very simple to use. A quick load testing output can be obtained in just one minute.
- Install this using below command
- Syntax:> apt-get install apache2-utils
- Install we need to get service IP to start load test first
- Syntax:> apt-get install apache2-utils
- Now increase load using below command
- Syntax:> ab -n 500000 -c 1000 http://10.110.197.198
- Here -n denotes Number of requests to perform & -c denotes Number of multiple requests to make at a time
- Here we can check HPA target % of CPU using below command

- Syntax:> kubectl get hpa

```
root@manager-1: ~
manager-1:~# kubectl get svc
  TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
etes   ClusterIP   10.96.0.1      <none>        443/TCP     6h40m
svc    ClusterIP   10.110.197.198   <none>        80/TCP      47m
manager-1:~# kubectl get hpa
  REFERENCE      TARGETS      MINPODS      MAXPODS      REPLICAS      AGE
hpa   Deployment/nginx  0%/10%      3            10           3          47m
manager-1:~# ab -n 500000 -c 1000 http://10.110.197.198/
  ApacheBench, Version 2.3 <$Revision: 1843412 $>
  Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
  http://www.apache.org/
  [idle] 10.110.197.198 (be patient)
  [idle] socket_recv: Connection refused (111)
  [idle] of 13722 requests completed
manager-1:~# kubectl get hpa
  REFERENCE      TARGETS      MINPODS      MAXPODS      REPLICAS      AGE
hpa   Deployment/nginx  70%/10%     3            10           3          48m
manager-1:~#
```

- We can check pods CPU usage goes high & automatically 10 replica of pods get created

```
root@manager-1: /
Every 2.0s: kubectl top pods                                         manager-1: Sun Jan 28 11:28:51 2024
NAME                  CPU(cores)      MEMORY(bytes)
nginx-66f66f54b5-42sbp  8m           2Mi
nginx-66f66f54b5-4txc4  0m           3Mi
nginx-66f66f54b5-77t54  7m           2Mi
nginx-66f66f54b5-887rq  13m          3Mi
nginx-66f66f54b5-gcjn8  16m          3Mi
nginx-66f66f54b5-vlc8m  8m           2Mi

root@manager-1: ~
Every 2.0s: kubectl get pods                                         manager-1: Sun Jan 28 11:29:18 2024
NAME      READY   STATUS    RESTARTS   AGE
nginx-66f66f54b5-42sbp  1/1     Running   0          60s
nginx-66f66f54b5-44rcg  1/1     Running   0          45s
nginx-66f66f54b5-4txc4  1/1     Running   0          13m
nginx-66f66f54b5-77t54  1/1     Running   0          60s
nginx-66f66f54b5-887rq  1/1     Running   0          55m
nginx-66f66f54b5-89xjq  1/1     Running   0          45s
nginx-66f66f54b5-gcjn8  1/1     Running   0          55m
nginx-66f66f54b5-12kbj  1/1     Running   0          45s
nginx-66f66f54b5-lqtdt  1/1     Running   0          45s
nginx-66f66f54b5-vlc8m  1/1     Running   0          60s
```

- Once load get reduce number of pods get down to min number 3

```
root@manager-1: /
Every 2.0s: kubectl top pods                                         manager-1: Sun Jan 28 11:30:00 2024
NAME                  CPU(cores)      MEMORY(bytes)
nginx-66f66f54b5-4txc4  0m           3Mi
nginx-66f66f54b5-887rq  0m           3Mi
nginx-66f66f54b5-gcjn8  0m           3Mi
```

29) What is dashboard in kubernetes? What are its usages?

- Dashboard is GUI software in kubernetes, which is used to manage kubernetes through GUI.
- It is also an image prepared by kubernetes.
- We can have many customized dashboards in kubernetes.
- We can install secure & insecure dashboards.
- We can install insecure dashboard using YAML as below.
- Syntax :> `kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/ui/dashboard/dashboard-insecure-v2.7.0.yaml`
- We can get all namespaces, using which we can get name of our dashboard namespace to check details of that namespace
- Syntax:> `kubectl get ns`
- Now using below commands we can get all components available in dashboard namespace
- Syntax:>`kubectl get all -n kubernetes-dashboard`

```
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/ui/dashboard/dashboard-insecure-v2.7.0.yaml
namespace/kubernetes-dashboard created
serviceaccount/kubernetes-dashboard created
service/kubernetes-dashboard created
secret/kubernetes-dashboard-certs created
secret/kubernetes-dashboard-csr created
secret/kubernetes-dashboard-key-holder created
configmap/kubernetes-dashboard-settings created
role.rbac.authorization.k8s.io/kubernetes-dashboard created
clusterrole.rbac.authorization.k8s.io/kubernetes-dashboard created
rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard created
clusterrolebinding.rbac.authorization.k8s.io/kubernetes-dashboard created
deployment.apps/kubernetes-dashboard created
service/dashboard-metrics-scraper created
deployment.apps/dashboard-metrics-scraper created
root@manager-1:~#
root@manager-1:~# kubectl get ns
NAME          STATUS   AGE
default        Active   7d7h
kube-node-lease Active   7d7h
kube-public    Active   7d7h
kube-system    Active   7d7h
kubernetes-dashboard Active  77s
sock-shop      Active   10h
root@manager-1:~# kubectl get all -n kubernetes-dashboard
NAME                           READY   STATUS    RESTARTS   AGE
pod/dashboard-metrics-scraper-6f669b9c9b-cjvrj   1/1     Running   0          112s
pod/kubernetes-dashboard-76958cccd64-lmjbh       1/1     Running   0          112s

NAME              TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/dashboard-metrics-scraper   ClusterIP  10.101.173.255 <none>        8000/TCP   113s
service/kubernetes-dashboard   NodePort    10.103.228.218  <none>        80:30009/TCP 113s

NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/dashboard-metrics-scraper   1/1      1           1           113s
deployment.apps/kubernetes-dashboard   1/1      1           1           113s

NAME           DESIRED   CURRENT   READY   AGE
replicaset.apps/dashboard-metrics-scraper-6f669b9c9b   1         1         1         113s
replicaset.apps/kubernetes-dashboard-76958cccd64        1         1         1         113s
root@manager-1:~#
```

- We can access dashboard application using node port service of dashboards
 - <http://34.72.240.122:30009/>
- We can access all resources of Kubernetes i.e. jobs, cron jobs, services, pods, replica sets, replication controllers, namespaces etc. in dashboard application.
- We can perform all tasks in dashboard, whatever we do through command line on all components of Kubernetes.
- But in organization we don't put edit option in dashboard, it is used to view only.

30) What is probes in kubernetes?

- If pod/container is showing running status, it doesn't mean it is serving request or it is healthy or working fine.
- Sometime pod/container is showing in running status, but internally it will be in hang status.
- To handle such situation in Kubernetes Probes will help us.
- Probe is nothing but health check condition of pod/container.
- There are 3 probes available in Kubernetes; **Liveness, Readiness and Startup Probes**

31) What is Liveness probes in kubernetes?

- The Liveness Probe serves as a diagnostic check to confirm if the container is alive.
- When the livenessProbe fails, Kubernetes considers the pod unhealthy, and then, attempts to restart it as a recovery measure.
- Let's take example of pod with nginx image in YML file(lp.yml), in which we put wrong condition to check connectivity on 800 port instead of 80 port

```
root@manager-1: ~
appVersion: v1
kind: Pod
metadata:
  name: lp-nginx-demo
  labels:
    app: nginx
spec:
  terminationGracePeriodSeconds: 0
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
      livenessProbe:
        tcpSocket:
          port: 800
        initialDelaySeconds: 5
        periodSeconds: 5
~
```

- Create pod with below command
- Syntax: > kubectl create -f lp.yml
- We can check pods status using below command, here we can see pod is getting auto restarts because of failing liveness condition
- Syntax: > kubectl get pods -o wide
- Syntax: > kubectl describe pod lp-nginx-demo

```
root@manager-1:~# kubectl create -f lp.yml
pod/lp-nginx-demo created
root@manager-1:~# kubectl get pods -o wide
error: the server doesn't have a resource type "pods"
root@manager-1:~# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS   GATES
lp-nginx-demo  1/1    Running   1 (15s ago)  35s   192.168.226.113  worker-1  <none>        <none>
root@manager-1:~# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS   GATES
lp-nginx-demo  1/1    Running   2 (11s ago)  51s   192.168.226.113  worker-1  <none>        <none>
```

Classification: Public

```
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS      RESTARTS   AGE    IP          NODE   NOMINATED-NODE   READINESS   GATES
lp-nginx-demo  0/1    CrashLoopBackOff  7 (79s ago)  8m49s  192.168.226.113  worker-1  <none>        <none>
root@manager-1:~# kubectl describe pod lp-nginx-demo
Name:         lp-nginx-demo
Namespace:    default
Priority:    0
Node:        worker-1/10.128.0.3
Start Time:  Mon, 29 Jan 2024 18:32:43 +0000
Labels:      app=nginx
Annotations:  cni.projectcalico.org/containerID: 99a7af7efb2b55d547d68ce68d000da84dc6a6f111917313637d8f67a274e132
              cni.projectcalico.org/podIP: 192.168.226.113/32
              cni.projectcalico.org/podIPs: 192.168.226.113/32
Status:      Running
IP:          192.168.226.113
IPs:
  IP: 192.168.226.113
Containers:
  nginx:
    Container ID: docker://758006ff76571113bec449803272fa3cb7ac001e52530a11cb7dc85cf729a856
    Image:        nginx
    Image ID:    docker-pullable://nginx@sha256:4c0fd8aa8b6341bfdeca5f18f7837462c80cff90527ee35ef185571e1c327beac
    Port:        80/TCP
    Host Port:   0/TCP
    State:       Waiting
      Reason:    CrashLoopBackOff
    Last State:  Terminated
      Reason:    Completed
      Exit Code:  0
      Started:   Mon, 29 Jan 2024 18:39:56 +0000
      Finished:  Mon, 29 Jan 2024 18:40:13 +0000
    Ready:      False
    Restart Count: 7
    Liveness:   tcp-socket :800 delay=5s timeout=1s period=5s #success=1 #failure=3
    Environment: <none>
Events:
  Type  Reason  Age            From          Message
  ----  -----  --            --          --
  Normal Scheduled  8m59s        default-scheduler  Successfully assigned default/lp-nginx-demo to worker-1
  Normal Pulled   8m58s        kubelet        Successfully pulled image "nginx" in 446.541386ms
  Normal Pulled   8m38s        kubelet        Successfully pulled image "nginx" in 357.875867ms
  Normal Created   8m18s (x3 over 8m58s)  kubelet        Created container nginx
  Normal Started   8m18s (x3 over 8m58s)  kubelet        Started container nginx
  Normal Pulled   8m18s        kubelet        Successfully pulled image "nginx" in 386.228379ms
  Warning Unhealthy 7m59s (x9 over 8m49s)  kubelet        Liveness probe failed: dial tcp 192.168.226.113:800: connect: connection refused
  Normal Killing   7m59s (x3 over 8m39s)  kubelet        Container nginx failed liveness probe, will be restarted
  Normal Pulling   7m59s (x4 over 8m58s)  kubelet        Pulling image "nginx"
  Warning BackOff  3m57s (x16 over 7m19s)  kubelet        Back-off restarting failed container
root@manager-1:~#
```

32) What is Readiness probs in kubernetes?

- Sending request to POD
- Suppose we create DB pod & its status is showing ready, but it is not confirm that it is ready to accept request because it might be its internal DB configuration are yet to start, so in that case we configure readiness prob.
- It will not restart container, it just confirm pod is ready to accept request.
- The kubelet uses readiness probes to know when a container is ready to start accepting traffic.
- A Pod is considered ready when all of its containers are ready.
- One use of this signal is to control which Pods are used as backends for Services. When a Pod is not ready, it is removed from Service load balancers.
- Let's create YML (rp.yaml) in which we have set up pod with nginx image & to check readiness prob we have places get request to 800 port instead of 80 port.

Classification: Public

Classification: Public

```
root@manager-1: ~
apiVersion : v1
kind: Pod
metadata:
  name: rp-nginx-demo
  labels:
    app : nginx
spec:
  terminationGracePeriodSeconds: 0
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
      readinessProbe:
        httpGet:
          path: /
          port: 800
      initialDelaySeconds: 5
      periodSeconds: 5
~
```

- Now we can create pod using below command:
- Syntax: > kubectl create -f rp.yml
- Here we can see that it is not ready to accept request, using below command
- Syntax: > kubectl get pods -o wide
- If we have expose this pod through service & check its end point, than we cannot get its end point
- Syntax:> kubectl expose pod rp-nginx-demo --name svc1 --port 80 --target-port 80
- Syntax:> kubectl describe svc svc1

```
root@manager-1: ~
root@manager-1:~# vim rp.yml
root@manager-1:~# kubectl create -f rp.yml
pod/rp-nginx-demo created
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
rp-nginx-demo  0/1     Running   0          13s    192.168.226.112  worker-1  <none>        <none>
root@manager-1:~# kubectl expose pod rp-nginx-demo --name svc1 --port 80 --target-port 80
service/svc1 exposed
root@manager-1:~# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP    2d
svc1       ClusterIP  10.111.69.204  <none>        80/TCP     17s
root@manager-1:~# kubectl describe svc svc1
Name:           svc1
Namespace:      default
Labels:         app=nginx
Annotations:   <none>
Selector:       app=nginx
Type:          ClusterIP
IP Family Policy: SingleStack
IP Families:   IPv4
IP:            10.111.69.204
IPs:           10.111.69.204
Port:          <unset>  80/TCP
TargetPort:    80/TCP
Endpoints:     <none>
Session Affinity: None
Events:        <none>
root@manager-1:~#
```

33) What is startup probs in kubernetes?

- The kubelet uses startup probes to know when a container application has started.
- If such a probe is configured, liveness and readiness probes do not start until it succeeds.
- It making sure above probes doesn't interfere with the application startup.

Classification: Public

34) How many ways we can schedule to run pod in kubernetes?

- We can schedule to run pod in below 3 different ways
 - NodeName: Schedule to run pod on specific node
 - NodeSelector: Schedule to run pod on specific node base on matching labels
 - NodeAffinity: Schedule to run pod base on if else condition.

35) What is NodeName? How can we schedule to run pod on specific node in kubernetes?

- By using **node name** we can deploy our pod in specific node.
- Suppose we have node with name worker-2 in which we want to deploy our pod, so we can write YML as below (spn.yml)

```
root@manager-1: ~
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubeserv
spec:
  replicas: 2
  selector:
    matchLabels:
      app: kubeserve
  template:
    metadata:
      name: kubeserve
      labels:
        app: kubeserve
    spec:
      nodeName: worker-2
      containers:
        - image: leaddevops/kubeserve:v1
          name: app
```

- Now we create deployment of this using below command
- Syntax:> kubectl create -f spn.yml
- Here we can see our pod is running on worker-2 only using below command
- Syntax:> kubectl get pod -o wide

```
Select root@manager-1: ~
root@manager-1:~# vim spn.yml
root@manager-1:~# kubectl create -f spn.yml
deployment.apps/kubeserv created
root@manager-1:~# kubectl get node -o wide
NAME           STATUS   ROLES      AGE     VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE       KERNEL-VERSION   CONTAINER-RUNTIME
manager-1      Ready    control-plane,master   12d    v1.23.10  10.128.0.2   <none>        Ubuntu 20.04.6 LTS  5.15.0-1048-gcp  docker://25.0.0
worker-1       Ready    <none>     12d    v1.23.10  10.128.0.3   <none>        Ubuntu 20.04.6 LTS  5.15.0-1048-gcp  docker://25.0.0
worker-2       Ready    <none>     5d22h   v1.23.10  10.128.0.5   <none>        Ubuntu 20.04.6 LTS  5.15.0-1049-gcp  docker://25.0.1
root@manager-1:~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED-NODE   READINESS   GATES
kubeserv-68c8d6f74c-69vxv  1/1    Running   0          2m53s   192.168.133.230  worker-2  <none>        <none>
kubeserv-68c8d6f74c-jhpfh  1/1    Running   0          2m53s   192.168.133.229  worker-2  <none>        <none>
nginx-demo     0/1    Running   0          43h     192.168.226.112  worker-1  <none>        <none>
```

- Here API server directly tells kubelete where to run pod, hence there is no use to scheduler in node name.
- IF node gets deleted, than that pod will be in pending state.

36) What is NodeSelector? How can we schedule to run pod on specific node based on some condition in kubernetes?

- In node name if node gets deleted, than that pod will be in pending state, because we wanted to go in specific pod.
- So to solve above issue **node selector** came into picture, suppose we have multiple nodes in kubernetes system with their own property such as environment like DEV, UAT & PROD, so we can assign label to each node as env:DEV, env:UAT, env:PROD & depending on label condition we can deploy our application using node selector.
- For example lets create label color=green on one node i.e. worker-1, first we can check labels of all node using below command

- Syntax:> kubectl get nodes --show-labels
- We can set label using below command:
- Syntax:>kubectl label node <node name> <label>
- Ex:> kubectl label node worker-1 color=green

```
root@manager-1:~# kubectl label node worker-1 color=green
node/worker-1 labeled
root@manager-1:~# kubectl get nodes --show-labels
NAME      STATUS   ROLES      AGE     VERSION   LABELS
manager-1  Ready    control-plane,master   13d    v1.23.10  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=manager-1,kubernetes.io/os=linux,node-role.kubernetes.io/control-planes,,node-role.kubernetes.io/masters=,node.kubernetes.io/exclude-from-external-load-balancers=
worker-1   Ready    <none>     13d    v1.23.10  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,color=green,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker-1,kubernetes.io/os=linux
worker-2   Ready    <none>     7d3h   v1.23.10  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker-2,kubernetes.io/os=linux
root@manager-1:~#
```

- We can create deployment with selector condition to run our pod on specific node using below YML file

```
root@manager-1: ~
apiVersion: apps/v1
kind: Deployment
metadata:
  name: kubeserv
spec:
  replicas: 2
  selector:
    matchLabels:
      app: kubeserv
  template:
    metadata:
      name: kubeserv
      labels:
        app: kubeserv
    spec:
      nodeSelector:
        color: "green"
      containers:
        - image: leaddevops/kubeserve:v1
          name: app
```

- Using below command we can run & check pods is running on which node
- Syntax:> kubectl create -f nodeSelector.yml
- Syntax:> kubectl get pod -o wide

```
root@manager-1: ~
root@manager-1:~# vim nodeSelector.yml
root@manager-1:~# kubectl create -f nodeSelector.yml
deployment.apps/kubeserv created
root@manager-1:~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
kubeserv-68dfd659df-42mjc  1/1    Running   0          61s    192.168.226.127  worker-1  <none>        <none>
kubeserv-68dfd659df-sq2zn  1/1    Running   0          61s    192.168.226.66  worker-1  <none>        <none>
root@manager-1:~#
```

- When we want to deploy our pod on multiple nodes we can use node selector.

37) What is NodeAffinity?

- Schedule to run pod base on if else condition is as NodeAffinity.
- We can apply condition as requiredDuringSchedulingIgnoredDuringExecution This means that the pod will get scheduled only on a node that satisfy condition
- Also we can apply condition as preferredDuringSchedulingIgnoredDuringExecution This means that the pod will get scheduled on a node that satisfy condition if not satisfied we can schedule on another node.

38) What is Taint & Toleration in Kubernetes?

39) What is Volumes in Kubernetes?

- Disk inside container are lasting for a very short time, due to which there are some of the problem for running application within container
- One problem when container crash or stopped all files inside that were created or modified during the lifetime of the container are lost because during a crash, kubelet restarts the container with a clean state
- Another problem occurs when multiple containers are running in a Pod and need to share files, it can be challenging to setup and access a shared filesystem across all of the containers.
- To resolve above problem Volumes concept is introduced.
- Kubernetes supports many types of volumes, **Ephemeral volume** types have a lifetime of a pod, but **persistent volumes** exist beyond the lifetime of a pod.
- At its core, a volume is a directory, possibly with some data in it, which is accessible to the containers in a pod.
- Below are types of volumes used in kubernetes
 - Empty Directory
 - Host Path
 - Persistent Volume(PV)
 - Persistent Volume Claim(PVC)
 - Config map
 - Secret

40) What is Empty Directory Volumes in Kubernetes?

- It is a type of volume that is created when a Pod is first assigned to a Node.
- It remains active as long as the Pod is running on that node.
- The volume is initially empty and the containers in the pod can read and write the files in the emptyDir volume.
- Each container in this volume can share files using mount path
- Once the Pod is removed from the node, the data in the emptyDir is erased.
- Let's create one pod having 3 different containers images
 - Centos having mount path name "/demo1"
 - Ubuntu having mount path name "/demo2"
 - Alpine having mount path name "/demo3"
- All 3 different containers share common mount, so that they can share files with each other
- Create pod using YML & below command
- Syntax:> kubectl create -f <https://raw.githubusercontent.com/lerndevops/kubernetes/master/5-storage/volumes/emptyDir-vol-ex2.yml>
- We can enter into each container & do read/write files of each other by entering into container using below command
- Syntax:> kubectl exec -it <pod name> -c <container name> -- bash
- First login in centos container & create file in share mount using below command
- Syntax:> exec -it myvolumes-pod myvolumes-container-1 -- bash
- Syntax:> echo centos >/demo1/t1.txt
- Syntax:> ls /demo1/
- Now login in Ubuntu container & create file in share mount using below command
- Syntax:> exec -it myvolumes-pod myvolumes-container-2 -- bash
- Syntax:> echo ubuntu >/demo2/t2.txt

- Syntax:> ls /demo2/
- Now login in Alpine container & create file in share mount using below command
- Syntax:> exec -it myvolumes-pod myvolumes-container-3 -- bash
- Syntax:> echo alpine >/demo3/t3.txt
- Syntax:> ls /demo3/

```
root@manager-1: ~
apiVersion: v1
kind: Pod
metadata:
  name: myvolumes-pod
  namespace: default
spec:
  containers:
    - command:
        - sh
        - -
        - echo The Bench Container 1 is Running ; sleep 3600
      image: centos
      imagePullPolicy: IfNotPresent
      name: myvolumes-container-1
      volumeMounts:
        - mountPath: /demo1
          name: demo-volume
    - command:
        - /bin/bash
        - -
        - echo The Bench Container 2 is Running ; sleep 3600
      image: ubuntu
      imagePullPolicy: IfNotPresent
      name: myvolumes-container-2
      volumeMounts:
        - mountPath: /demo2
          name: demo-volume
    - command:
        - sh
        - -
        - echo The Bench Container 3 is Running ; sleep 3600
      image: alpine
      imagePullPolicy: IfNotPresent
      name: myvolumes-container-3
      volumeMounts:
        - mountPath: /demo3
          name: demo-volume
  volumes:
    - emptyDir: {}
      name: demo-volume
root@manager-1: ~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED-NODE   READINESS   GATES
myvolumes-pod  3/3    Running   213 (6m40s ago)  2d23h  192.168.226.126  worker-1  <none>        <none>
root@manager-1: ~# kubectl exec -it myvolumes-pod -- bash
Defaulted container "myvolumes-container-1" out of: myvolumes-container-1, myvolumes-container-2, myvolumes-container-3
[root@myvolumes-pod /]# ls /demo1/
/demo1/: file not recognized: Is a directory
[root@myvolumes-pod /]# ls /demo1/
[root@myvolumes-pod /]# echo cento >t1.txt
[root@myvolumes-pod /]# ls /demo1/
[root@myvolumes-pod /]# echo cento > /demo1/t1.txt
[root@myvolumes-pod /]# ls /demo1/
t1.txt
[root@myvolumes-pod /]# exit
root@manager-1: ~# kubectl exec -it myvolumes-pod -c myvolumes-container-2 -- bash
root@myvolumes-pod:/# ls /demo2/
t1.txt
root@myvolumes-pod:/# echo cento > /demo1/t2.txt
bash: /demo1/t2.txt: No such file or directory
root@myvolumes-pod:/# echo cento > /demo2/t2.txt
root@myvolumes-pod:/# ls /demo2/
t1.txt t2.txt
```

41) What is Host Path Volumes in Kubernetes?

- A hostPath volume mounts a file or directory from the host node's file system into your Pod.
- To use hostPath volume we need to configure 2 properties :
 - Path : directory location on host machine
 - Type: this is optional, if we have not specified any value, kubernetes will automatically create folder on given path of host machine where pod will run.
 - If we have specified type: Directory, it is must to have directory on given path
 - If we have specified type: DirectoryOrCreate, if nothing exists at the given path, an empty directory will be created on given path of host machine where pod will run.
 - Other Type are File, Socket, CharDevice, BlockDevice
- Some uses for a hostPath are:

- Running a container that needs access to node-level system components (such as a container that transfers system logs to a central location, accessing those logs using a read-only mount of `/var/log`).
- Making a configuration file stored on the host system available read-only to a static pod; unlike normal Pods, static Pods cannot access ConfigMaps.
- let's create nginx container whose log folder path we mount in host machine path as below YML

```
root@manager-1: ~
apiVersion: v1
kind: Pod
metadata:
  name: hostpath-vol-pod1
spec:
  containers:
    - name: test-container
      image: nginx
      volumeMounts:
        - name: hostpath-volume
          mountPath: /var/log/nginx
  volumes:
    - name: hostpath-volume
      hostPath:
        path: /tmp/logs
```

- Now create POD using below command

```
root@manager-1:~# vim hp.yml
root@manager-1:~# kubectl create -f hp.yml
pod/hostpath-vol-pod1 created
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
hostpath-vol-pod1   1/1    Running   0          15s    192.168.133.240   worker-2   <none>        <none>
myvolumes-pod     3/3    Running   6 (46m ago) 166m   192.168.226.65   worker-1   <none>        <none>
root@manager-1:~# kubectl exec -it hostpath-vol-pod1 -- bash
root@hostpath-vol-pod1:~# ls /var/log/nginx/
access.log  error.log
root@hostpath-vol-pod1:~#
```



```
root@worker-2: ~
root@worker-2:~# ls /tmp/logs/
access.log  error.log
root@worker-2:~#
```

- We can see pod is running on worker-2 node & check all logs available in container location(`/var/log/nginx`) by entering into pod container are mounted in host machine(worker-2) path (`/tmp/logs`)

42) What is Persistent Volume (PV), Persistent Volume (PVC) & Storage Class in Kubernetes?

PersistentVolume (PV)

- A **PersistentVolume (PV)** is a piece of storage in the cluster that has been created by an administrator or dynamically provisioned using Storage Classes.
- It is a resource/object in the cluster just like a node is a cluster resource.
- Persistent volumes are independent of the lifecycle of the pod that uses it, meaning that even if the pod shuts down, the data in the volume is not erased.
- In order for a Pod to start using these volumes, it must request a volume by issuing a **persistent volume claim (PVC)**.
- **PVCs** describe the storage capacity and characteristics a pod requires, and the cluster attempts to match the request and provision the desired persistent volume.
- The PV having below specification attributes
- **Capacity:** this is used to set PV capacity
- **VolumeMode:** Kubernetes supports two volume modes of persistent volumes. A valid value for volume mode can be either Filesystem or Block. Filesystem is the default mode if the volume mode is not defined.

- **AccessMode:**
 - ReadOnlyMany(ROX) allows being mounted by multiple nodes in read-only mode.
 - ReadWriteOnce(RWO) allows being mounted by a single node in read-write mode.
 - ReadWriteMany(RWX) allows multiple nodes to be mounted in read-write mode.
- A volume can only be mounted using one access mode at a time, even if it supports many access modes.
- **Class:** A PV can specify a StorageClass to dynamically bind the PV and PVC, where the specific StorageClass is specified via the storageClassName property. If no PV is specified with this property, it can only bind to PVC that does not require a specific class.
- **Reclaim Policy:** When the node no longer needs persistent storage, the reclaiming strategies that can be used include:
 - **Retain** - meaning the PV, until deleted, is kept alive.
 - **Recycle** - meaning the data can be restored later after getting scrubbed.
 - **Delete** - associated storage assets (such as AWS EBS, GCE PD, Azure Disk, and OpenStack Cinder volumes) are deleted.
- **Mount Option:** Kubernetes administrators can specify mount options for mounting persistent volumes on a node. Not all PV types support mount options. Common types of mount options supported are:
 - gcePersistentDisk
 - awsElasticBlockStore
 - AzureDisk
 - NFS
 - RBD (Rados Block Device)
 - CephFS
 - Cinder (OpenStack volume storage)
 - Glusterfs

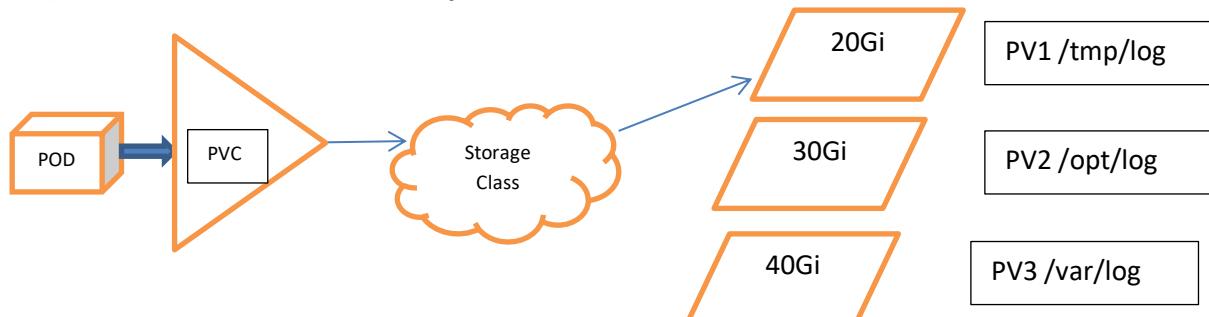
PersistentVolumeClaim (PVC)

- This is a request sent by a Kubernetes node for storage.
- Persistent volume claim is a request for storage usage by a Kubernetes developer.
- The claim can include specific storage parameters required by the application—for example an amount of storage, or a specific type of access (read/write, read-only, etc.).

Storage Class

- It is an object in kubernates which is help to manage PV allocation to pods.
- The StorageClass object allows cluster administrators to define PVs with different properties, like performance, size or access parameters.
- It lets you expose persistent storage to users while abstracting the details of storage implementation.
- There are many predefined StorageClasses in Kubernetes (see the following section), or you can create your own.

43) What is Persistent Volume lifecycle in Kubernetes?



- The interaction between PVs and PVCs follows a distinct lifecycle, starting with provisioning and including binding, using, and reclaiming.

Provisioning

Here are the two main types of provisioning:

- Static - PVs created by a cluster administrator. These PVs are located in the Kubernetes API and contain information about storage resources available to cluster users.
- Dynamic - to meet PVC demands, clusters can attempt to dynamically provision a volume. This typically occurs when available static PVs do not match the PVC. In this case, the cluster can optionally be configured to use the default StorageClasses.

Binding

- A user creates a PVC request, specifying a certain amount of storage and the required access modes.
- A control loop, located in the master, looks for new PVCs, and then:
 - **Static** - the control loop attempts to find a matching PV and then binds it to the PVC.
 - **Dynamic** - if a PV was already dynamically provisioned for the PVC, the control loop binds them together. After PVC and a PV are bound, they remain exclusive.
- PVC to PV binding is a one-to-one mapping. The process uses a ClaimRef, which creates a bi-directional binding between the PV and the PVC.

Using

- Pods use claims as volumes.
- The cluster inspects the claim to find the bound volume and mounts that volume for a Pod.
- Once a user has a claim and that claim is bound, the bound PV belongs to the user for as long as they need it.

Reclaiming

- A reclaim policy for PVs tells the cluster what to do with the volume after the claim is released, and volumes can be Retained, Recycled, or Deleted.
- Once users do not need their volume anymore, they can delete the PVC objects from the API that allows reclamation of the resource.

Example

- Let's create one sample application which is having tomcat container, write log on PV using PVC
- Create PV using below YML(my-pv.yml)
- Syntax:> kubectl create -f my-pv.yml
- Check PV using below command
- Syntax:> kubectl get pv

```
root@manager-1: ~
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  storageClassName: standard
  persistentVolumeReclaimPolicy: Recycle
  capacity:
    storage: 1000Mi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /tmp/pvdata
```

```
root@manager-1: ~
root@manager-1:~# vim my-pv.yml
root@manager-1:~# kubectl create -f my-pv.yml
persistentvolume/my-pv created
root@manager-1:~# kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS     CLAIM   STORAGECLASS   REASON   AGE
my-pv    1000Mi      RWO            Recycle        Available     standard   3m50s
root@manager-1:~#
```

- Create PVC using below YML(my-pvc.yml), in which we have ask(request) for minimum 500Mi size storage

```
root@manager-1: ~
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mv-pvc
spec:
  storageClassName: standard
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
```

- Syntax:> kubectl create -f my-pvc.yml
- We can check pvc using below command
- Syntax:> kubectl get pvc
- We can check both pv & pvc using below command
- Syntax:> kubectl get pv,pvc

```
root@manager-1: ~
root@manager-1:~# vi my-pvc.yml
root@manager-1:~# kubectl create -f my-pvc.yml
persistentvolumeclaim/my-pvc created
root@manager-1:~# kubectl get pvc
NAME      STATUS   VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE
my-pvc   Bound    my-pv    1000Mi     RWO          standard       11s
root@manager-1:~# kubectl get pv,pvc
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS     CLAIM   STORAGECLASS   REASON   AGE
persistentvolume/my-pv  1000Mi      RWO            Recycle        Bound     default/my-pvc standard   95s
NAME      STATUS   VOLUME   CAPACITY   ACCESS MODES   STORAGECLASS   AGE
persistentvolumeclaim/my-pvc Bound    my-pv    1000Mi     RWO          standard       21s
root@manager-1:~#
```

- Now let's run pod using below yml (pod-pv.yml) which is having tomcat container that uses this PVC (my-pvc), which is already bound to PV (my-pv).
- So on whatever node this pod gets run, PV will bind logs of tomcat on host path (/tmp/pvdata).

```
root@manager-1: ~
apiVersion: v1
kind: Pod
metadata:
  name: pvc-pod
spec:
  containers:
    - name: my-pvc-container
      image: tomcat:8.5
      volumeMounts:
        - name: pv
          mountPath: "/usr/local/tomcat/logs"
  volumes:
    - name: pv
      persistentVolumeClaim:
        claimName: my-pvc
```

- Now run pod & check on host where this pod gets run.
- Syntax: > kubectl create -f pod-pv.yml

```
root@manager-1: ~
root@manager-1:~# vim pod-pv.yml
root@manager-1:~# kubectl create -f pod-pv.yml
pod/pvc-pod created
root@manager-1:~# kubectl get pod -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
pvc-pod   1/1     Running   0          11s    192.168.133.244   worker-2   <none>        <none>
root@manager-1:~#
```

- We can see pod is running on worker-2, so we can check on host path (/tmp/pvdata) on that to get log by login in worker-2

```
root@worker-2: /#
root@worker-2:~/# ls -ltrh /tmp/pvdata/
total 8.0K
-rw-r----- 1 root root    0 Feb 18 08:30 manager.2024-02-18.log
-rw-r----- 1 root root    0 Feb 18 08:30 localhost.2024-02-18.log
-rw-r----- 1 root root    0 Feb 18 08:30 host-manager.2024-02-18.log
-rw-r----- 1 root root    0 Feb 18 08:30 localhost_access_log.2024-02-18.txt
-rw-r----- 1 root root 5.0K Feb 18 08:30 catalina.2024-02-18.log
```

- Here in above example we have created PV which uses host path storage, it has many disadvantages :
- As pod run on multiple nodes, we cannot access all pod data at single point.
- If pod is reschedule to run on other node its previous node data is gone.
- So to overcome above it is recommended to use common storage i.e. NFS(Network File Storage)
- We need to create NFS server → create PV by using that NFS mount instead of host path → create PVC → create POD

44) How can we supply dynamic data to pod in Kubernetes?

- By using **ConfigMap & Secret** we can supply dynamic data in kubernetes.
- Both are kubernetes object

45) What is ConfigMap in Kubernetes? How can we create & use it?

- Many applications rely on configuration which is used during either application initialization or runtime.
- Most of the times there is a requirement to adjust values assigned to configuration parameters.
- ConfigMaps** is the kubernetes way to inject application pods with configuration data.
- ConfigMaps allow you to decouple configuration artifacts from image content to keep containerized applications portable.

- A ConfigMap is an API object used to store non-confidential data in key-value pairs in the cluster store (ETCD).
- Pods can consume ConfigMaps as
 - environment variables,
 - command-line arguments
 - As configuration files in a volume.
- We can use ConfigMap to setup different environment i.e. DEV/QC/UAT/PROD with same image but different configuration file.
- We can create ConfigMap using below command
 - Syntax:> kubectl create configmap <map-name> <data-source>
 - Syntax:> kubectl create cm <map-name> <data-source>
- We can create 2 different configMap for 2 different environment as below
 - Ex:> kubectl create cm cm1 --from-literal=username=san-dev
 - Ex:> kubectl create cm cm2 --from-literal=username=san-uat

```
c# root@manager-1: ~
root@manager-1:~# kubectl create cm cm1 --from-literal=username=san-dev
configmap/cm1 created
root@manager-1:~# kubectl create cm cm2 --from-literal=username=san-uat
configmap/cm2 created
root@manager-1:~# kubectl get cm
NAME        DATA   AGE
cm1          1    42s
cm2          1    5s
kube-root-ca.crt 1    64s
```

- We can create multiple properties configuration in file(config) & create configMap as below

```
c# root@manager-1: ~
username: san-dev
port: 1234
host: www.myhost.com
~
```

- Ex:> kubectl create cm cm3 --from-file=config
- We can describe ConfigMap as below
- Ex:> kubectl describe cm cm3

```
c# root@manager-1: ~
root@manager-1:~# vim config
root@manager-1:~# kubectl create cm cm3 --from-file=config
configmap/cm3 created
root@manager-1:~# kubectl describe cm cm3
Name:          cm3
Namespace:     default
Labels:        <none>
Annotations:   <none>

Data
====

config:
-----
username: san-dev
port: 1234
host: www.myhost.com

BinaryData
=====

Events:  <none>
```

- We can create ConfigMap using YML as below

```
c# Select root@manager-1: ~
apiVersion: v1
data:
  config: |
    username: san-dev
    port: 1234
    host: www.myhost.com
kind: ConfigMap
metadata:
  creationTimestamp: null
  name: cm3
  type: "Opaque"
```

Classification: Public

- Let's create nginx container in which we initialize index.html file using config map using below YML file

```
c$ root@manager-1: ~
apiVersion: v1
kind: ConfigMap
metadata:
  name: dev-html
data:
  index.html: |
    <H1> Hello Dev ENV ConfigMap </H1>
---
kind: Pod
apiVersion: v1
metadata:
  name: cm-pod-ex1
spec:
  containers:
    - name: cm-pod-cont
      image: nginx
      ports:
        - containerPort: 80
      volumeMounts:
        - name: html-from-cm
          mountPath: /usr/share/nginx/html/
          readOnly: true
  volumes:
    - name: html-from-cm
      configMap:
        name: dev-html
```

- We can create ConfigMap & pod using below command

```
c$ Select root@manager-1: ~
root@manager-1:~# vim cm-demo1.yml
root@manager-1:~# kubectl create -f cm-demo1.yml
configmap/dev-html created
pod/cm-pod-ex1 created
root@manager-1:~# kubectl get cm
NAME      DATA   AGE
cm1       1      58m
cm2       1      58m
cm3       1      52m
dev-html   1      47s
kube-root-ca.crt 1      59m
root@manager-1:~# kubectl describe cm dev-html
error: unknown command "describe" for "kubectl"

Did you mean this?
  describe
root@manager-1:~# kubectl describe cm dev-html
Name:      dev-html
Namespace:  default
Labels:    <none>
Annotations: <none>

Data
=====
index.html:
-----
<H1> Hello Dev ENV ConfigMap </H1>

BinaryData
=====

Events:  <none>
```

- We can check in POD, our index.html is used or not by entering into it as below

```
c$ root@manager-1: ~
root@manager-1:~# kubectl get pod -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
cm-pod-ex1  1/1    Running   0          17m    192.168.133.246  worker-2   <none>        <none>
root@manager-1:~# kubectl exec -it cm-pod-ex1 -- bash
root@cm-pod-ex1:~/# cat /usr/share/nginx/html/index.html
<H1> Hello Dev ENV ConfigMap </H1>
```

Classification: Public

Classification: Public

- Let's create nginx container in which we use some environment variables using config map using below YML file & command

```
root@manager-1:~  
apiVersion: v1  
kind: ConfigMap  
metadata:  
  name: env-cm  
data:  
  database: mongodb  
  database_uri: mongodb://localhost:27017  
---  
apiVersion: v1  
kind: Pod  
metadata:  
  name: cm-env-pod  
spec:  
  containers:  
    - name: env-var-configmap  
      image: nginx  
      envFrom:  
        - configMapRef:  
            name: env-cm
```

```
root@manager-1:~# vim cm-demo2.yml  
root@manager-1:~# kubectl create -f cm-demo2.yml  
configmap/env-cm created  
pod/cm-env-pod created  
root@manager-1:~# kubectl get cm  
NAME          DATA   AGE  
cm1           1      94m  
cm2           1      94m  
cm3           1      87m  
dev-html       1      36m  
env-cm        2      19s  
kube-root-ca.crt 1      95m  
root@manager-1:~# kubectl describe cm env-cm  
Name:         env-cm  
Namespace:    default  
Labels:       <none>  
Annotations:  <none>  
  
Data  
====  
database:  
----  
mongodb  
database_uri:  
----  
mongodb://localhost:27017  
  
BinaryData  
====  
  
Events:  <none>
```

- We can check environment variable value by entering inside containers & executing commands as below

```
root@manager-1:~# kubectl get pod -o wide  
NAME        READY   STATUS    RESTARTS   AGE     IP          NODE    NOMINATED NODE   READINESS GATES  
cm-env-pod  1/1     Running   0          16s    192.168.133.248  worker-2  <none>        <none>  
cm-pod-ex1  1/1     Running   0          45m    192.168.133.246  worker-2  <none>        <none>  
root@manager-1:~# kubectl exec -it cm-env-pod -- bash  
root@cm-env-pod:/# env  
KUBERNETES_SERVICE_PORT_HTTPS=443  
database_uri=mongodb://localhost:27017  
KUBERNETES_SERVICE_PORT=443  
HOSTNAME=cm-env-pod  
PWD=/  
PKG_RELEASE=1~bookworm  
HOME=/root  
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443  
database=mongodb  
NJS_VERSION=0.8.3  
TERM=xterm  
SHELL=/bin/sh  
KUBERNETES_PORT_443_TCP_PROTO=tcp  
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1  
KUBERNETES_SERVICE_HOST=10.96.0.1  
KUBERNETES_PORT=tcp://10.96.0.1:443  
KUBERNETES_PORT_443_TCP_PORT=443  
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin  
NGINX_VERSION=1.25.4  
_= /usr/bin/env  
root@cm-env-pod:/# echo jdbc:odbc:$database_uri  
jdbc:odbc:mongodb://localhost:27017  
root@cm-env-pod:/#
```

Classification: Public

46) What is secret in Kubernetes? How can we create & use it?

- A Secret is an object that contains a small amount of sensitive data such as a password, a token, or a key.
- Such information might otherwise be put in a Pod specification or in a container image.
- Using a Secret means that you don't need to include confidential data in your application code.
- By default in kubernetes it used base 64 encoding to store secrets, once we used secret in pod it will automatically decode this value.
- We can create secrets by command line as below

- Syntax:> kubectl create secret generic <secret name> <data source key-value>
- Ex:> kubectl create secret generic sc1 --from-literal=pass=abc@123
- Ex:> kubectl get secret
- Ex:> kubectl describe secret sc1

```
root@manager-1: ~
root@manager-1:~# kubectl create secret generic sc1 --from-literal=pass=abc@123
secret/sc1 created
root@manager-1:~# kubectl get secret
NAME          TYPE        DATA   AGE
default-token-ssgb6  kubernetes.io/service-account-token  3      2m4s
sc1           Opaque      1      22s
root@manager-1:~# kubectl describe secret sc1
Name:         sc1
Namespace:    default
Labels:       <none>
Annotations: <none>
Type:        Opaque
Data
=====
pass: 7 bytes
root@manager-1:~#
```

- We can create secrets & use that by below YML & command

```
root@manager-1: ~
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: ZXhhbXBsZS11c2Vy
  password: ZXhhbXBsZS1wYXNzd29yZA==

apiVersion: v1
kind: Pod
metadata:
  name: sec-pod-demo
spec:
  containers:
  - name: mycontainer
    image: redis
    volumeMounts:
    - name: myvol
      mountPath: "/apps/secrets"
      readOnly: true
  volumes:
  - name: myvol
    secret:
      secretName: mysecret
root@manager-1: ~
```

```
root@manager-1:~# vim sec-demo.yml
root@manager-1:~# kubectl create -f sec-demo.yml
secret/mysecret created
pod/sec-pod-demo created
root@manager-1:~# kubectl get secret
NAME          TYPE        DATA   AGE
default-token-kzn7t  kubernetes.io/service-account-token  3      3m43s
mysecret      Opaque      2      32s
root@manager-1:~# kubectl explain secret mysecret
error: We accept only this format: explain RESOURCE
root@manager-1:~# kubectl describe secret mysecret
Name:         mysecret
Namespace:    default
Labels:       <none>
Annotations: <none>
Type:        Opaque
Data
=====
password: / 16 bytes
username: / 12 bytes
root@manager-1:~#
```

- We can check secret value by login in side container as below

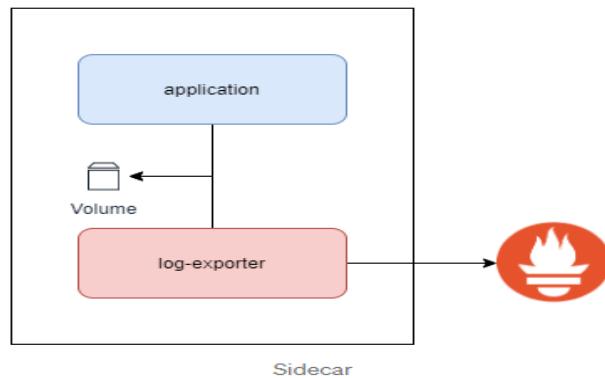
```
root@manager-1: ~
root@manager-1:~# kubectl exec -it sec-pod-demo -- bash
root@sec-pod-demo:/data# cd /
root@sec-pod-demo:/# ls /apps/secrets/
password  username
root@sec-pod-demo:/# cat /apps/secrets/username
example-userroot@sec-pod-demo:#
root@sec-pod-demo:/# cat /apps/secrets/password
example-passwordroot@sec-pod-demo:#
root@sec-pod-demo:/#
```

47) What is multi container Pod in Kubernetes? How can we create & use it?

- Pods can contain multiple containers, for some excellent reasons — primarily, the fact that containers in a pod get scheduled in the same node in a multi-node cluster. This makes communication between them faster and more secure.
- Containers within a pod can interact with each other in various ways:
- **Network:** Containers can access any listening ports on containers within the same pod, even if those ports are not exposed outside the pod.
- **Shared Storage Volumes:** Containers in the same pod can be given the same mounted storage volume, which allows them to interact with the same files.
- **Shared Process Namespace:** Process namespace sharing can be enabled by setting `shareProcessNamespace true` in the pod spec. This allows containers within the pod to interact with, and signal, one another's processes.
- We can run multi container pod using below pattern :
 - Sidecar
 - Ambassador
 - Adapter

48) What is Side Car Pattern in multi container Pod in Kubernetes? How can we create & use it?

- Sidecars derive their name from motorcycle sidecars. While your motorcycle can work fine without the sidecar, having one enhances or extends the functionality of your bike, by giving it an extra seat.
- Similarly, in Kubernetes, a sidecar pattern is used to enhance or extend the existing functionality of the container.
- Your container works perfectly well without the sidecar, but with it, it can perform some extra functions.
- So, a sidecar container enhances the main container in some way, adding functionality to it.
- Some examples are :
 - Using a sidecar for monitoring and logging and adding an agent for these purposes
 - A sidecar periodically syncs files in a webserver container's file system from a Git repository.
- Let's create nginx container as main container which displays HTML page with current date time, and its side car container debian OS whose role is to create HTML file with current date time,
- Debian OS will share HTML file to Nginx using emptyDir shared volume concept.



```
root@manager-1: ~
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mc1
spec:
  volumes:
  - name: html
    emptyDir: {}
  containers:
  - name: 1st
    image: nginx
    volumeMounts:
    - name: html
      mountPath: /usr/share/nginx/html
  - name: 2nd
    image: debian
    volumeMounts:
    - name: html
      mountPath: /html
    command: ["/bin/sh", "-c"]
    args:
    - while true; do
      date >/html/index.html
      sleep 1;
    done
```

- Create pod using below command:
- Syntax: kubectl create -f sidecar.yml
- Create svc to access nginx using below command
- Syntax: kubectl expose pod mc1 --name svc1 --port 1234 --target-port 80 --type NodePort

```
root@manager-1: ~
```

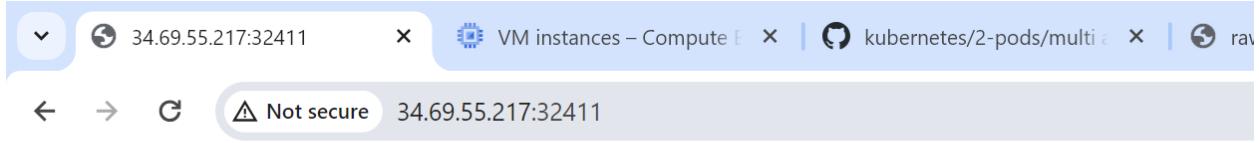
```
root@manager-1:~# vim sidecar.yml
26L, 436C written
root@manager-1:~# kubectl create -f sidecar.yml
pod/mc1 created
root@manager-1:~# kubectl get pod -o wide
NAME    READY  STATUS   RESTARTS   AGE     IP          NODE   NOMINATED-NODE   READINESS   GATES
mc1    2/2    Running   0          30s    192.168.226.84   worker-1   <none>        <none>
root@manager-1:~# kubectl label pod mc1 app=myapp
pod/mc1 labeled
root@manager-1:~# kubectl expose pod mc1 --name svc1 --port 1234 --target-port 80 --type NodePort
service/svc1 exposed
root@manager-1:~# kubectl get svc
NAME    TYPE      CLUSTER-IP    EXTERNAL-IP    PORT(S)        AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP       5m26s
svc1    NodePort  10.101.116.38  <none>        1234:32411/TCP 32s
root@manager-1:~#
```

- We can login into debian container & create "index.html" file in shared path /html/index.html

```
root@manager-1: ~
```

```
root@manager-1:~# kubectl exec -it mc1 -c 2nd -- bash
root@mc1:/# echo "<h1>Hello Kubernetes</h1>" >> /html/index.html
root@mc1:/#
```

- We can check this html page create by debins OS is used to display nginx page by accessing in browser by using port & host IP <http://34.69.55.217:32411/>

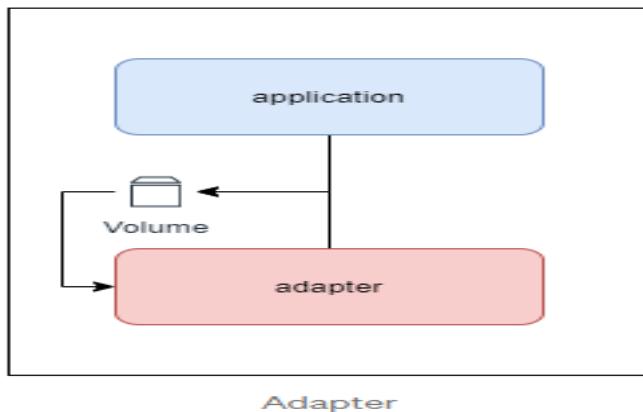


Hello Kuber

Hello Kernet

49) What is Adapter Pattern in multi container Pod in Kubernetes? How can we create & use it?

- The Adapter is another pattern that you can implement with multiple containers.
- The adapter pattern helps you standardize something heterogeneous in nature.
- For example, you're running multiple applications within separate containers, but every application has a different way of outputting log files.
- Now, if you have a centralized logging system that accepts logs in a particular format only, what can you do in such a situation? Well, you can either change the source code of each application to output a standard log format or use an adapter to standardize the logs before sending it to your central server. That's where the adapter pattern comes in.
- So, an adapter container transforms the output of the main container.
- Some great examples are: An adapter container reads log output from the main container and transforms it.



- Let's create pod using below YML having 2 apline container 1st one i.e. **app-container** will generate log, 2nd one i.e. **adapter-container** take log of 1st & format in stander format.
- After running this pod we can login in **adapter-container** & check log in location /var/log, we can found **app-container** logs available in **adapter-container**

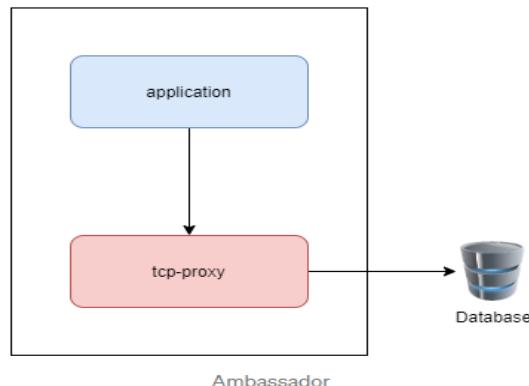
```

root@manager-1:~#
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-adapter
spec:
  volumes:
    - name: shared-logs
      emptyDir: {}
  containers:
    - name: app-container
      image: alpine
      command: ["/bin/sh"]
      args: ["-c", "while true; do date > /var/log/top.txt && top -n 1 -b >> /var/log/top.txt; sleep 5;done"]
      volumeMounts:
        - name: shared-logs
          mountPath: /var/log
    - name: adapter-container
      image: alpine
      command: ["/bin/sh"]
      args: ["-c", "while true; do (cat /var/log/top.txt | head -1 > /var/log/status.txt) && (cat /var/log/top.txt | head -2 | tail -1 | grep -o '\d+\w*' | head -1 >> /var/log/status.txt)
&& (cat /var/log/top.txt | head -3 | tail -1 | grep -o '\d+\w*' | head -1 >> /var/log/status.txt); sleep 5; done"]
      volumeMounts:
        - name: shared-logs
          mountPath: /var/log
root@manager-1:~#
root@manager-1:~# vim adapter.yml
root@manager-1:~# kubectl create -f adapter.yml
pod/pod-with-adapter created
root@manager-1:~# kubectl get pod -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
pod-with-adapter   2/2    Running   0          9s   192.168.226.69   worker-1   <none>        <none>
root@manager-1:~# kubectl exec -it pod-with-adapter -c adapter-container -- sh
/ # ls /var/log/
status.txt  top.txt
/ # cat /var/log/status.txt
Sun Mar  3 11:57:38 UTC 2024
1346768K
0%
/ #

```

50) What is Ambassador Pattern in multi container Pod in Kubernetes? How can we create & use it?

- The ambassador pattern derives its name from an Ambassador, who is an envoy and a person a country chooses to represent their country and connect with the rest of the world.
- Similarly, in the Kubernetes perspective, an Ambassador pattern implements a proxy to the external world.
- Let's look at an example — if you build an application that needs to connect with a database server, the server configuration, etc, changes with the environment.
- Now, the official recommendation to handle these is to use Config Maps, but what if you have legacy code that is already using another way of connecting to the database. Maybe, a properties file, or even worse, a hardcoded set of values. What if you want to communicate with localhost, and you can leave the rest to the admin? You can use the Ambassador pattern for these kinds of scenarios.

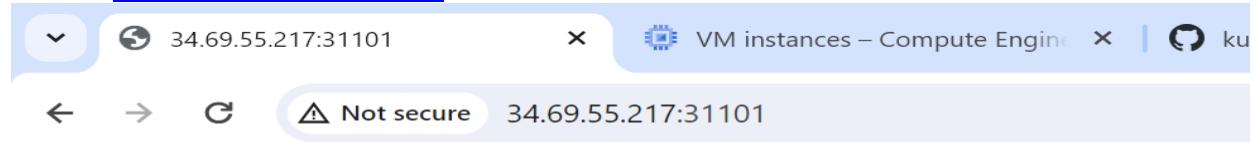


51) What is Service Load Balancing in Kubernetes?

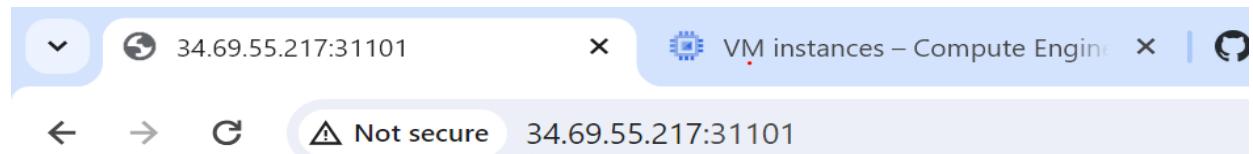
- Load Balancer service is an extension of NodePort service. NodePort and ClusterIP Services, to which the external load balancer routes are automatically created.
- It integrates NodePort with cloud-based load balancers.
- It exposes the Service externally using a cloud provider's load balancer.
- Each cloud provider (AWS, Azure, GCP, etc) has its own native load balancer implementation. The cloud provider will create a load balancer, which then automatically routes requests to your Kubernetes Service.
- Traffic from the external load balancer is directed at the backend Pods. The cloud provider decides how it is load balanced.
- The actual creation of the load balancer happens asynchronously.
- Every time you want to expose a service to the outside world, you have to create a new LoadBalancer and get an IP address.
- Let's create one deployment having some python application image & we scale this application up to 4 replicas
- Now expose this application using Load balancer service, if this application could have been running on Google Kuberne Engine at External IP instead of pending we could seen virtual IP of Google cloud.

```
c4 root@manager-1: ~
root@manager-1:~# kubectl create deploy dep1 --image lerndevops/samplepyapp:v1
deployment.apps/dep1 created
root@manager-1:~# kubectl scale deploy dep1 --replicas 4
deployment.apps/dep1 scaled
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE   READINESS GATES
dep1-5b5f6c8f8d-2jpt9  1/1    Running   0          8s    192.168.133.255  worker-2  <none>        <none>
dep1-5b5f6c8f8d-b4ds2  1/1    Running   0          8s    192.168.133.193  worker-2  <none>        <none>
dep1-5b5f6c8f8d-t4crg  1/1    Running   0          39s   192.168.226.83   worker-1  <none>        <none>
dep1-5b5f6c8f8d-vlwg4  1/1    Running   0          8s    192.168.226.88   worker-1  <none>        <none>
root@manager-1:~# kubectl expose deploy dep1 --name lb --port 2345 --target-port 3000 --type LoadBalancer
service/lb exposed
root@manager-1:~# kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP      2m6s
lb        LoadBalancer  10.109.214.248  <pending>      2345:31101/TCP  11s
root@manager-1:~#
```

- But we can access this application using our master/worker nodes IP & highlighted port <http://34.69.55.217:31101>



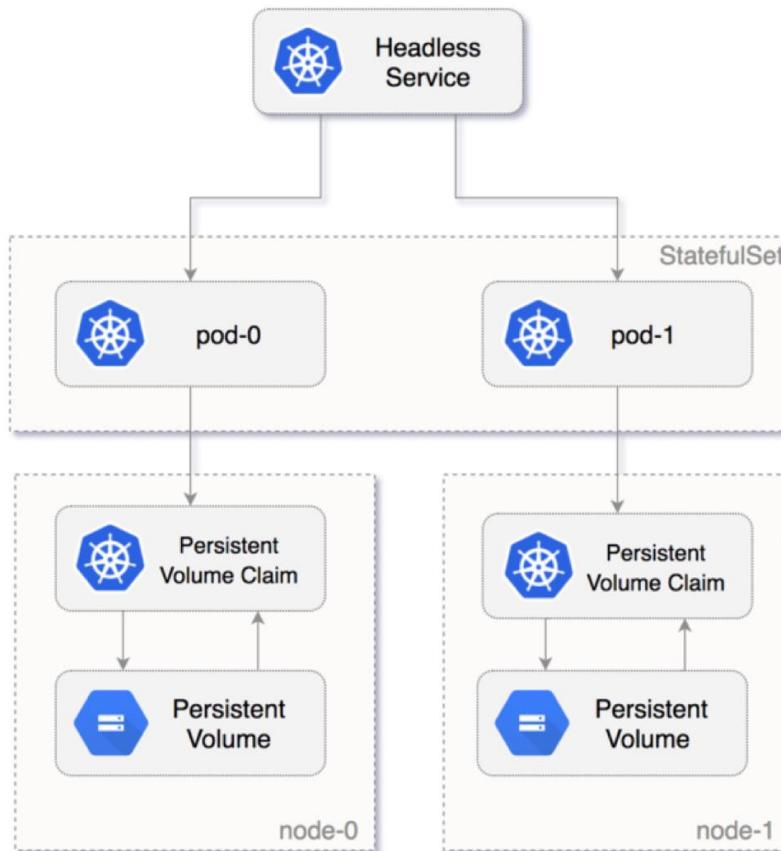
this is my V1 app - Container ID: dep1-5b5f6c8f8d-2jpt9



this is my V1 app - Container ID: dep1-5b5f6c8f8d-b4ds2

52) What is Statefull sets in Kubernetes?

- A stateful application is a program that saves client data from the activities of one session for use in the next session. The data that is saved is called the application's state. OR
- Stateful applications require that data that is used or generated by the application is persisted, retained, backed up and accessible outside of the particular hosts that run the application
- Statefull applications may expect to have static hostname of its replicas
- Statefull applications may expect to start from a start point not as a whole.
- All databases, such as MySQL, Oracle, and PostgreSQL, are examples of stateful applications.
- Kubernetes is a container orchestration tool that uses many controllers to run applications as containers (Pods). One of these controllers is called StatefulSet, which is used to run stateful applications. i.e. Data Base
- A StatefulSet is the Kubernetes controller used to run the stateful application as containers (Pods) in the Kubernetes cluster.
- StatefulSets assign a sticky identity—an ordinal number starting from zero—to each Pod instead of assigning random IDs for each replica Pod.
- A new Pod is created by cloning the previous Pod's data. If the previous Pod is in the pending state, then the new Pod will not be created.
- If you delete a Pod, it will delete the Pod in reverse order, not in random order. For example, if you had four replicas and you scaled down to three, it will delete the Pod numbered 3.



When to Use StatefulSets :

- There are several reasons to consider using StatefulSets. Here are two examples:
- Assume you deployed a MySQL database in the Kubernetes cluster and scaled this to three replicas, and a frontend application wants to access the MySQL cluster to read and write data. The read request will be forwarded to three Pods. However, the write request will only be forwarded to the first (primary) Pod, and the data will be synced with the other Pods. You can achieve this by using StatefulSets.
- Deleting or scaling down a StatefulSet will not delete the volumes associated with the stateful application. This gives you your data safety. If you delete the MySQL Pod or if the MySQL Pod restarts, you can have access to the data in the same volume.

53) How Pod to Pod communications work in Kubernetes?

Networking in Kubernetes:

- Since a Kubernetes cluster consists of various nodes and pods, understanding how they communicate between them is essential. The Kubernetes networking model supports different types of open source implementations.
- The Kubernetes networking model, on the other hand, natively supports multi-host networking in which pods are able to communicate with each other by default, regardless of which host they live in.
- Kubernetes does not provide a default network implementation; it only enforces a model for third-party tools to implement.

Pod to Pod communication:

- Kubernetes follows an “IP-per-pod” model where each pod get assigned an IP address and all containers in a single pod share the same network namespaces and IP address. Containers in the same pod can therefore reach each other’s ports via localhost.
- However, it is not recommended to communicate directly with a pod via its IP address due to pod’s volatility (a pod can be killed and replaced at any moment).
- Instead, use a Kubernetes service which represents a group of pods acting as a single entity to the outside. Services get allocated their own IP address in the cluster and provide a reliable entry point.
- Kubernetes services allow grouping pods under a common access policy (for example, load-balanced). The service gets assigned a virtual IP which pods outside the service can communicate with. Those requests are then transparently proxied (via the **kube-proxy** component that runs on each node) to the pods inside the service.
- Different proxy-modes are supported:
- **iptables**: kube-proxy installs iptables rules trap access to service IP addresses and redirect them to the correct pods.
- **userspace**: kube-proxy opens a port (randomly chosen) on the local node. Requests on this “proxy port” get proxied to one of the service’s pods (as retrieved from Endpoints API).
- **ipvs** (from Kubernetes 1.9): calls netlink interface to create ipvs rules and regularly synchronizes them with the Endpoints API.
- These kube-proxy must be running in all node of cluster
- We can see all kube proxy pod running in system namespace of cluster using below command
- Syntax:> kubectl get pods -o wide -n kube-system

```
root@manager-1:~# kubectl get pods -o wide -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED NODE   READINESS GATES
calico-kube-controllers-66966888c4-fkfk8   1/1    Running   3 (8d ago)  48d   192.168.11.141  manager-1  <none>        <none>
calico-node-jfnlr                   1/1    Running   3 (8d ago)  48d   10.128.0.2    manager-1  <none>        <none>
calico-node-n7598                  1/1    Running   3 (6d9h ago) 41d   10.128.0.5    worker-2   <none>        <none>
calico-node-q5pwm                 1/1    Running   5 (6d9h ago) 48d   10.128.0.3    worker-1   <none>        <none>
coredns-64897985d-lsv6p            1/1    Running   3 (8d ago)  48d   192.168.11.139  manager-1  <none>        <none>
coredns-64897985d-zz55n            1/1    Running   3 (6d9h ago) 48d   192.168.11.140  manager-1  <none>        <none>
etcd-manager-1                    1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-apiserver-manager-1          1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-controller-manager-manager-1  1/1    Running   5 (6d9h ago) 48d   10.128.0.3    worker-1   <none>        <none>
kube-proxy-82x7l                 1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-proxy-jjcqw                1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-proxy-sm5qc                1/1    Running   3 (6d9h ago) 41d   10.128.0.5    worker-2   <none>        <none>
kube-scheduler-manager-1         1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
metrics-server-7cb5455c67-8wjdv  1/1    Running   3 (8d ago)   41d  192.168.133.251  worker-2  <none>        <none>
root@manager-1:~#
```

54) What is DNS in Kubernetes? How it is used for service & pods?

- Kubernetes provides its own DNS service to resolve domain names inside the cluster in order for pods to communicate with each other.
- This is implemented by deploying a regular Kubernetes service which does name resolution inside the cluster, and configuring individual containers to contact the DNS service to resolve domain names.
- Note that this “internal DNS” is compatible and expected to run along with the cloud provider’s DNS service.
- Every service gets assigned a DNS name which resolves to the cluster IP of the service. The naming convention includes the service name and its namespace.
- For example: my-service.my-namespace.svc.cluster.local
- A pod inside the same namespace as the service does not need to fully qualify its name, for example a pod in my-namespace could lookup this service with a DNS query for my-service, while a pod outside my-namespace would have to query for my-service.my-namespace.
- We can see DNS software (pod) in side system namespace, which is used for all DNS communication within cluster i.e. pod name resolution & service name resolution.

```
root@manager-1:~# kubectl get pods -o wide -n kube-system
NAME                               READY   STATUS    RESTARTS   AGE     IP           NODE      NOMINATED NODE   READINESS GATES
calico-kube-controllers-66966888c4-fkfk8   1/1    Running   3 (8d ago)  48d   192.168.11.141  manager-1  <none>        <none>
calico-node-jfnlr                   1/1    Running   3 (8d ago)  48d   10.128.0.2    manager-1  <none>        <none>
calico-node-n7598                  1/1    Running   3 (6d9h ago) 41d   10.128.0.5    worker-2   <none>        <none>
calico-node-q5pwm                 1/1    Running   5 (6d9h ago) 48d   10.128.0.3    worker-1   <none>        <none>
coredns-64897985d-lsv6p            1/1    Running   3 (8d ago)  48d   192.168.11.139  manager-1  <none>        <none>
coredns-64897985d-zz55n            1/1    Running   3 (6d9h ago) 48d   192.168.11.140  manager-1  <none>        <none>
etcd-manager-1                    1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-apiserver-manager-1          1/1    Running   3 (8d ago)   48d  10.128.0.2    manager-1  <none>        <none>
kube-controller-manager-manager-1  1/1    Running   5 (6d9h ago) 48d   10.128.0.3    worker-1   <none>        <none>
kube-proxy-82x7l                 1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-proxy-jjcqw                1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
kube-proxy-sm5qc                1/1    Running   3 (6d9h ago) 41d   10.128.0.5    worker-2   <none>        <none>
kube-scheduler-manager-1         1/1    Running   3 (6d9h ago) 48d   10.128.0.2    manager-1  <none>        <none>
metrics-server-7cb5455c67-8wjdv  1/1    Running   3 (8d ago)   41d  192.168.133.251  worker-2  <none>        <none>
root@manager-1:~# kubectl get rs -o wide -n kube-system
NAME              DESIRED   CURRENT   READY   AGE     CONTAINERS   IMAGES                                     SELECTOR
calico-kube-controllers-66966888c4   1         1         1       48d    calico-kube-controllers   docker.io/calico/kube-controllers:v3.24.1   k8s-app=calico-kube-controllers,pod-template-hash=64897985d
coredns-64897985d                     2         2         2       48d    coredns               k8s.gcr.io/coredns:coredns:v1.8.6        k8s-app=kube-dns,pod-template-hash=64897985d
metrics-server-7cb5455c67             1         1         1       41d    metrics-server       k8s.gcr.io/metrics-server/metrics-server:v0.5.0   k8s-app=metrics-server,pod-template-hash=7cb5455c67
root@manager-1:~# kubectl get deploy -o wide -n kube-system
NAME              READY   UP-TO-DATE   AVAILABLE   AGE     CONTAINERS   IMAGES                                     SELECTOR
calico-kube-controllers   1/1    1           1       48d    calico-kube-controllers   docker.io/calico/kube-controllers:v3.24.1   k8s-app=calico-kube-controllers
coredns             2/2    2           2       48d    coredns               k8s.gcr.io/coredns:coredns:v1.8.6        k8s-app=kube-dns
metrics-server      1/1    1           1       41d    metrics-server       k8s.gcr.io/metrics-server/metrics-server:v0.5.0   k8s-app=metrics-server
root@manager-1:~#
```

- Let's take example to check how service & pod can be access :
- Step1: create pod nginx & service nginx in a namespace called "testns"

```
root@manager-1: ~
root@manager-1:~# kubectl create namespace testns
namespace/testns created
root@manager-1:~# kubectl -n testns run nginx --image=nginx --expose --port=80
service/nginx created
pod/nginx created
root@manager-1:~# kubectl get all -n testns -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
pod/nginx  1/1     Running   0          24s    192.168.133.195  worker-2  <none>        <none>
NAME        TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)   AGE     SELECTOR
service/nginx ClusterIP  10.98.252.58  <none>       80/TCP    24s    run=nginx
root@manager-1:~#
```

- Step2: create pod mynginx & service mynginx in a namespace called "default"

```
root@manager-1: ~
root@manager-1:~# kubectl -n default run mynginx --image=nginx --expose --port=80
service/mynginx created
pod/mynginx created
root@manager-1:~# kubectl get all -n default -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
pod/mynginx  1/1     Running   0          15s    192.168.226.87  worker-1  <none>        <none>
NAME        TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)   AGE     SELECTOR
service/kubernetes ClusterIP  10.96.0.1    <none>       443/TCP  34h    <none>
service/mynginx ClusterIP  10.105.57.127  <none>       80/TCP   15s    run=mynginx
root@manager-1:~#
```

- Step3: create another pod in any namespace which has netutils so that we can run a nslookup command

```
root@manager-1: ~
root@manager-1:~# kubectl -n testns run dnstest --image=lerndevops/samples:netshoot --rm -it -- /bin/bash
If you don't see a command prompt, try pressing enter.
bash-5.1#
```

- The above command will create new pod called dnstest in testns namespace & get you inside it, from inside pod we can try accessing pods using their service dns & pod dns
- Step4: test Service DNS:
 - if we try to access **nginx** service using name, we can see it will work because **nginx** service is available in local

```
root@manager-1: ~
root@manager-1:~# kubectl -n testns run dnstest --image=lerndevops/samples:netshoot --rm -it -- /bin/bash
If you don't see a command prompt, try pressing enter.
bash-5.1# cat /etc/resolv.conf
nameserver 10.96.0.10
search testns.svc.cluster.local svc.cluster.local cluster.local us-central1-a.c.googleapis.com
options ndots:5
bash-5.1# nslookup nginx
Server: 10.96.0.10
Address: 10.96.0.10#53
Name: nginx.testns.svc.cluster.local
Address: 10.98.252.58
```

- if we try to access **mynginx** service using name which is in some other namespace i.e. default, we can see it will not work because **mynginx** service is not available in local, so to access this service we need to use namespace name also i.e. **mynginx.default.svc**

```
bash-5.1# nslookup mynginx
Server: 10.96.0.10
Address: 10.96.0.10#53
** server can't find mynginx: NXDOMAIN
bash-5.1# nslookup mynginx.default.svc
Server: 10.96.0.10
Address: 10.96.0.10#53
Name: mynginx.default.svc.cluster.local
Address: 10.105.57.127
bash-5.1#
```

- Step4: test Pod DNS:
 - to make pod DNS we need get the podip & replace "." with "-" then add .namespace.pod
 - Example: convert 1.2.3.4 to 1-2-3-4 & add .namespace.pod
 - > 1-2-3-5.testns.pod
 - So we can access local **nginx** pod whose IP is : 192.168.133.195 we can use DNS as **192-168-133-195.testns.pod**
 - Similarly we can access default namespace's pod **mynginx** whose IP is : 192.168.226.87 using DNS as **192-168-226-87.default.pod**

```
root@manager-1: ~
bash-5.1# nslookup 192-168-133-195.testns.pod
Server: 10.96.0.10
Address: 10.96.0.10#53

Name: 192-168-133-195.testns.pod.cluster.local
Address: 192.168.133.195

bash-5.1# nslookup 192-168-226-88.default.pod
Server: 10.96.0.10
Address: 10.96.0.10#53

Name: 192-168-226-88.default.pod.cluster.local
Address: 192.168.226.88
```

55) What is network policy in Kubernetes? How it is used for pods communication?

- A network policy is a specification of how groups of pods are allowed to communicate with each other and other network endpoints.
- Network policies allow you to specify which pods can talk to other pods.
- This helps when securing communication between pods, allowing you to identify **ingress** (incoming traffic) and **egress** (outgoing traffic) rules.
- Network policy is created at namespace level, not cluster level.
- We can apply a network policy to a pod by using **labels** or namespace **selectors**.
- We can even choose a CIDR block range to apply the network policy.
- By Default, all pods in the cluster can communicate with any other pod and reach out to any available IP.
- Network Policies allow you to limit what network traffic is allowed to and from pods in your cluster.
- Let's create one nginx pod

```
root@manager-1: ~
apiVersion: v1
kind: Pod
metadata:
  name: network-policy-secure-pod
  labels:
    app: secure-app
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
```

Classification: Public

```
c:\ root@manager-1: ~
root@manager-1:~# vim a.yml
root@manager-1:~# kubectl create -f a.yml
pod/network-policy-secure-pod created
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE   READINESS GATES
network-policy-secure-pod   1/1    Running   0          4m11s   192.168.226.89   worker-1   <none>        <none>
```

- Now create test pod which will connect to above pod

```
c:\ root@manager-1: ~
apiVersion: v1
kind: Pod
metadata:
  name: network-policy-test-pod
spec:
  containers:
    - name: busybox
      image: radial/busyboxplus:curl
      command: ["/bin/sh", "-c", "while true; do sleep 3600; done"]
root@manager-1:~# kubectl create -f te.yaml
pod/network-policy-test-pod created
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE   READINESS GATES
network-policy-secure-pod   1/1    Running   0          10m     192.168.226.89   worker-1   <none>        <none>
network-policy-test-pod     1/1    Running   0          5s      192.168.226.91   worker-1   <none>        <none>
root@manager-1:~#
```

- We can check by login inside **test-pod** & do curl to IP of **secure-pod (10.168.226.89)**, it will work fine

```
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE   READINESS GATES
network-policy-secure-pod   1/1    Running   0          10m     192.168.226.89   worker-1   <none>        <none>
network-policy-test-pod     1/1    Running   0          5s      192.168.226.91   worker-1   <none>        <none>
root@manager-1:~# kubectl exec -it network-policy-test-pod -- sh
sh: shopt: not found
| root@network-policy-test-pod:/ ]$ curl 192.168.226.89
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
| root@network-policy-test-pod:/ ]$
```

- Create a network policy that restricts all access to the secure pod, except to and from pods which bear the allow-access: "true" label:
- We can check all network policy using below command
- Syntax: > kubectl get networkpolicies
- Syntax: > kubectl get netpol
- We can get any network policy details using below command
- Syntax:> kubectl describe networkpolicy <my-network-policy name>

Classification: Public

Classification: Public

```
root@manager-1: ~
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: my-network-policy
spec:
  podSelector:
    matchLabels:
      app: secure-app
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          allow-access: "true"
    ports:
    - protocol: TCP
      port: 80
  egress:
  - to:
    - podSelector:
        matchLabels:
          allow-access: "true"
    ports:
    - protocol: TCP
      port: 80
```

```
root@manager-1: ~
root@manager-1:~# vim np.yml
27L, 479C written
root@manager-1:~# kubectl create -f np.yml
networkpolicy.networking.k8s.io/my-network-policy created
root@manager-1:~# kubectl get netpol
NAME           POD-SELECTOR   AGE
my-network-policy  app=secure-app  7m3s
root@manager-1:~# kubectl describe netpol my-network-policy
Name:           my-network-policy
Namespace:      default
Created on:    2024-03-10 08:20:07 +0000 UTC
Labels:         <none>
Annotations:    <none>
Spec:
  PodSelector:    app=secure-app
  Allowing ingress traffic:
    To Port: 80/TCP
    From:
      PodSelector: allow-access=true
  Allowing egress traffic:
    To Port: 80/TCP
    To:
      PodSelector: allow-access=true
  Policy Types: Ingress, Egress
```

- Now we can check curl again by login in test-pod, we cannot connect to secure-pod now.

```
root@manager-1: ~
root@manager-1:~# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
network-policy-secure-pod  1/1    Running   0          57m   192.168.226.89  worker-1  <none>       <none>
network-policy-test-pod   1/1    Running   0          46m   192.168.226.91  worker-1  <none>       <none>
root@manager-1:~# kubectl exec -it network-policy-test-pod -- sh
sh: shopt: not found
| root@network-policy-test-pod:/ ]$ curl 192.168.226.89
curl: (7) Failed to connect to 192.168.226.89 port 80: Connection timed out
| root@network-policy-test-pod:/ ]$
```

- To connect to secure-pod we need to add label to test-pod & again login & check with curl command
- Syntax:> kubectl label pod network-policy-test-pod allow-access=true

Classification: Public

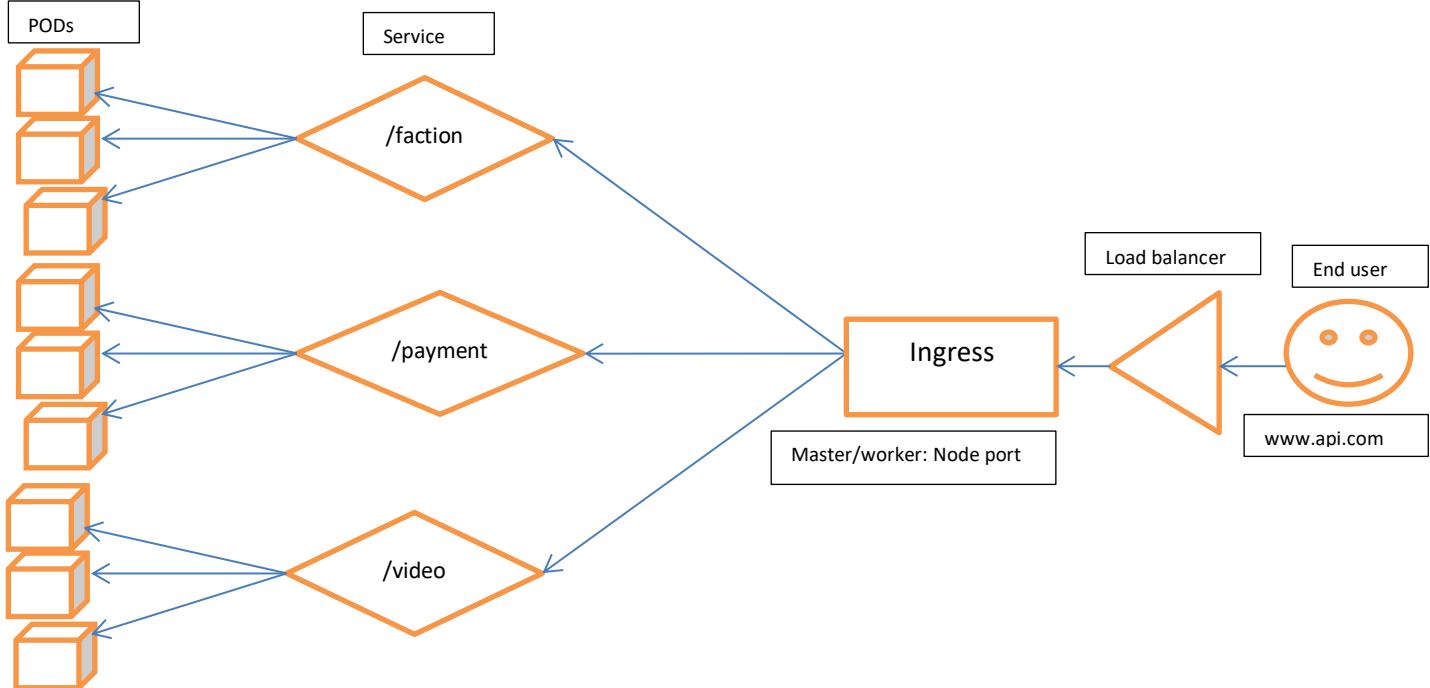
```

root@manager-1: ~
root@manager-1:~# kubectl get pods --show-labels
NAME          READY   STATUS    RESTARTS   AGE   LABELS
network-policy-secure-pod 1/1     Running   0          68m   app=secure-app
network-policy-test-pod   1/1     Running   0          57m   <none>
root@manager-1:~# kubectl label pod network-policy-test-pod allow-access=true
pod/network-policy-test-pod labeled
root@manager-1:~# kubectl get pods --show-labels
NAME          READY   STATUS    RESTARTS   AGE   LABELS
network-policy-secure-pod 1/1     Running   0          69m   app=secure-app
network-policy-test-pod   1/1     Running   0          59m   allow-access=true
root@manager-1:~# kubectl get pods --show-labels -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES   LABELS
network-policy-secure-pod 1/1     Running   0          69m   192.168.226.89   worker-1   <none>        <none>      app=secure-app
network-policy-test-pod   1/1     Running   0          59m   192.168.226.91   worker-1   <none>        <none>      allow-access=true
root@manager-1:~# kubectl exec -it network-policy-test-pod -- sh
sh: shopt: not found
[ root@network-policy-test-pod:/ ]$ curl 192.168.226.89
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style> <!-->
html { color-scheme: light-dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>
<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
[ root@network-policy-test-pod:/ ]$ 

```

56) What is ingress in Kubernetes?

- It is an API object that manages external access to the services in a cluster, typically HTTP.
- Ingress can also provide load balancing, SSL termination and name-based virtual hosting.
- Ingress exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.
- Traffic routing is controlled by rules defined on the Ingress resource
- Here is a simple example where an Ingress sends all its traffic to multiple Service:



- To use ingress we need to install software which is called **ingress controller**, there are many ingress controller available in market i.e. nginx, HA-proxy, Envoy, traefik etc.
- Best used ingress controller in industry is nginx ingress controller.
- To route traffic to specific service depends on requirement, we need to configure set of rules in ingress controller, which is called as **ingress resources/rules**
- We can install nginx ingress using below, all ingress object are install in namespace “**ingress-nginx**”
- Syntax:> `kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.3.0/deploy/static/provider/baremetal/deploy.yaml`
- We can check what all object are available in “**ingress-nginx**” namespace
- Syntax:> `kubectl get all -n ingress-nginx`
- Here we can see many pods & service are created

```
root@manager-1:~# kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.3.0/deploy/static/provider/baremetal/deploy.yaml
namespace/ingress-nginx created
serviceaccount/ingress-nginx created
serviceaccount/ingress-admission created
role.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrole.rbac.authorization.k8s.io/ingress-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-admission created
rolebinding.rbac.authorization.k8s.io/ingress-admission created
clusterrolebinding.rbac.authorization.k8s.io/ingress-nginx created
clusterrolebinding.rbac.authorization.k8s.io/ingress-admission created
configmap/ingress-nginx-controller created
service/ingress-nginx-controller created
service/ingress-nginx-controller-admission created
deployment.apps/ingress-nginx-controller created
job.batch/ingress-nginx-admission-create created
ingressClass.networking.k8s.io/nginx created
validatingwebhookconfiguration.admissionregistration.k8s.io/ingress-nginx-admission created
root@manager-1:~# kubectl get all -n ingress-nginx
NAME                                     READY   STATUS    RESTARTS   AGE
pod/ingress-admission-create-g2x84        0/1     Completed  0          46s
pod/ingress-admission-patch-p8cgd         0/1     Completed  1          46s
pod/ingress-nginx-controller-77cb5dbf4d-smgnt 1/1     Running   0          46s

NAME              TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)      AGE
service/ingress-nginx-controller   NodePort   10.106.133.51   <none>       80:31697/TCP,443:32755/TCP   46s
service/ingress-nginx-controller-admission ClusterIP  10.109.191.72   <none>       443/TCP     46s

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/ingress-nginx-controller 1/1     1           1           46s

NAME          DESIRED  CURRENT  READY   AGE
replicaset.apps/ingress-nginx-controller-77cb5dbf4d 1        1        1        46s

NAME          COMPLETIONS  DURATION   AGE
job.batch/ingress-nginx-admission-create 1/1        8s        46s
job.batch/ingress-nginx-admission-patch  1/1        8s        46s
```

- Let's take example where we have created 3 application with 3 different URL i.e. /app1, /app2 & /app3 using <https://raw.githubusercontent.com/lerndevops/kubernetes/master/6-networking/ingress/deploy-app.yml>

```
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/6-networking/ingress/deploy-app.yaml
deployment.apps/app1 created
service/app1-svc created
deployment.apps/app2 created
service/app2-svc created
deployment.apps/app3 created
service/app3-svc created
root@manager-1:~# kubectl get deploy,svc
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/app1  2/2     2           2           32s
deployment.apps/app2  2/2     2           2           32s
deployment.apps/app3  2/2     2           2           32s

NAME              TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)      AGE
service/app1-svc  ClusterIP  10.105.131.209  <none>       80/TCP      32s
service/app2-svc  ClusterIP  10.111.127.156  <none>       80/TCP      32s
service/app3-svc  ClusterIP  10.110.72.121   <none>       80/TCP      32s
service/kubernetes ClusterIP  10.96.0.1     <none>       443/TCP     12m
```

- Now create ingress rule using YML file
- Syntax:> `kubectl apply -f https://raw.githubusercontent.com/lerndevops/educka/master/6-networking/ingress/ingress-rule1.yaml`
- We can check all ingress rule using below command
- Syntax:> `kubectl get ing`

Classification: Public

```
root@manager-1: ~
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress-rule1
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - http:
    paths:
      - path: /app1
        pathType: Prefix
        backend:
          service:
            name: app1-svc
            port:
              number: 80
      - path: /app2
        pathType: Prefix
        backend:
          service:
            name: app2-svc
            port:
              number: 80
      - path: /app3
        pathType: Prefix
        backend:
          service:
            name: app3-svc
            port:
              number: 80
```

```
root@manager-1: ~
root@manager-1:~# vim ing.yml
root@manager-1:~# kubectl create -f ing.yml
ingress.networking.k8s.io/ingress-rule1 created
root@manager-1:~# kubectl get ing
NAME      CLASS   HOSTS   ADDRESS      PORTS   AGE
ingress-rule1  nginx  *       10.128.0.3  80      2m23s
root@manager-1:~#
```

- Now to access application, we need to get nginx ingress controller service node port using below command
- Syntax:> `kubectl get services -n ingress-nginx`

```
root@manager-1:~# kubectl get pods -n ingress-nginx -o wide
NAME                               READY   STATUS    RESTARTS   AGE     IP           NODE   NOMINATED-NODE   READINESS   GATES
ingress-nginx-admission-create-g2x84  0/1    Completed  0          43m    192.168.226.93  worker-1  <none>        <none>
ingress-nginx-admission-patch-p8cgd   0/1    Completed  1          43m    192.168.226.90  worker-1  <none>        <none>
ingress-nginx-controller-77cb5dbf4d-smngt  1/1    Running   0          43m    192.168.226.81  worker-1  <none>        <none>
root@manager-1:~# kubectl get services -n ingress-nginx
NAME        TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
ingress-nginx-controller   NodePort   10.106.133.51 <none>        80:31697/TCP,443:32755/TCP   44m
ingress-nginx-controller-admission ClusterIP  10.109.191.72 <none>        443/TCP    44m
```

- By using IP of master/worker node using below command & ingress node port we can access our application.
- If we will not specified any path <http://34.69.55.217:31697/> it will give 404, because in ingress controller we have not specified default path (/) rout.



Classification: Public

- So if we use ingress controller configured path as below we can see out put
- <http://34.69.55.217:31697/app1>

The image consists of three vertically stacked screenshots of a web browser. Each screenshot shows a single tab with the URL <http://34.69.55.217:31697/>. The browser interface includes a back button, forward button, refresh button, address bar, and various toolbar icons.

- Screenshot 1:** The address bar shows <http://34.69.55.217:31697/app1>. The page content says "hey there, you reached app1". Below it, a list item shows the URL <http://34.69.55.217:31697/app2>.
- Screenshot 2:** The address bar shows <http://34.69.55.217:31697/app2>. The page content says "hey there, You reached app2". Below it, a list item shows the URL <http://34.69.55.217:31697/app3>.
- Screenshot 3:** The address bar shows <http://34.69.55.217:31697/app3>. The page content says "hey there, You reached app3".

57) Why we need monitoring? How many types of monitoring available in Kubernetes?

- We need monitoring to make sure everything is running fine. If something is not working we should get alert.
- Monitoring is divide into 3 type
 - **System Level** : Monitoring CPU, RAM, speed & disk utilization
 - **Service/Process Level**: Monitoring software application which is needed to run application for ex. Docker software, DB software, and application server software.
 - **Application Level**: Monitoring application is running fine, by hitting application URL or login in application using some dummy credentials, also by using some monitoring tools.

58) How we can monitor Kubernetes?

- To monitor utilization of Pods we need to install software name as matrix server. It shows current time utilization of CPU & RAM.
- We can install matrix server using below command

```
Syntax:> kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/metrics-server/metrics-server.yaml
```
- If we want to see utilization over period of time & top utilization we need to use "Prometheus"
- Prometheus is an open-source software application used for event monitoring and alerting. It records real-time metrics in a time series database (allowing for high dimensionality) built using a HTTP pull model, with flexible queries and real-time alerting.

- The Prometheus server works on the principle of scraping, i.e., invoking the metrics endpoints of the various nodes that it is configured to monitor. It collects these metrics at regular intervals and stores them locally.
- It collects data & with the help of Grafana GUI software it displays on UI, in the form of Graph, chart etc.
- Prometheus can configure rules to trigger alerts using PromQL, alert manager will be in charge of managing alert notification, grouping, inhibition, etc.
- The alert manager component configures the receivers, gateways to deliver alert notifications.
- Grafana can pull metrics from any number of Prometheus servers and display panels and Dashboards.
- To work with Prometheus we need to install kube-state-metrics first
 - <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/cluster-role-binding.yaml>
 - <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/cluster-role.yaml>
 - <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/deployment.yaml>
 - <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/service-account.yaml>
 - <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/service.yaml>
- Now we need to install Prometheus using below YML
 - <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/prometheus.yaml>
- We can see namespace monitoring is created in which pod, service, deployment & replicaset is created.
- We can access Prometheus using service port & IP of cluster <http://146.148.36.178:30855/>

```
root@manager-1:~#
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/cluster-role-binding.yaml
clusterrolebinding.rbac.authorization.k8s.io/kube-state-metrics created
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/cluster-role.yaml
clusterrole.rbac.authorization.k8s.io/kube-state-metrics created
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/deployment.yaml
deployment.apps/kube-state-metrics created
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/service-account.yaml
serviceaccount/kube-state-metrics created
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/kube-state-metrics/standard/service.yaml
service/kube-state-metrics created
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/prometheus/prometheus.yaml
namespace/monitoring created
configmap/prom-server-conf created
clusterrole.rbac.authorization.k8s.io/prometheus created
clusterrolebinding.rbac.authorization.k8s.io/prometheus created
deployment.apps/prometheus-deployment created
service/prometheus-service created
root@manager-1:~# kubectl get all -n monitoring
NAME                      READY   STATUS    RESTARTS   AGE
pod/prometheus-deployment-7878bc4fd4-95mmr   1/1     Running   0          2m14s

NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/prometheus-service   NodePort   10.111.49.59   <none>       8080:30855/TCP   2m14s

NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/prometheus-deployment   1/1     1           1          2m14s

NAME                           DESIRED   CURRENT   READY   AGE
replicaset.apps/prometheus-deployment-7878bc4fd4   1         1         1        2m14s
```

- But this GUI is not so useful, so we need to install **Grafana**
- kubectl create -f <https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/prometheus/grafana.yaml>
- This will also create pod, service, deployment & replicaset.
- We can access this **Grafana** GUI using service port & cluster IP <http://146.148.36.178:32346/login>
- We can login using default credentials
 - user : admin
 - Password: admin

Classification: Public

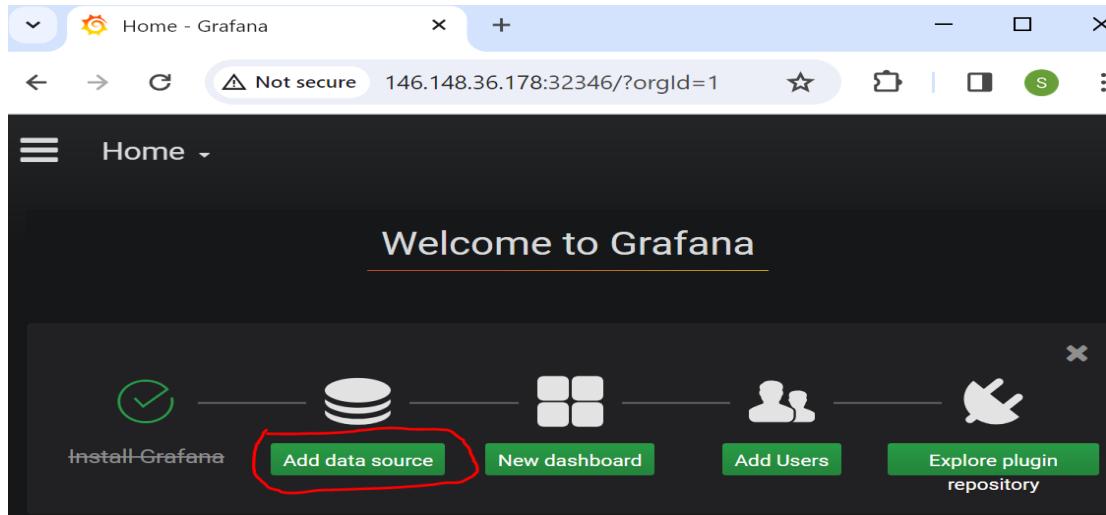
```
root@manager-1: ~
root@manager-1:~# kubectl create -f https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/prometheus/grafana.yml
deployment.apps/grafana created
service/grafana-ui-svc created
Error from server (AlreadyExists): error when creating "https://raw.githubusercontent.com/lerndevops/kubernetes/master/9-monitoring/prometheus/grafana.yml": exists
root@manager-1:~# kubectl get all -n monitoring
NAME                                         READY   STATUS    RESTARTS   AGE
pod/grafana-65cc8d46dd-zdh6h                1/1     Running   0          23s
pod/prometheus-deployment-7878bc4fd4-95mmr   1/1     Running   0          19m

NAME           TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
service/grafana-ui-svc  NodePort   10.99.141.51   <none>          3000:32346/TCP  23s
service/prometheus-service  NodePort   10.111.49.59   <none>          8080:30855/TCP  19m

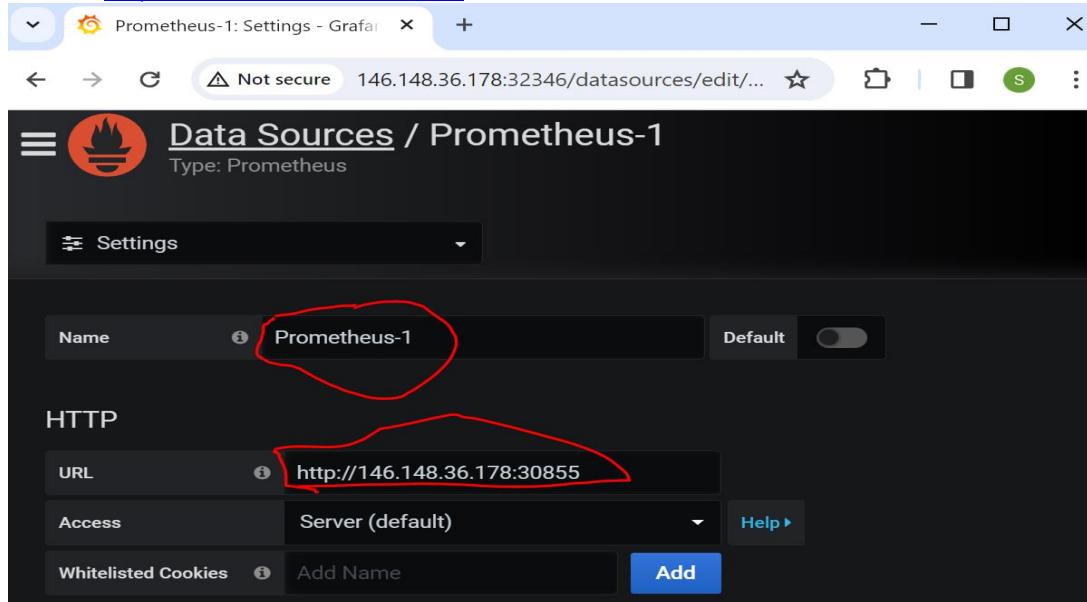
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/grafana   1/1     1           1           23s
deployment.apps/prometheus-deployment  1/1     1           1           19m

NAME          DESIRED  CURRENT  READY   AGE
replicaset.apps/grafana-65cc8d46dd       1        1        1        23s
replicaset.apps/prometheus-deployment-7878bc4fd4  1        1        1        19m
root@manager-1:~#
```

- After login we need to configure data source as **Prometheus** in **Grafana**



- Go to add data source → select Prometheus → configure URL for Prometheus <http://146.148.36.178:30855> → click on save & test



Classification: Public

Classification: Public

- Now we can create new dashboard using Dashboard → click on New Dashboard → Add Query → copy query from **Prometheus & paste**

The screenshot shows the 'New dashboard' screen in Grafana. At the top, there is a chart titled 'Panel Title' showing data over a 6-hour period. Below the chart is a query editor. The 'Query' dropdown is set to 'Prometheus-1'. The 'Metrics' dropdown has 'kube_namespace_created' selected. A red box highlights the 'Prometheus-1' dropdown, the 'Add Query' button, and the 'kube_namespace_created' dropdown.

- We can create use ready-made dashboard by taking from **Grafana** official web site <https://grafana.com/grafana/dashboards/6417-kubernetes-cluster-prometheus/>
- Copy ID of Dashboard → go to Garafana → New Dashboard → Import Dashboard → paste ID → select name of prometheus → click on import

The screenshot shows the 'Import' dialog in Grafana. It displays information about a dashboard published by 'sekka1' on June 7, 2018. The 'Options' section allows setting the 'Name' (Kubernetes Cluster (Prometheus)), 'Folder' (General), and 'Unique identifier (uid)' (value set). A dropdown for 'prometheus' is set to 'Prometheus-1'. At the bottom are 'Import' and 'Cancel' buttons.

- It will give you GUI for your Kurnet cluster

The screenshot shows the imported 'Kubernetes Cluster (Prometheus)' dashboard. It features several panels: 'Cluster Health' with 'Cluster Pod Usage' (6.364%), 'Cluster CPU Usage' (30.00%), 'Cluster Memory Usage' (4.639%), and 'Cluster Disk Usage' (N/A); 'Deployments' with tables for 'Deployment Replicas - Up To Date' and 'Deployment Replicas - Updated'; and a 'Deployment Replicas' summary panel showing 8 available replicas.

- We can also give dashboard JSON to create Dashboard

Classification: Public

59) What is Kubernetes Security?

- Kubernetes Security is based on the 4C's of cloud native security: Cloud, Cluster, Container, and Code
- **Cloud (or Corporate Datacenter/Colocation facility):** The underlying physical infrastructure is the basis of Kubernetes security. Whether the cluster is built on one's own datacenter or a cloud provider, basic cloud provider (or physical security) best practices must be observed.
- **Cluster:** Securing a Kubernetes cluster involves both the configurable components such as the Kubernetes API and security of all the applications that are part of the cluster. Since most cloud-native applications are designed around micro services and APIs, applications are only as secure as the weakest link in the chain of services that comprise the entire application.
- **Container:** Container design best practices consist of: starting with the smallest code base possible (excluding unnecessary libraries or functions), avoiding granting unnecessary privileges to users in the container, and ensuring that containers are scanned for vulnerabilities at build time.
- **Code:** Code presents a major attack surface for any Kubernetes environment. Simple policies such as encrypting TCP using TLS handshakes, not exposing unused ports, scanning, and testing regularly can help prevent security issues from arising in a production environment.